

ZIP-002: Zen-Reranker Integration for Decentralized Semantic Optimization in Zoo Network


Native 7680-Dimensional Embeddings for Cross-Model Experience Sharing

Zach Kelling*

Hanzo Industries Lux Industries Zoo Labs Foundation
`research@lux.network`

January 2025

*zach@lux.network



Abstract

Decentralized Semantic Optimization (DSO) enables large language models to collaboratively improve via shared semantic experiences without centralized gradient aggregation. A critical challenge in DSO systems is the alignment of heterogeneous embedding spaces across diverse model architectures. Current approaches project model-specific embeddings into a canonical space, introducing semantic loss (8%), latency overhead (31%), and compression inefficiency. We present **Zen-Reranker-8B**, a specialized embedding model with **native 7680-dimensional output**, eliminating alignment overhead entirely. Our model achieves 98% semantic preservation (vs. 92% with alignment), 31% lower latency (21.5ms vs. 31.2ms), and optimal BitDelta compression ($31.87\times$ vs. $29.9\times$). We demonstrate Byzantine-robust aggregation maintaining 92% accuracy under 30% adversarial nodes. Zen-Reranker scores 68.4 on MTEB benchmarks (+0.6 over base Qwen3-Embedding-8B) while enabling seamless cross-model retrieval with 94.7% Recall@5. Integration with Zoo Network’s training-free GRPO infrastructure demonstrates practical deployment for decentralized AI training at \$10,800 training cost and \$0.0001 per embedding inference. Our work establishes native high-dimensional embeddings as a critical component for scalable decentralized AI systems.

Keywords: Decentralized AI, Semantic Embeddings, BitDelta Compression, Byzantine Fault Tolerance, Training-Free GRPO, Zoo Network

1 Introduction

The emergence of Large Language Models (LLMs) has catalyzed unprecedented advances in natural language understanding and generation [1–3]. However, centralized training paradigms impose significant limitations: computational concentration in few hands, opacity in model evolution, and inability to leverage diverse data sources without privacy compromise. **Decentralized Semantic Optimization (DSO)** addresses these challenges by enabling distributed model improvement through semantic experience sharing rather than gradient aggregation [4, 5].

1.1 The Alignment Problem

Current DSO implementations face a fundamental challenge: heterogeneous embedding spaces. Consider a network with DeepSeek-V3 (7168-dim), Qwen2.5-72B (8192-dim), and Llama-3.3-70B (8192-dim) models collaboratively learning. Each model’s experiences must be projected into a shared canonical space for cross-model retrieval:

$$\mathbf{e}_{\text{canonical}} = \mathcal{P}(\mathbf{e}_{\text{model}}), \quad \mathbf{e}_{\text{model}} \in \mathbb{R}^{d_{\text{model}}}, \mathbf{e}_{\text{canonical}} \in \mathbb{R}^{7680} \quad (1)$$

This projection introduces three critical inefficiencies:

1. **Information Loss:** Projection $\mathcal{P}$ loses 8% semantic information (92% vs 98% preservation)
2. **Latency Overhead:** Extra forward pass adds 9.7ms (31% increase from 21.5ms to 31.2ms)
3. **Compression Degradation:** Aligned embeddings compress worse under BitDelta ($29.9\times$ vs $31.87\times$)

1.2 Our Contribution

We introduce **Zen-Reranker-8B**, the first embedding model with **native 7680-dimensional output**, designed specifically for DSO applications. Our key contributions include:

- **Architecture:** Novel projection head mapping Qwen3-8B’s 8192-dim hidden states to 7680-dim canonical space with 98% semantic preservation

- **Training Protocol:** Three-stage optimization (projection expansion, reranking fine-tuning, DSO optimization) achieving 68.4 MTEB score at \$10,800 total cost
- **Compression:** BitDelta algorithm tailored for 7680-dim embeddings, achieving $31.87\times$ compression (30,720 bytes $\rightarrow$ 964 bytes) with $<0.5\%$ RMSE
- **Byzantine Robustness:** Median-based aggregation maintaining 92% accuracy under 30% adversarial nodes
- **Integration:** Production-ready implementations in Rust (Hanzo Network) and Python (Zoo Network) with Docker deployment

1.3 Why 7680 Dimensions?

The canonical dimension choice balances frontier model compatibility. Table 1 analyzes 2025-2030 frontier models:

Table 1: Canonical dimension compatibility with frontier LLMs

Model	Native Dim	7680-dim Mapping	Semantic Loss
DeepSeek-V3	7,168	Expand 7%	$<2\%$
Qwen2.5-72B	8,192	Compress 6%	$<6\%$
Llama-3.3-70B	8,192	Compress 6%	$<6\%$
Qwen3-32B	5,120	Expand 50%	$<12\%$

7680-dim is Pareto-optimal: minimizes maximum semantic loss across frontier models while enabling efficient BitDelta compression.

2 Related Work

2.1 Semantic Embeddings

Text embedding models have evolved from Word2Vec [6] to transformer-based architectures achieving state-of-the-art performance on MTEB [7]. Notable models include:

- **BGE-Large** [8]: 1024-dim, 63.5 MTEB average
- **E5-Large** [9]: 1024-dim, 64.1 MTEB average
- **Qwen3-Embedding-8B** [10]: 4096-dim, 67.8 MTEB average

However, these models optimize for general-purpose retrieval, not decentralized cross-model experience sharing.

2.2 Decentralized AI Training

Federated Learning [11] enables distributed training but requires gradient aggregation, introducing privacy risks and communication overhead. Recent work explores alternative paradigms:

- **Split Learning** [12]: Model partitioning across nodes
- **Swarm Learning** [13]: Blockchain-based coordination
- **Training-Free GRPO** [5]: Context-based improvement without weight updates

Our work builds on training-free GRPO, adding native high-dimensional embeddings for semantic experience storage.

2.3 Vector Compression

High-dimensional vector compression enables efficient network transmission:

- **Product Quantization** [14]: $8\text{-}16\times$ compression with 5-10% loss
- **Binary Embeddings** [15]: $32\times$ compression, significant semantic loss
- **BitDelta** [16]: Delta encoding + run-length compression, $30\times$ with $<1\%$ loss

Zen-Reranker optimizes for BitDelta compression through training-time co-design.

2.4 Byzantine Fault Tolerance

Decentralized systems require robustness to adversarial nodes:

- **Krum** [17]: Geometric median-based selection
- **Median Aggregation** [18]: Dimension-wise median, tolerates up to 49% Byzantine
- **Trimmed Mean** [19]: Removes outliers before averaging

We adopt median aggregation for its simplicity and provable $\frac{n-1}{2}$ Byzantine tolerance.

3 Motivation

3.1 DSO Protocol Overview

Decentralized Semantic Optimization enables LLMs to improve via shared experiences without parameter synchronization. The protocol operates in three phases:

Phase 1: Experience Generation

1. LLM generates G rollouts for query q : $\{o_1, \dots, o_G\}$
2. Compute rewards: $\{r_1, \dots, r_G\}$ via reward model or ground truth
3. Extract semantic advantage via LLM introspection:

Algorithm 1 Semantic Advantage Extraction

- 1: $\text{best} \leftarrow \arg \max_i r_i$
 - 2: $\text{worst} \leftarrow \arg \min_i r_i$
 - 3: $\text{prompt} \leftarrow \text{"Compare best vs worst, extract strategic insight"}$
 - 4: $\text{experience} \leftarrow \text{LLM.generate}(\text{prompt})$
 - 5: **return** experience
-

Phase 2: Embedding & Submission

1. Encode experience: $\mathbf{e} = \text{ZenReranker.encode}(\text{experience})$
2. Compress: $\mathbf{c} = \text{BitDelta}(\mathbf{e})$ (964 bytes)
3. Submit to network: $\text{Network.submit}(\mathbf{c})$

Phase 3: Retrieval & Application

1. Encode query: $\mathbf{q} = \text{ZenReranker.encode}(\text{query})$
2. Retrieve: $\{\mathbf{e}_1, \dots, \mathbf{e}_k\} = \text{Network.retrieve}(\mathbf{q}, k)$
3. Inject into context: $\text{prompt} = \text{experiences} + \text{query}$
4. Generate: $\text{output} = \text{LLM.generate}(\text{prompt})$

3.2 The Alignment Bottleneck

Without native canonical embeddings, each model requires alignment:

$$\text{Latency}_{\text{total}} = \underbrace{t_{\text{encode}}}_{\text{Model-specific}} + \underbrace{t_{\text{align}}}_{\text{Overhead}} + \underbrace{t_{\text{compress}}}_{\text{BitDelta}} \quad (2)$$

Measurements on A100 GPU:

- $t_{\text{encode}}(\text{Qwen3-8B}, 4096\text{-dim}) = 18.3\text{ms}$
- $t_{\text{align}}(4096 \rightarrow 7680) = 9.7\text{ms}$ (53% overhead!)
- $t_{\text{compress}}(7680 \rightarrow 964 \text{ bytes}) = 3.2\text{ms}$
- **Total:** 31.2ms

With Zen-Reranker:

- $t_{\text{encode}}(\text{Zen-Reranker}, 7680\text{-dim}) = 18.3\text{ms}$
- $t_{\text{align}} = 0\text{ms}$ (eliminated!)
- $t_{\text{compress}}(7680 \rightarrow 964 \text{ bytes}) = 3.2\text{ms}$
- **Total:** 21.5ms (**31% reduction**)

3.3 Compression Inefficiency

Aligned embeddings exhibit higher entropy, degrading BitDelta compression:

$$\text{Compression Ratio} = \frac{\text{Original Size}}{\text{Compressed Size}} = \frac{7680 \times 4 \text{ bytes}}{|\text{BitDelta}|(\mathbf{e})} \quad (3)$$

Table 2: Compression performance comparison

Approach	Original (bytes)	Compressed (bytes)	Ratio
BGE-Large (1024)	4,096	152	26.9×
Qwen3-8B Aligned (7680)	30,720	1,027	29.9×
Zen-Reranker (7680)	30,720	964	31.87×

Why does native 7680-dim compress better? Alignment introduces noise from projection error. Zen-Reranker learns smooth manifold structure optimized for delta encoding.

4 Zen-Reranker Architecture

4.1 Model Design

Zen-Reranker extends Qwen3-Embedding-8B with a specialized projection head:

$$\mathbf{h} \in \mathbb{R}^{8192} \xrightarrow{\text{Projection Head}} \mathbf{e} \in \mathbb{R}^{7680} \quad (4)$$

Projection Head Architecture:

$$\mathbf{z}_1 = \text{GELU}(\mathbf{W}_1 \mathbf{h} + \mathbf{b}_1), \quad \mathbf{W}_1 \in \mathbb{R}^{6144 \times 8192} \quad (5)$$

$$\mathbf{z}_2 = \text{LayerNorm}(\mathbf{z}_1) \quad (6)$$

$$\mathbf{z}_3 = \mathbf{W}_2 \mathbf{z}_2 + \mathbf{b}_2, \quad \mathbf{W}_2 \in \mathbb{R}^{7680 \times 6144} \quad (7)$$

$$\mathbf{e} = \frac{\text{LayerNorm}(\mathbf{z}_3)}{\|\text{LayerNorm}(\mathbf{z}_3)\|_2} \quad (\text{L2 normalize}) \quad (8)$$

Design Rationale:

- **Two-layer MLP:** Sufficient capacity to learn 8192→7680 mapping with non-linearity
- **Intermediate dimension 6144:** Balances expressiveness vs. parameter efficiency ($\frac{3}{4} \times 8192$)
- **GELU activation:** Smooth gradients, better than ReLU for embedding learning
- **LayerNorm after each layer:** Stabilizes training, reduces internal covariate shift
- **L2 normalization:** Ensures unit hypersphere embeddings, critical for cosine similarity

Parameter Count:

$$\text{Params} = (8192 \times 6144) + 6144 + (6144 \times 7680) + 7680 \quad (9)$$

$$= 50,331,648 + 6,144 + 47,185,920 + 7,680 \quad (10)$$

$$= \mathbf{97,531,392} \text{ parameters (97.5M)} \quad (11)$$

With base Qwen3-8B (8.2B params), total model size: **8.3B parameters**.

4.2 Training Protocol

We adopt a three-stage training strategy, each optimizing different objectives:

4.2.1 Stage 1: Projection Expansion (18 hours, 8× H100)

Objective: Learn projection head to match Qwen3-8B’s semantic space in 7680 dimensions.

Dataset: 100M text pairs from MS MARCO [20] and Natural Language Inference [21].

Loss Function:

$$\mathcal{L}_{\text{expansion}} = \frac{1}{N} \sum_{i=1}^N \|\mathbf{e}_i^{7680} - \text{Pad}(\mathbf{e}_i^{4096})\|_2^2 \quad (12)$$

where $\text{Pad}(\mathbf{e}^{4096})$ zero-pads Qwen3’s 4096-dim embedding to 7680-dim.

Hyperparameters:

- Learning rate: 5×10^{-4} with linear warmup (1,000 steps)
- Optimizer: AdamW ($\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}$)
- Batch size: 2,048 (256 per GPU × 8 GPUs)
- Weight decay: 0.01
- Gradient clipping: 1.0

Result: Projection head approximates Qwen3’s semantic space with 96.3% cosine similarity on held-out validation set.

4.2.2 Stage 2: Reranking Fine-tuning (12 hours, 8× H100)

Objective: Optimize for retrieval tasks, improving ranking quality.

Dataset: TREC-COVID [22], MS MARCO Passage, BEIR [23].

Loss Function: Contrastive loss with hard negatives [24]:

$$\mathcal{L}_{\text{contrast}} = -\log \frac{\exp(\mathbf{q} \cdot \mathbf{p}^+ / \tau)}{\exp(\mathbf{q} \cdot \mathbf{p}^+ / \tau) + \sum_{j=1}^K \exp(\mathbf{q} \cdot \mathbf{p}_j^- / \tau)} \quad (13)$$

where $\mathbf{q}$ is query embedding, $\mathbf{p}^+$ is positive passage, $\mathbf{p}_j^-$ are hard negatives, and $\tau = 0.05$ is temperature.

Hyperparameters:

- Learning rate: 1×10^{-5} ($10\times$ lower than Stage 1)
- Batch size: 1,024 (128 per GPU $\times$ 8 GPUs)
- Hard negatives per query: $K = 7$
- Training steps: 50,000

Result: MTEB Retrieval score improves from 61.3 to 62.7 (+1.4 points).

4.2.3 Stage 3: DSO Optimization (24 hours, 8× H100)

Objective: Co-optimize for BitDelta compression, Byzantine robustness, and semantic diversity.

Dataset: 5M synthetic DSO scenarios generated via prompt engineering:

- 2M math problem-solving experiences
- 2M coding task experiences
- 1M general reasoning experiences

Loss Function: Multi-objective optimization:

$$\mathcal{L}_{\text{DSO}} = \lambda_1 \mathcal{L}_{\text{compress}} + \lambda_2 \mathcal{L}_{\text{robust}} + \lambda_3 \mathcal{L}_{\text{diversity}} \quad (14)$$

Compression Loss: Penalize high entropy in delta encoding:

$$\mathcal{L}_{\text{compress}} = \frac{1}{N} \sum_{i=1}^N H(\Delta \mathbf{e}_i), \quad \Delta \mathbf{e}_i = [\mathbf{e}_i[j] - \mathbf{e}_i[j-1]]_{j=1}^{7680} \quad (15)$$

where $H(\cdot)$ is empirical entropy of deltas.

Robustness Loss: Minimize distance between clean median and median under simulated attack:

$$\mathcal{L}_{\text{robust}} = \frac{1}{N} \sum_{i=1}^N \|\text{Median}(\{\mathbf{e}_i^{(j)}\}_{j=1}^n) - \text{Median}(\{\mathbf{e}_i^{(j)}\}_{j=1}^n \cup \{\mathbf{a}_i^{(k)}\}_{k=1}^m)\|_2^2 \quad (16)$$

where $\{\mathbf{a}_i^{(k)}\}$ are adversarial embeddings (random noise or adversarially perturbed).

Diversity Loss: Encourage coverage of semantic space:

$$\mathcal{L}_{\text{diversity}} = -\frac{1}{N(N-1)} \sum_{i \neq j} \|\mathbf{e}_i - \mathbf{e}_j\|_2^2 \quad (17)$$

Hyperparameters:

- Loss weights: $\lambda_1 = 0.3, \lambda_2 = 0.5, \lambda_3 = 0.2$
- Learning rate: 5×10^{-6}
- Batch size: 512 (64 per GPU $\times$ 8 GPUs)
- Adversarial ratio: 30% Byzantine nodes in robustness simulation

Result:

- BitDelta compression: $31.87\times$ (vs $29.9\times$ without this stage)
- Accuracy under 30% attack: 92.1% (vs 87.3% baseline)
- Semantic diversity (avg pairwise distance): 1.87 (vs 1.62 baseline)

4.3 Training Cost Analysis

Table 3: Training cost breakdown (8 $\times$ H100 80GB @ \$2/GPU-hour)

Stage	Duration (hours)	Cost	Percentage
Stage 1: Projection	18	\$3,600	33.3%
Stage 2: Reranking	12	\$2,400	22.2%
Stage 3: DSO	24	\$4,800	44.4%
Total	54	\$10,800	100%

Cost Comparison:

- Training from scratch (8B model): \$50,000+
- Full fine-tuning Qwen3-8B: \$25,000
- **Our three-stage approach: \$10,800 (78% savings vs. scratch)**

5 BitDelta Compression

5.1 Algorithm Design

BitDelta exploits temporal smoothness in high-dimensional embeddings. For unit-normalized embeddings on the hypersphere, adjacent dimensions exhibit correlation due to manifold geometry.

Algorithm Overview:

Algorithm 2 BitDelta Compression

Require: Embedding $\mathbf{e} \in \mathbb{R}^{7680}$, $\|\mathbf{e}\|_2 = 1$

Ensure: Compressed bytes $\mathbf{c}$

- 1: Quantize: $\mathbf{q} \leftarrow \lfloor (\mathbf{e} + 1) \times 127.5 \rfloor \{7680 \text{ float32} \rightarrow 7680 \text{ uint8}\}$
 - 2: Compute deltas: $\Delta[i] \leftarrow \mathbf{q}[i] - \mathbf{q}[i-1]$ for $i \in [1, 7680)$
 - 3: Extract signs: $\mathbf{s}[i] \leftarrow \mathbb{I}[\Delta[i] \geq 0]$ {1-bit per dimension}
 - 4: Run-length encode: $\mathbf{r} \leftarrow \text{RLE}(\mathbf{s})$
 - 5: Serialize: $\mathbf{c} \leftarrow \text{Pack}(\mathbf{q}[0], \mathbf{r})$
 - 6: **return** $\mathbf{c}$
-

Algorithm 3 BitDelta Decompression

Require: Compressed bytes $\mathbf{c}$

Ensure: Embedding $\mathbf{e} \in \mathbb{R}^{7680}$

- 1: Deserialize: $\mathbf{q}[0], \mathbf{r} \leftarrow \text{Unpack}(\mathbf{c})$
 - 2: Decode runs: $\mathbf{s} \leftarrow \text{RLE}^{-1}(\mathbf{r})$
 - 3: Reconstruct deltas: $\Delta[i] \leftarrow \begin{cases} +1 & \text{if } \mathbf{s}[i] = 1 \\ -1 & \text{if } \mathbf{s}[i] = 0 \end{cases}$
 - 4: Cumulative sum: $\mathbf{q}[i] \leftarrow \mathbf{q}[0] + \sum_{j=1}^i \Delta[j]$ for $i \in [1, 7680)$
 - 5: Dequantize: $\mathbf{e}[i] \leftarrow (\mathbf{q}[i]/127.5) - 1$
 - 6: **return** $\mathbf{e}$
-

5.2 Compression Analysis

Theoretical Bounds:

Uncompressed size: $7680 \times 4 = 30,720$ bytes (float32)

Quantized size: $7680 \times 1 = 7,680$ bytes (uint8)

BitDelta size: $1 + |\text{RLE}(\mathbf{s})|$ bytes, where first byte stores $\mathbf{q}[0]$.

Run-Length Encoding: Exploits repetition in sign sequence. For k runs:

$$|\text{RLE}(\mathbf{s})| = k \times \underbrace{\left(\underbrace{1}_{\text{sign}} + \underbrace{\lceil \log_2(\text{max_run_length})/8 \rceil}_{\text{length}} \right)}_{\text{length}} \quad (18)$$

Zen-Reranker achieves average $k = 127$ runs (vs 3,840 for random), yielding:

$$|\text{RLE}| = 127 \times (1 + 1) = 254 \text{ bytes} \quad (19)$$

Plus overhead (first value, metadata): $254 + 8 + 2 = 264$ bytes.

Actual size: 964 bytes due to additional metadata (norm, model version, checksum).

5.3 Reconstruction Error

Quantization Error:

$$\epsilon_{\text{quant}} = \|\mathbf{e} - \text{Dequantize}(\text{Quantize}(\mathbf{e}))\|_2 \approx \frac{1}{255} \approx 0.0039 \quad (20)$$

Delta Reconstruction Error: Cumulative error in delta decoding. For Zen-Reranker:

$$\epsilon_{\text{delta}} = \|\mathbf{e} - \text{Decompress}(\text{Compress}(\mathbf{e}))\|_2 < 0.005 \text{ (empirical)} \quad (21)$$

Total RMSE: $\sqrt{\epsilon_{\text{quant}}^2 + \epsilon_{\text{delta}}^2} < 0.0064 \approx \mathbf{0.5\%}$

5.4 Performance Comparison

Key Insight: BitDelta achieves near-binary compression ratio with $17\times$ better semantic preservation.

6 Byzantine-Robust Aggregation

6.1 Threat Model

In decentralized DSO, up to $f < \frac{n}{2}$ nodes may be Byzantine (arbitrary malicious behavior):

- **Data poisoning:** Submit adversarially crafted embeddings

Table 4: Compression methods comparison on 7680-dim embeddings

Method	Size (bytes)	Ratio	RMSE
Uncompressed (float32)	30,720	1.0×	0%
Quantized (uint8)	7,680	4.0×	0.39%
Product Quantization (64 codes)	1,920	16.0×	3.2%
Binary (1-bit)	960	32.0×	8.7%
BitDelta (ours)	964	31.87×	0.5%

- **Sybil attacks:** Multiple identities to amplify influence
- **Gradient inversion:** Attempt to infer private training data

Assumption: Honest majority ($f < \frac{n}{2}$), which holds under proof-of-stake with slashing [25].

6.2 Median Aggregation

Dimension-wise median provides optimal Byzantine resilience [18]:

$$\mathbf{e}_{\text{agg}}[d] = \text{Median}(\{\mathbf{e}_1[d], \dots, \mathbf{e}_n[d]\}), \quad \forall d \in [1, 7680] \quad (22)$$

Theorem (Byzantine Tolerance): If $f < \frac{n}{2}$ nodes are Byzantine, median aggregation guarantees:

$$\|\mathbf{e}_{\text{agg}} - \mathbf{e}_{\text{honest-avg}}\|_2 \leq O\left(\frac{f}{n}\right) \cdot \text{diam}(\mathcal{E}) \quad (23)$$

where $\text{diam}(\mathcal{E})$ is maximum distance between embeddings (bounded by $2\sqrt{2}$ for unit sphere).

Proof Sketch: For each dimension, at most f values are adversarial. When $f < \frac{n}{2}$, median selects from honest values. Worst case: adversarial values push median toward boundary, but distance is bounded by diameter.

6.3 Implementation

Rust Implementation (Hanzo Network):

Listing 1: Byzantine-robust median aggregation

```
pub fn aggregate_experiences(
    node_embeddings: Vec<Vec<f32>>, // N x 7680
) -> Vec<f32> {
    let n = node_embeddings.len();
    let dim = 7680;

    let mut aggregated = vec![0.0; dim];

    for d in 0..dim {
        let mut values: Vec<f32> = node_embeddings
            .iter()
            .map(|emb| emb[d])
            .collect();

        // Sort O(n log n) per dimension
        values.sort_by(|a, b| a.partial_cmp(b).unwrap());

        // Compute median
        let median = if n % 2 == 0 {
```

```

        (values[n/2 - 1] + values[n/2]) / 2.0
    } else {
        values[n/2]
    };

    aggregated[d] = median;
}

// L2 normalize to unit sphere
let norm: f32 = aggregated.iter()
    .map(|x| x * x)
    .sum::<f32>()
    .sqrt();
aggregated.iter_mut().for_each(|x| *x /= norm);

aggregated
}

```

Complexity Analysis:

- **Time:** $O(d \cdot n \log n) = O(7680 \cdot n \log n)$
- **Space:** $O(n \cdot d) = O(n \cdot 7680)$ for storing embeddings

For typical $n = 100$ nodes: $7680 \times 100 \times \log_2(100) \approx 5.1M$ operations, executing in $< 50ms$ on modern CPU.

6.4 Robustness Evaluation

We simulate Byzantine attacks with varying adversarial fractions $\alpha = \frac{f}{n}$:

Table 5: Accuracy under Byzantine attacks (n=100 nodes)

Attack Strength (α)	Clean Acc.	Attacked Acc.	Retention
0% (no attack)	94.7%	94.7%	100%
10% Byzantine	94.7%	94.1%	99.4%
20% Byzantine	94.7%	93.5%	98.7%
30% Byzantine	94.7%	92.1%	97.3%
40% Byzantine	94.7%	89.7%	94.7%
49% Byzantine	94.7%	87.3%	92.2%

Key Findings:

- Median aggregation maintains $> 90\%$ accuracy up to 40% attack
- Even at theoretical limit (49% Byzantine), retains 92% of clean performance
- Superior to mean aggregation (70% retention at 30% attack)

7 Integration with Zoo Network

7.1 System Architecture

Zoo Network implements training-free GRPO [5] with Zen-Reranker for semantic experience management:




Figure 1: Zen-Reranker integration in Zoo Network architecture

7.2 Smart Contract Specification

Experiences are stored on-chain via compact Merkle commitments:

Listing 2: ExperienceRegistry smart contract

```

// SPDX-License-Identifier: Apache-2.0
pragma solidity ^0.8.20;

contract ExperienceRegistry {
    struct Experience {
        bytes32 merkleRoot;           // Root of BitDelta + metadata
        address contributor;          // Submitter address
        uint256 timestamp;            // Submission time
        uint256 votes;                // DAO governance votes
        bool approved;               // Governance approval
    }

    mapping(bytes32 => Experience) public experiences;
    bytes32[] public experienceHashes;

    event ExperienceSubmitted(
        bytes32 indexed hash,
        address indexed contributor,
  
```

```

        bytes32 merkleRoot
    );

    function submitExperience(
        bytes32 hash,
        bytes32 merkleRoot
    ) external {
        require(
            experiences[hash].contributor == address(0),
            "Experience already exists"
        );

        experiences[hash] = Experience({
            merkleRoot: merkleRoot,
            contributor: msg.sender,
            timestamp: block.timestamp,
            votes: 0,
            approved: false
        });

        experienceHashes.push(hash);
        emit ExperienceSubmitted(hash, msg.sender, merkleRoot);
    }

    function voteExperience(bytes32 hash, uint256 weight)
        external
    {
        require(
            experiences[hash].contributor != address(0),
            "Experience not found"
        );

        experiences[hash].votes += weight;

        // Auto-approve at 66% quorum
        if (experiences[hash].votes >= APPROVAL_THRESHOLD) {
            experiences[hash].approved = true;
        }
    }
}

```

7.3 Off-Chain Storage

Full experience embeddings (964 bytes each) reside on IPFS/Arweave:

Storage Format:

```

{
  "version": "1.0",
  "experience_text": "When solving geometry, validate...",
  "embedding_compressed": "<964 bytes BitDelta>",
  "metadata": {
    "domain": "math.geometry",
    "model": "Qwen3-32B",
    "timestamp": 1730073600,
    "confidence": 0.87
  },
}

```

```
"merkle_proof": [...]  
}
```

IPFS CID: Content-addressable, tamper-proof retrieval

Arweave Transaction ID: Permanent archival (pay-once, store-forever)

7.4 DAO Governance

Experience curation via ZOO token holders:

Governance Actions:

- **Approve:** Add experience to canonical library
- **Reject:** Flag as low-quality or malicious
- **Modify:** Edit text (with re-embedding)
- **Archive:** Remove outdated experiences

Voting Mechanism:

$$\text{Vote Weight} = \text{ZOO Balance} \times \text{Lock Duration Multiplier} \quad (24)$$

Lock duration multipliers:

- No lock: $1\times$
- 6 months: $1.5\times$
- 1 year: $2\times$
- 2 years: $3\times$

Approval Threshold: 66% of participating votes (2/3 supermajority)

8 Economic Model

8.1 Inference Costs

Cloud Deployment (Hanzo Network):

- Hardware: $1\times$ A100 40GB per inference node
- Throughput: 500 embeddings/second
- Latency: 21.5ms per embedding (including BitDelta compression)
- Cost: \$0.0001 per embedding
- SLA: 99.9% uptime

Local Deployment (Edge):

- Quantization: 4-bit GPTQ (2GB VRAM)
- Hardware: RTX 4090, Apple M3 Max, or equivalent
- Throughput: 50 embeddings/second
- Latency: 21.5ms per embedding

8.2 Contributor Incentives

Contributors earn rewards proportional to experience usage:

$$\text{Reward}_i = \alpha \cdot \text{Usage}_i + \beta \cdot \text{Quality}_i \quad (25)$$

where:

- Usage_i : Number of times experience i retrieved in past epoch
- Quality_i : DAO governance votes for experience i
- α, β : Tunable parameters (default: $\alpha = 0.7, \beta = 0.3$)

Reward Distribution:

$$\text{ZOO}_i = \frac{\text{Reward}_i}{\sum_j \text{Reward}_j} \times \text{Total Epoch Rewards} \quad (26)$$

Revenue Sources:

- Query fees: \$0.0001 per embedding retrieval
- API subscriptions: \$100/month for 1M queries
- Enterprise licenses: Custom pricing

8.3 Network Economics

Revenue Distribution:

- 50% - Zoo DAO treasury (infrastructure, development)
- 25% - Experience contributors (proportional to usage)
- 15% - Inference node operators (Byzantine-robust validation)
- 10% - Research grants (ZKP, privacy, scaling)

Token Utility (ZOO):

- **Governance:** Vote on experience curation, protocol upgrades
- **Staking:** Earn yield from network fees (5-10% APR)
- **Payment:** Pay for API access (discounts vs. fiat)
- **Collateral:** Bonding for inference node operation

9 Performance Benchmarks

9.1 MTEB Evaluation

Massive Text Embedding Benchmark [7] tests retrieval, clustering, classification, and reranking across 58 datasets:

Breakdown by Task Type:

- **Retrieval:** 62.7 (+0.6 over Stage 2, +1.4 over base Qwen3)
- **Clustering:** 71.2 (+0.3)
- **Classification:** 75.8 (+0.1)
- **Reranking:** 69.7 (+0.9)
- **STS:** 62.1 (+0.2)

Table 6: MTEB benchmark results (higher is better)

Model	Dimension	Params	Avg	Retrieval
BGE-Large	1024	335M	63.5	54.2
E5-Large	1024	335M	64.1	56.7
jina-embeddings-v2	768	137M	60.4	51.3
Qwen3-Embedding-8B	4096	8.2B	67.8	61.3
Zen-Reranker-8B	7680	8.3B	68.4	62.7

9.2 DSO Cross-Model Retrieval

Critical test: Can Model A retrieve experiences from Model B?

Setup:

- Model A: DeepSeek-V3 (7168-dim native, aligned to 7680)
- Model B: Qwen2.5-72B (8192-dim native, aligned to 7680)
- Model C: Zen-Reranker (7680-dim native)
- Library: 10,000 math experiences from training-free GRPO

Query: “How to solve quadratic equations with complex coefficients?”

Table 7: Cross-model retrieval performance

Approach	Recall@5	Recall@10	Latency (ms)
Aligned DeepSeek-V3	84.2%	89.7%	29.8
Aligned Qwen2.5-72B	87.3%	92.1%	31.2
Aligned BGE-Large	79.5%	85.8%	28.4
Zen-Reranker (native)	94.7%	97.9%	21.5

Analysis:

- +7.4% Recall@5 improvement over best aligned model
- 31% latency reduction (critical for real-time inference)
- Consistent performance across different query models

9.3 Compression & Reconstruction

Table 8: BitDelta compression metrics on 10,000 experiences

Metric	Aligned Qwen3	Zen-Reranker	Delta
Original Size (MB)	292.97	292.97	-
Compressed Size (MB)	9.80	9.19	-6.2%
Compression Ratio	29.9×	31.87×	+6.6%
Avg RMSE	0.68%	0.51%	-25.0%
Max RMSE	1.23%	0.87%	-29.3%

Key Insights:

- Native 7680-dim compresses 6.6% better than aligned
- 25% lower reconstruction error on average
- Consistent across diverse experience domains (math, coding, reasoning)

9.4 Latency Breakdown

Table 9: End-to-end DSO pipeline latency (A100 GPU)

Operation	Aligned (ms)	Zen-Reranker (ms)
Tokenization	2.1	2.1
Model Forward Pass	18.3	18.3
Alignment Projection	9.7	0.0
L2 Normalization	0.4	0.4
BitDelta Compression	3.5	3.2
Total	34.0	24.0
Speedup	-	29.4%

10 Security Analysis

10.1 Model Security

Weight Integrity:

- SHA-256 hash published on-chain: `abc123...def789`
- Deterministic inference (fixed random seeds)
- Cryptographic verification during download

Adversarial Robustness:

- FGSM attack: 91.3% accuracy (vs 94.7% clean)
- PGD attack: 88.7% accuracy
- Certified robustness via randomized smoothing [26]

10.2 Network Security

DDoS Protection:

- Rate limiting: 100 submissions/hour per node
- Proof-of-work for spam prevention (Hashcash-style)
- Reputation-based priority queuing

Sybil Resistance:

- Proof-of-stake bonding (minimum 10,000 ZOO per node)
- Slashing for Byzantine behavior (50% bond forfeiture)

- Identity verification via Zero-Knowledge proofs [27]

Data Privacy:

- Experiences are semantic summaries (not raw data)
- Optional homomorphic encryption for sensitive domains [28]
- Zero-knowledge proofs for experience provenance

10.3 Privacy Analysis

Gradient Inversion Attack: Can adversary reconstruct training data from embeddings?
Analysis:

- Embeddings are 7680-dim averages of 8192 tokens
- Information-theoretic bound: at most $\log_2(2^{7680})/8192 \approx 0.94$ bits/token
- Empirical attack success: <5% token recovery (random guess: 0.01%)

Membership Inference Attack: Can adversary determine if text was in training set?
Defense:

- Differential privacy during training ($\epsilon = 8, \delta = 10^{-5}$) [29]
- Membership advantage: <2% above random

11 Comparison with Alternatives

11.1 BAAI bge-reranker-v2-m3

Table 10: Comparison with BAAI bge-reranker-v2-m3

Metric	bge-reranker-v2-m3	Zen-Reranker-8B
Dimension	1024	7680
Parameters	568M	8.3B
MTEB Avg	64.7	68.4
MTEB Retrieval	58.1	62.7
DSO Recall@5	81.3%	94.7%
Compression Ratio	26.9×	31.87×
Latency (ms)	16.2	21.5
Training Cost	\$8,000	\$10,800

Trade-offs:

- BGE: Faster inference, lower cost
- **Zen-Reranker: Higher accuracy, better DSO performance, optimal compression**

11.2 jina-embeddings-v3

Key Difference: Jina supports variable dimensions via truncation/padding, but not optimized for 7680-dim DSO. Zen-Reranker is purpose-built.

Table 11: Comparison with jina-embeddings-v3

Metric	jina-embeddings-v3	Zen-Reranker-8B
Dimension (configurable)	768-8192	7680 (native)
Parameters	570M	8.3B
MTEB Avg	66.2	68.4
DSO Cross-Model	85.7%	94.7%
Byzantine Robustness	Not evaluated	92.1% @ 30% attack
License	Apache 2.0	Apache 2.0

12 Implementation

12.1 Rust Client (Hanzo Network)

Located at: `/hanzo/node/crates/hanzo-zen-reranker/`

Features:

- Candle-based inference (native Rust ML)
- Zero-copy tensor operations
- SIMD-optimized BitDelta compression
- Async multi-threaded API

Example Usage:

```
use hanzo_zen_reranker::{ZenReranker, BitDelta};

let zen = ZenReranker::load("zoo/zen-reranker-8b")?;

// Encode experience
let experience = "When solving geometry, validate bounds";
let embedding = zen.encode(experience)?; // Vec<f32>, 7680

// Compress
let compressed = BitDelta::compress(&embedding)?; // 964 bytes
assert_eq!(compressed.len(), 964);

// Submit to network
hanzo_network::submit_experience(compressed)?;
```

12.2 Python Client (Zoo Network)

Located at: `/zoo/gym/src/gym/train/grpo/zen_integration.py`

Features:

- HuggingFace Transformers integration
- PyTorch-based inference
- FAISS vector index for fast retrieval
- Training-free GRPO compatibility

Example Usage:

```

from gym.train.grpo.zen_integration import ZenRerankerDSO

zen = ZenRerankerDSO("zoo/zen-reranker-8b")

# Encode experience
experience = "When solving geometry, validate bounds"
embedding = zen.encode_experience(experience) # [7680]

# Retrieve similar
query = "How to avoid extraneous solutions in geometry?"
similar = zen.retrieve_similar(query, library, k=5)

# Inject into context
context = "\n".join([exp.text for exp in similar])
prompt = f"{context}\n\n{query}"

```

12.3 Docker Deployment

CPU Inference:

```

docker pull zoo/zen-reranker:cpu-latest
docker run -p 8080:8080 zoo/zen-reranker:cpu-latest

```

GPU Inference (CUDA 12.1):

```

docker pull zoo/zen-reranker:gpu-latest
docker run --gpus all -p 8080:8080 zoo/zen-reranker:gpu-latest

```

Quantized (4-bit GPTQ):

```

docker pull zoo/zen-reranker:quantized-latest
docker run -p 8080:8080 zoo/zen-reranker:quantized-latest

```

API Endpoint:

```

POST /embed
{
  "texts": ["Experience 1", "Experience 2", ...],
  "compress": true
}

Response:
{
  "embeddings": ["<964 bytes>", "<964 bytes>", ...],
  "latency_ms": 21.5
}

```

13 Conclusion

We presented **Zen-Reranker-8B**, the first embedding model with native 7680-dimensional output optimized for Decentralized Semantic Optimization. By eliminating alignment overhead, Zen-Reranker achieves:

1. **98% semantic preservation** (vs 92% with alignment)
2. **31% latency reduction** (21.5ms vs 31.2ms per embedding)

3. **Optimal compression** ($31.87\times$ vs $29.9\times$ BitDelta)
4. **Byzantine robustness** (92% accuracy under 30% attack)
5. **State-of-the-art MTEB** (68.4 average, 62.7 retrieval)

Our three-stage training protocol (projection expansion, reranking fine-tuning, DSO optimization) achieves production-ready performance at \$10,800 cost. Integration with Zoo Network’s training-free GRPO enables collaborative LLM improvement without gradient aggregation, addressing privacy, centralization, and computational efficiency challenges.

13.1 Future Work

- **Dynamic dimensionality:** Adaptive 1920/3840/7680-dim based on query complexity
- **Hierarchical compression:** Multi-scale encoding for network efficiency
- **Multi-granularity retrieval:** Coarse-to-fine experience matching
- **Federated continual learning:** Update Zen-Reranker from DSO feedback without centralized retraining
- **Zero-knowledge proofs:** Private experience verification via zk-SNARKs
- **Cross-chain deployment:** Integrate with Ethereum, Solana, Cosmos ecosystems

13.2 Open Source Release

All code and models are released under Apache 2.0 license:

- **Model:** <https://huggingface.co/zoo/zen-reranker-8b>
- **Code:** <https://github.com/zoolabs/zen-reranker>
- **Zoo Network:** <https://github.com/zoolabs/gym>
- **Hanzo Network:** <https://github.com/hanzoai/node>

We invite the research community to:

- Benchmark on new DSO scenarios
- Propose improvements to BitDelta compression
- Explore alternative Byzantine-robust aggregation methods
- Deploy on production decentralized AI systems

Acknowledgments

This work was supported by Zoo Labs Foundation (501(c)(3) non-profit) and Hanzo AI. We thank the Qwen team for open-sourcing Qwen3-Embedding-8B, Tencent youtu-agent team for training-free GRPO, and the broader open-source AI community for tools and datasets enabling this research.

References

- [1] Brown, T., et al. (2020). Language models are few-shot learners. *NeurIPS*.
- [2] Touvron, H., et al. (2023). LLaMA: Open and efficient foundation language models. *arXiv:2302.13971*.
- [3] Qwen Team. (2024). Qwen3 technical report. *arXiv:2505.09388*.
- [4] Zoo Labs Foundation. (2025). HLLM: Hamiltonian Large Language Models for decentralized semantic optimization. *Technical Report*.
- [5] Tencent youtu-agent Team. (2025). Training-Free Group Relative Policy Optimization. *arXiv:2510.08191v1*.
- [6] Mikolov, T., et al. (2013). Efficient estimation of word representations in vector space. *ICLR*.
- [7] Muennighoff, N., et al. (2023). MTEB: Massive text embedding benchmark. *arXiv:2210.07316*.
- [8] Xiao, S., et al. (2024). C-Pack: Packaged resources to advance general Chinese embedding. *arXiv:2309.07597*.
- [9] Wang, L., et al. (2024). Text embeddings by weakly-supervised contrastive pre-training. *arXiv:2212.03533*.
- [10] Qwen Team. (2024). Qwen3-Embedding technical report. *arXiv:2506.05176*.
- [11] McMahan, B., et al. (2017). Communication-efficient learning of deep networks from decentralized data. *AISTATS*.
- [12] Gupta, O., & Raskar, R. (2018). Distributed learning of deep neural network over multiple agents. *J. Network and Computer Applications*.
- [13] Warnat-Herresthal, S., et al. (2021). Swarm Learning for decentralized and confidential clinical machine learning. *Nature*.
- [14] Jégou, H., et al. (2011). Product quantization for nearest neighbor search. *IEEE TPAMI*.
- [15] Shen, F., et al. (2015). Hashing on nonlinear manifolds. *IEEE TIP*.
- [16] Hanzo AI. (2025). BitDelta: Delta encoding for high-dimensional embeddings. *Technical Report*.
- [17] Blanchard, P., et al. (2017). Machine learning with adversaries: Byzantine tolerant gradient descent. *NeurIPS*.
- [18] Yin, D., et al. (2018). Byzantine-robust distributed learning: Towards optimal statistical rates. *ICML*.
- [19] Yin, D., et al. (2018). Defending against saddle point attack in Byzantine-robust distributed learning. *ICML*.
- [20] Nguyen, T., et al. (2016). MS MARCO: A human generated machine reading comprehension dataset. *arXiv:1611.09268*.
- [21] Bowman, S. R., et al. (2015). A large annotated corpus for learning natural language inference. *EMNLP*.

- [22] Voorhees, E., et al. (2020). TREC-COVID: Constructing a pandemic information retrieval test collection. *arXiv:2005.04474*.
- [23] Thakur, N., et al. (2021). BEIR: A heterogeneous benchmark for zero-shot evaluation of information retrieval models. *arXiv:2104.08663*.
- [24] Karpukhin, V., et al. (2020). Dense passage retrieval for open-domain question answering. *EMNLP*.
- [25] Buterin, V., & Griffith, V. (2017). Casper the Friendly Finality Gadget. *arXiv:1710.09437*.
- [26] Cohen, J., et al. (2019). Certified adversarial robustness via randomized smoothing. *ICML*.
- [27] Groth, J. (2016). On the size of pairing-based non-interactive arguments. *EUROCRYPT*.
- [28] Gentry, C. (2009). Fully homomorphic encryption using ideal lattices. *STOC*.
- [29] Abadi, M., et al. (2016). Deep learning with differential privacy. *CCS*.