



Threshold Fully  
Homomorphic Encryption  
on Zoo Network:  
GPU-Accelerated  
Confidential Computation  
Privacy-Preserving Smart  
Contracts via TFHE and CKKS

Zach Kelling

**Version 2026.04**

*Zoo Labs Foundation*

2026

## Abstract

We present ZOO Network’s integration of threshold fully homomorphic encryption (TFHE and CKKS) into a Layer 2 blockchain, enabling arbitrary computation on encrypted on-chain state. ZOO deploys programmable bootstrapping (TFHE) for boolean and integer circuits and approximate arithmetic (CKKS) for privacy-preserving machine-learning inference directly within smart contracts. By mapping the Number Theoretic Transform butterfly network onto GPU thread blocks, we achieve a **40× speedup** over CPU-only bootstrapping, reducing single-gate bootstrap latency from 12 ms to 0.3 ms. Ciphertext operations are distributed across validators via an Encrypt-to-Shares (E2S) threshold decryption protocol with a **67-of-100** default quorum, ensuring that no individual validator ever possesses a complete decryption key. Key shares are protected in transit with ML-KEM-768 post-quantum key encapsulation and anchored on LUX K-CHAIN. Formal correctness of the CKKS rescaling pipeline, TFHE noise budget tracking, and GPU scaling bounds are verified in Lean 4 (proofs reside in `Crypto/FHE/CKKS.lean`, `Crypto/FHE/TFHE.lean`, and `GPU/FHEScaling.lean`). We benchmark confidential DeFi order matching, sealed-bid auctions, and encrypted neural-network inference, demonstrating practical throughput for production deployment on ZOO T-CHAIN.

## Contents

|          |                                          |          |
|----------|------------------------------------------|----------|
| <b>1</b> | <b>Introduction</b>                      | <b>4</b> |
| <b>2</b> | <b>TFHE on Zoo</b>                       | <b>4</b> |
| 2.1      | LWE Encryption . . . . .                 | 4        |
| 2.2      | Boolean Gate Evaluation . . . . .        | 5        |
| 2.3      | Programmable Bootstrapping . . . . .     | 5        |
| 2.4      | Noise Management . . . . .               | 5        |
| <b>3</b> | <b>CKKS on Zoo</b>                       | <b>5</b> |
| 3.1      | Approximate Arithmetic . . . . .         | 5        |
| 3.2      | Homomorphic Operations . . . . .         | 5        |
| 3.3      | Rescaling and Level Management . . . . . | 6        |
| 3.4      | ML Inference on Encrypted Data . . . . . | 6        |
| <b>4</b> | <b>GPU Acceleration</b>                  | <b>6</b> |
| 4.1      | NTT Butterfly Parallelism . . . . .      | 6        |
| 4.2      | TFHE Bootstrap on GPU . . . . .          | 6        |
| 4.3      | Batch Amortization . . . . .             | 7        |
| 4.4      | Memory Bounds . . . . .                  | 7        |
| <b>5</b> | <b>Threshold Decryption</b>              | <b>7</b> |
| 5.1      | Motivation . . . . .                     | 7        |
| 5.2      | Encrypt-to-Shares Protocol . . . . .     | 7        |
| 5.3      | 67-of-100 Default Threshold . . . . .    | 8        |
| 5.4      | Key Share Rotation . . . . .             | 8        |
| <b>6</b> | <b>Integration with T-Chain</b>          | <b>8</b> |
| 6.1      | FHE Opcodes . . . . .                    | 8        |
| 6.2      | GPU Acceleration Configuration . . . . . | 8        |
| 6.3      | Ciphertext Storage . . . . .             | 9        |

|           |                                             |           |
|-----------|---------------------------------------------|-----------|
| <b>7</b>  | <b>Formal Verification</b>                  | <b>9</b>  |
| 7.1       | CKKS Correctness . . . . .                  | 9         |
| 7.2       | TFHE Noise Tracking . . . . .               | 9         |
| 7.3       | GPU Scaling . . . . .                       | 10        |
| <b>8</b>  | <b>Applications</b>                         | <b>10</b> |
| 8.1       | Confidential DeFi . . . . .                 | 10        |
| 8.2       | Private AI Inference . . . . .              | 10        |
| 8.3       | Sealed-Bid Auctions . . . . .               | 10        |
| <b>9</b>  | <b>Benchmarks</b>                           | <b>10</b> |
| 9.1       | Bootstrap Latency . . . . .                 | 10        |
| 9.2       | CKKS Multiplication Throughput . . . . .    | 11        |
| 9.3       | Threshold Decryption Latency . . . . .      | 11        |
| 9.4       | End-to-End Application Benchmarks . . . . . | 11        |
| <b>10</b> | <b>Related Work</b>                         | <b>11</b> |
| <b>11</b> | <b>Conclusion</b>                           | <b>12</b> |

# 1 Introduction

Public blockchains expose every transaction and every byte of contract state to all participants. This transparency, while valuable for auditability, precludes entire classes of applications: sealed-bid auctions, private order matching, confidential voting, and machine-learning inference over sensitive data.

Two cryptographic families address this gap:

1. **Zero-knowledge proofs** (ZKPs) allow a prover to convince a verifier of a statement’s truth without revealing witness data. ZKPs prove facts about plaintext but do not enable *computation* on encrypted data.
2. **Fully homomorphic encryption** (FHE) permits arbitrary computation directly on ciphertexts. The result, once decrypted, equals the output of the same computation on plaintexts [1].

FHE is strictly more expressive than ZKPs for on-chain privacy: a smart contract can evaluate functions over encrypted inputs without any party—including validators—ever observing plaintext values. ZOO Network combines two complementary FHE schemes:

- **TFHE** (Torus FHE): bit-level and small-integer operations with fast programmable bootstrapping, ideal for branching logic, comparisons, and access-control circuits.
- **CKKS** (Cheon–Kim–Kim–Song): approximate fixed-point arithmetic over vectors of real numbers, ideal for neural-network inference, statistical aggregation, and financial modelling.

ZOO’s approach differs from prior work in three ways: (1) GPU-accelerated bootstrapping that reduces latency by  $40\times$ , making FHE practical inside block times; (2) threshold decryption integrated with ZOO T-CHAIN validators so that decryption requires 67-of-100 consensus; and (3) formal verification in Lean 4 of noise budgets, rescaling correctness, and GPU parallelism bounds.

This paper is organized as follows. Sections 2–3 describe the two FHE schemes and their on-chain encoding. Section 4 details GPU acceleration. Section 5 presents the threshold decryption protocol. Section 6 explains integration with T-CHAIN. Section 7 covers formal verification. Section 8 surveys applications, Section 9 provides benchmarks, and Section 10 discusses related work.

## 2 TFHE on Zoo

### 2.1 LWE Encryption

TFHE operates over the Learning With Errors (LWE) problem. A plaintext bit  $m \in \{0, 1\}$  is encrypted as:

$$\text{ct} = (\mathbf{a}, b) \in \mathbb{Z}_q^n \times \mathbb{Z}_q, \quad b = \langle \mathbf{a}, \mathbf{s} \rangle + m \cdot \lfloor q/2 \rfloor + e \quad (1)$$

where  $\mathbf{s} \in \mathbb{Z}_q^n$  is the secret key,  $\mathbf{a} \stackrel{\$}{\leftarrow} \mathbb{Z}_q^n$  is uniform, and  $e$  is a small Gaussian noise term with standard deviation  $\sigma$ . ZOO fixes  $n = 630$ ,  $q = 2^{32}$ , and  $\sigma = 2^{-15}$ , yielding 128-bit security under standard LWE hardness assumptions [4].

## 2.2 Boolean Gate Evaluation

Homomorphic boolean gates (AND, OR, XOR, NOT) on LWE ciphertexts follow the gate-bootstrapping paradigm of Chillotti et al. [2]. For two input ciphertexts  $\text{ct}_0, \text{ct}_1$ :

$$\text{ct}_{\text{AND}} = \text{Bootstrap}((0, \lfloor q/8 \rfloor) - \text{ct}_0 - \text{ct}_1) \quad (2)$$

The bootstrap operation simultaneously evaluates a lookup table and refreshes noise, producing a ciphertext with bounded noise regardless of input noise level.

## 2.3 Programmable Bootstrapping

ZOO uses programmable bootstrapping (PBS) to evaluate arbitrary functions  $f : \mathbb{Z}_p \rightarrow \mathbb{Z}_p$  during the bootstrap step itself. The test polynomial  $v(X) \in \mathbb{Z}_q[X]/(X^N + 1)$  encodes  $f$  in its coefficients:

$$v(X) = \sum_{i=0}^{N-1} f\left(\left\lfloor \frac{i \cdot p}{2N} \right\rfloor\right) X^i \quad (3)$$

This eliminates post-bootstrap computation for common operations (sign extraction, ReLU, step functions), reducing circuit depth and noise accumulation.

## 2.4 Noise Management

Each homomorphic operation adds noise to the ciphertext. ZOO tracks noise budgets per ciphertext in the VM state:

**Definition 1** (Noise Budget). For ciphertext  $\text{ct} = (\mathbf{a}, b)$  with secret  $\mathbf{s}$ , the noise  $e = b - \langle \mathbf{a}, \mathbf{s} \rangle - m \cdot \lfloor q/2 \rfloor$ . The noise budget is  $\beta = \lfloor \log_2(q/2) \rfloor - \lceil \log_2 |e| \rceil$  bits. Decryption succeeds when  $\beta > 0$ .

The VM rejects execution when a ciphertext's projected noise budget would drop below 1 bit, forcing the contract to insert a bootstrap before proceeding. Gas metering accounts for bootstrap cost.

# 3 CKKS on Zoo

## 3.1 Approximate Arithmetic

CKKS [3] encodes a vector  $\mathbf{z} \in \mathbb{C}^{N/2}$  into a polynomial ring  $R_q = \mathbb{Z}_q[X]/(X^N + 1)$  via the canonical embedding inverse, scaled by  $\Delta$ :

$$m(X) = \lfloor \Delta \cdot \sigma^{-1}(\mathbf{z}) \rfloor \in R_q \quad (4)$$

where  $\Delta = 2^{40}$  is the scaling factor and  $\sigma^{-1}$  is the inverse canonical embedding. Encryption produces  $\text{ct} = (\text{ct}_0, \text{ct}_1) \in R_q^2$  under RLWE.

## 3.2 Homomorphic Operations

Addition of two CKKS ciphertexts at the same level is component-wise. Multiplication produces a degree-2 ciphertext that must be relinearized:

$$\text{ct}_{\text{mult}} = (\text{ct}_0^{(a)} \cdot \text{ct}_0^{(b)}, \text{ct}_0^{(a)} \cdot \text{ct}_1^{(b)} + \text{ct}_1^{(a)} \cdot \text{ct}_0^{(b)}, \text{ct}_1^{(a)} \cdot \text{ct}_1^{(b)}) \quad (5)$$

$$\text{ct}_{\text{relin}} = \text{Relin}(\text{ct}_{\text{mult}}, \text{evk}) \quad (6)$$

### 3.3 Rescaling and Level Management

After multiplication, the scaling factor doubles to  $\Delta^2$ . Rescaling divides by  $\Delta$  and drops one prime from the ciphertext modulus chain:

$$\text{Rescale}(\text{ct}, q_\ell) \rightarrow \text{ct}' \text{ at level } \ell - 1, \quad q_{\ell-1} = q_\ell / p_\ell \quad (7)$$

ZOO provisions a modulus chain of  $L = 15$  levels by default, supporting up to 14 sequential multiplications—sufficient for a 4-layer neural network with ReLU approximations. The VM tracks levels alongside gas, rejecting operations that would exhaust the chain.

### 3.4 ML Inference on Encrypted Data

A neural-network forward pass on encrypted inputs uses CKKS slot packing. For a weight matrix  $W \in \mathbb{R}^{m \times n}$  and encrypted input  $\text{Enc}(\mathbf{x})$ :

$$\text{Enc}(W\mathbf{x} + \mathbf{b}) = W \otimes \text{Enc}(\mathbf{x}) \oplus \text{Enc}(\mathbf{b}) \quad (8)$$

where  $\otimes$  denotes plaintext–ciphertext multiplication and  $\oplus$  denotes ciphertext addition. Activation functions are approximated by low-degree polynomials (degree-3 Chebyshev for ReLU) consuming one multiplicative level each.

## 4 GPU Acceleration

### 4.1 NTT Butterfly Parallelism

The computational bottleneck in both TFHE and CKKS is polynomial multiplication in  $R_q$ , implemented via the Number Theoretic Transform (NTT). The NTT butterfly network of depth  $\log_2 N$  maps naturally to GPU thread blocks:

- Each butterfly stage is a parallel map over  $N/2$  independent pairs.
- Stages synchronize via `__syncthreads()` within a block or global barrier across blocks.
- For  $N = 2^{14} = 16384$  (Zoo’s default polynomial degree), a single SM processes 8192 butterflies per stage.

### 4.2 TFHE Bootstrap on GPU

A single TFHE bootstrap consists of:

1. **Blind rotation:**  $n = 630$  external products, each requiring two NTT forward/inverse transforms on degree- $N$  polynomials.
2. **Key switching:** one matrix–vector product in  $\mathbb{Z}_q^{n' \times n}$ .
3. **Modulus switching:** rounding from  $q$  to  $q'$ .

On CPU (single-threaded, AVX-512), one bootstrap takes  $\sim 12$  ms. On an NVIDIA H100 GPU:

- Blind rotation: 630 external products execute as a batched NTT kernel, 0.22 ms.
- Key switching: cuBLAS GEMV, 0.05 ms.

- Modulus switching: element-wise kernel, 0.03 ms.
- **Total: 0.3 ms** ( $40\times$  speedup).

### 4.3 Batch Amortization

When multiple independent bootstraps are required (e.g., evaluating a 256-bit comparison circuit), Zoo batches them into a single GPU dispatch:

$$T_{\text{batch}}(k) = T_{\text{single}} + \alpha \cdot k, \quad \alpha \approx 0.015 \text{ ms per additional bootstrap} \quad (9)$$

for  $k \leq 4096$  bootstraps. Beyond this, L2 cache pressure causes throughput degradation. At  $k = 256$  (a typical 8-bit integer comparison), total latency is  $\sim 4.1$  ms.

### 4.4 Memory Bounds

FHE ciphertexts are large. A single CKKS ciphertext at level  $\ell = 15$  with  $N = 16384$  occupies:

$$|\text{ct}| = 2 \cdot N \cdot \ell \cdot 8 = 2 \cdot 16384 \cdot 15 \cdot 8 \approx 3.75 \text{ MB} \quad (10)$$

An H100 with 80 GB HBM3 can hold  $\sim 21,000$  such ciphertexts simultaneously. Zoo validators must provision at least 40 GB GPU memory for FHE workloads. The VM enforces per-contract ciphertext storage limits via gas metering.

## 5 Threshold Decryption

### 5.1 Motivation

In a blockchain setting, no single entity should hold the decryption key. Zoo distributes the TFHE/CKKS secret key across validators using Shamir secret sharing with a 67-of-100 reconstruction threshold.

### 5.2 Encrypt-to-Shares Protocol

Key generation proceeds as follows:

1. A trusted dealer (the genesis ceremony) generates  $\mathbf{s} \in \mathbb{Z}_q^n$  and computes Shamir shares  $\mathbf{s}_i$  for  $i \in \{1, \dots, 100\}$  with threshold  $t = 67$ .
2. Each share  $\mathbf{s}_i$  is encapsulated under validator  $i$ 's ML-KEM-768 public key and published to LUX K-CHAIN as a sealed transaction.
3. Validators decrypt their share locally and store it in a hardware security module (HSM) or TEE enclave.

**Definition 2** (E2S Decryption). To decrypt ciphertext  $\text{ct} = (\mathbf{a}, b)$ , each participating validator  $i$  computes a partial decryption:

$$d_i = \langle \mathbf{a}, \mathbf{s}_i \rangle + e_i \quad (11)$$

where  $e_i$  is a smudging noise term with  $\|e_i\| \ll q/4$  to prevent share leakage. The aggregator reconstructs:

$$m = b - \sum_{i \in S} \lambda_i \cdot d_i \pmod{q}, \quad |S| \geq 67 \quad (12)$$

where  $\lambda_i$  are Lagrange interpolation coefficients for the set  $S$ .

### 5.3 67-of-100 Default Threshold

The threshold  $t = 67$  is chosen to satisfy:

- **Safety:** an adversary controlling fewer than 34 validators cannot decrypt any ciphertext.
- **Liveness:** the system tolerates up to 33 offline or Byzantine validators while maintaining decryption capability.
- **Alignment with BFT:** the  $2/3 + 1$  threshold matches the Byzantine fault tolerance bound used in ZOO T-CHAIN consensus.

### 5.4 Key Share Rotation

Validator set changes trigger proactive secret sharing (PSS). New shares are computed via a re-sharing protocol without reconstructing  $\mathbf{s}$ :

$$\mathbf{s}_i^{(\text{new})} = \sum_{j \in S_{\text{old}}} \lambda_j \cdot \mathbf{s}_{j \rightarrow i} \quad (13)$$

where  $\mathbf{s}_{j \rightarrow i}$  is a re-share from old validator  $j$  to new validator  $i$ , again encapsulated under ML-KEM-768. This ensures forward secrecy: compromised old shares cannot decrypt ciphertexts encrypted under the new epoch’s public key.

## 6 Integration with T-Chain

ZOO T-CHAIN (Threshold Chain) combines threshold signing [13] with FHE in a unified VM execution environment.

### 6.1 FHE Opcodes

The T-CHAIN VM extends the EVM with FHE-specific opcodes:

| Opcode      | Operation                                 | Gas Cost  |
|-------------|-------------------------------------------|-----------|
| FHE.ENC     | Encrypt plaintext to TFHE/CKKS ciphertext | 50,000    |
| FHE.ADD     | Homomorphic addition                      | 10,000    |
| FHE.MUL     | Homomorphic multiplication (+ relin)      | 200,000   |
| FHE.BOOT    | Programmable bootstrap (TFHE)             | 500,000   |
| FHE.RESCALE | Rescale CKKS ciphertext                   | 80,000    |
| FHE.ROTATE  | Rotate CKKS slots                         | 150,000   |
| FHE.DEC.REQ | Request threshold decryption              | 1,000,000 |

Table 1: FHE opcodes added to the T-CHAIN VM

Gas costs reflect computational load. FHE.DEC.REQ is the most expensive because it triggers a multi-round threshold decryption protocol across 67+ validators.

### 6.2 GPU Acceleration Configuration

Validators configure GPU FHE acceleration in their node configuration:

- `fhe.gpu.enabled`: boolean, default `true` for T-CHAIN validators.

- `fhe.gpu.device`: CUDA device index (supports multi-GPU).
- `fhe.gpu.max_batch`: maximum bootstrap batch size (default 4096).
- `fhe.gpu.memory_limit`: GPU memory reserved for FHE (default 40 GB).

Validators without GPU support may still participate in consensus and threshold decryption but cannot execute FHE opcodes. Block producers for FHE-heavy blocks are selected preferentially from the GPU-enabled validator subset.

### 6.3 Ciphertext Storage

TFHE ciphertexts (LWE,  $\sim 5$  KB each) and CKKS ciphertexts ( $\sim 3.75$  MB at full level) are stored in a dedicated ciphertext trie separate from the account state trie. Content-addressing (Blake3 hash) enables deduplication. State rent is charged per ciphertext per epoch to bound storage growth.

## 7 Formal Verification

ZOO maintains Lean 4 proofs for critical FHE subsystems, following the broader ZOO formal verification initiative [12].

### 7.1 CKKS Correctness

The file `Crypto/FHE/CKKS.lean` proves:

**Theorem 1** (Rescaling Precision). *For a CKKS ciphertext at level  $\ell$  with scaling factor  $\Delta$  and modulus  $q_\ell$ , rescaling to level  $\ell - 1$  introduces approximation error bounded by:*

$$\|\epsilon_{\text{rescale}}\|_\infty \leq \frac{N \cdot B_{\text{key}}}{p_\ell} \quad (14)$$

where  $B_{\text{key}}$  bounds the secret key coefficients and  $p_\ell$  is the dropped prime.

### 7.2 TFHE Noise Tracking

The file `Crypto/FHE/TFHE.lean` verifies:

**Theorem 2** (Bootstrap Noise Bound). *After programmable bootstrapping with parameters  $(n, N, q, \sigma, B_g, d_g)$ , the output noise satisfies:*

$$\|e_{\text{out}}\| \leq n \cdot d_g \cdot (N + 1) \cdot B_g \cdot \sigma_{\text{bsk}} + (1 + N \cdot \|\mathbf{s}\|_1) \cdot \frac{B_g^{d_g}}{2} \quad (15)$$

*This bound is strictly less than  $q/4$  for ZOO's parameter choices, guaranteeing correct decryption.*

### 7.3 GPU Scaling

The file `GPU/FHEScaling.lean` proves:

**Proposition 3** (NTT Parallelism). *The  $\log_2 N$ -stage butterfly network for degree- $N$  NTT admits a work-depth decomposition with:*

$$W = \frac{N}{2} \log_2 N, \quad D = \log_2 N \quad (16)$$

*Given  $P \geq N/2$  parallel processors, the time complexity is  $\Theta(\log_2 N)$ , achieving linear speedup up to  $P = N/2$ .*

## 8 Applications

### 8.1 Confidential DeFi

**Encrypted order matching.** Users submit encrypted limit orders  $\text{Enc}(\text{price}, \text{quantity})$ . The matching engine evaluates comparison circuits homomorphically (TFHE for price comparison, CKKS for quantity arithmetic). Matched orders are decrypted via threshold decryption; unmatched orders remain encrypted and invisible to other participants.

This eliminates front-running, sandwich attacks, and information leakage from the order book.

### 8.2 Private AI Inference

A model owner publishes encrypted weights  $\text{Enc}(W_i)$  and a user submits encrypted input  $\text{Enc}(\mathbf{x})$ . The contract evaluates the forward pass entirely under CKKS encryption:

$$\text{Enc}(\hat{y}) = f_{\text{NN}}(\text{Enc}(W_1), \dots, \text{Enc}(W_L), \text{Enc}(\mathbf{x})) \quad (17)$$

Neither the model weights nor the input are revealed to any party. The user decrypts  $\hat{y}$  locally. This enables monetization of proprietary models without exposing intellectual property.

### 8.3 Sealed-Bid Auctions

Bidders submit  $\text{Enc}(\text{bid}_i)$  to an auction contract. The contract computes the maximum via a homomorphic comparison tree (TFHE with  $\lceil \log_2 k \rceil$  rounds for  $k$  bids). Only the winning bid is decrypted; losing bids remain sealed.

## 9 Benchmarks

All benchmarks run on Zoo testnet validators equipped with NVIDIA H100 80 GB GPUs, AMD EPYC 9654 CPUs (96 cores), and 512 GB DDR5.

### 9.1 Bootstrap Latency

Batch amortization yields superlinear speedup because GPU occupancy increases with batch size until the 4096-element sweet spot.

| Operation                  | CPU (ms) | GPU (ms) | Speedup |
|----------------------------|----------|----------|---------|
| Single TFHE bootstrap      | 12.0     | 0.30     | 40×     |
| 8-bit integer comparison   | 96.0     | 4.10     | 23×     |
| 256-bit comparison circuit | 3,072    | 52.0     | 59×     |
| Batch 4096 bootstraps      | 49,152   | 62.0     | 793×    |

Table 2: TFHE bootstrap latency: CPU vs GPU

| Metric                                  | CPU | GPU   |
|-----------------------------------------|-----|-------|
| Single ciphertext multiply + relin (ms) | 8.5 | 0.6   |
| Throughput (multiplications/sec)        | 118 | 1,667 |
| Rescale (ms)                            | 2.1 | 0.15  |
| Slot rotation (ms)                      | 6.2 | 0.45  |

Table 3: CKKS operation throughput at  $N = 16384$ ,  $L = 15$

## 9.2 CKKS Multiplication Throughput

## 9.3 Threshold Decryption Latency

The dominant cost is network latency for collecting 67 partial decryptions. Computation is negligible relative to communication.

## 9.4 End-to-End Application Benchmarks

# 10 Related Work

**Lux TFHE.** The LUX `lux-tfhe` paper [9] introduced GPU-accelerated TFHE for the Lux L0 base layer. ZOO extends this work with threshold decryption, CKKS support, and deeper VM integration via FHE-specific opcodes on T-CHAIN.

**Zama fhEVM.** Zama’s `fhEVM` [6] pioneered FHE smart contracts on EVM-compatible chains. ZOO differs in three respects: (1) CKKS support for approximate arithmetic, enabling ML inference; (2) threshold decryption rather than single-operator key management; (3) formal verification of noise bounds and GPU scaling in Lean 4.

**Sunscreen FHE.** The Sunscreen compiler [8] provides a Rust-based FHE toolkit targeting BFV and CKKS. ZOO builds on similar compiler infrastructure but integrates it directly into the blockchain VM with gas metering and threshold key management.

**Concrete/TFHE-rs.** Zama’s `tfhe-rs` library [7] provides a Rust implementation of TFHE with GPU support. ZOO’s GPU kernels build on this foundation, extending it with batched bootstrapping optimizations and T-CHAIN integration.

**Threshold FHE in MPC.** Boneh et al. [5] formalized threshold FHE for multi-party computation. ZOO adapts this framework to the blockchain setting, replacing interactive MPC rounds with asynchronous partial-decryption collection anchored by BFT consensus.

| Phase                               | Latency (ms) |
|-------------------------------------|--------------|
| Partial decryption (per validator)  | 0.8          |
| Network collection (67 shares, p95) | 180          |
| Lagrange aggregation                | 2.1          |
| Smudging noise verification         | 0.3          |
| <b>Total end-to-end</b>             | <b>~183</b>  |

Table 4: Threshold decryption latency for LWE ciphertext (67-of-100)

| Application                      | Latency (ms) | Gas Cost |
|----------------------------------|--------------|----------|
| Encrypted order match (2 orders) | 48           | 1.8M     |
| Sealed-bid auction (64 bids)     | 312          | 42M      |
| 4-layer NN inference (CKKS)      | 580          | 68M      |

Table 5: End-to-end application benchmarks on GPU-enabled validators

## 11 Conclusion

We have presented ZOO Network’s threshold FHE system, combining TFHE for boolean circuits and CKKS for approximate arithmetic with GPU acceleration and validator-distributed threshold decryption. Key results:

1. **40× GPU speedup** for single TFHE bootstrap (12 ms  $\rightarrow$  0.3 ms), with up to 793× for batched operations.
2. **67-of-100 threshold decryption** aligned with BFT consensus, with end-to-end latency of  $\sim$ 183 ms dominated by network communication.
3. **Formal verification** of CKKS rescaling precision, TFHE noise bounds, and GPU parallelism in Lean 4.
4. **Practical applications** demonstrated: encrypted order matching (48 ms), sealed-bid auctions (312 ms), and private neural-network inference (580 ms).

FHE on ZOO T-CHAIN enables a new class of privacy-preserving smart contracts where computation proceeds on encrypted state without any party—including validators—observing plaintext values. Combined with post-quantum key encapsulation (ML-KEM-768) for share distribution, the system provides long-term confidentiality against both classical and quantum adversaries.

Future work includes bootstrapping-free FHE schemes for reduced latency, hardware-accelerated key switching via custom FPGA pipelines, and integration with ZOO’s experience ledger [10] for privacy-preserving model training.

## Acknowledgments

The authors thank the ZOO Labs Foundation engineering team for GPU kernel development, the LUX cryptography team for the `lux-tfhe` foundation, and the HANZO infrastructure team for validator provisioning.

## References

- [1] Craig Gentry. *Fully Homomorphic Encryption Using Ideal Lattices*. STOC, 2009.
- [2] Ilaria Chillotti, Nicolas Gama, Mariya Georgieva, and Malika Izabachène. *TFHE: Fast Fully Homomorphic Encryption over the Torus*. Journal of Cryptology, 33(1):34–91, 2020.
- [3] Jung Hee Cheon, Andrey Kim, Miran Kim, and Yongsoo Song. *Homomorphic Encryption for Arithmetic of Approximate Numbers*. ASIACRYPT, 2017.
- [4] Martin R. Albrecht, Rachel Player, and Sam Scott. *On the Concrete Hardness of Learning with Errors*. Journal of Mathematical Cryptology, 9(3):169–203, 2015.
- [5] Dan Boneh, Rosario Gennaro, Steven Goldfeder, Aayush Jain, Sam Kim, Peter M.R. Rasmussen, and Amit Sahai. *Threshold Cryptosystems from Threshold Fully Homomorphic Encryption*. CRYPTO, 2018.
- [6] Zama. *fhEVM: Confidential Smart Contracts on the EVM Using Fully Homomorphic Encryption*. Zama Whitepaper, 2023.
- [7] Zama. *TFHE-rs: A Pure Rust Implementation of the TFHE Scheme*. <https://github.com/zama-ai/tfhe-rs>, 2023.
- [8] Sunscreen. *Sunscreen: A Compiler for FHE Programs*. Sunscreen Tech Report, 2023.
- [9] Lux Industries. *GPU-Accelerated TFHE for Lux Network*. Lux Research, 2025.
- [10] Zach Kelling. *Zoo Network Tokenomics: Fair Distribution Through Complete Airdrop*. Zoo Labs Foundation, 2025.
- [11] Zach Kelling. *Zoo EVM: A Post-Quantum Safe Layer 2 Blockchain with AI-Native Precompiles*. Zoo Labs Foundation, 2026.
- [12] Zach Kelling. *Zoo Network: Decentralized AI Training and Inference Layer with Semantic Optimization*. Zoo Labs Foundation, 2025.
- [13] Zach Kelling. *Threshold Signatures on Zoo T-Chain* (forthcoming). Zoo Labs Foundation, 2026.