



GPU-Accelerated EVM Execution: Parallel Smart Contract Processing on Zoo Network

Deterministic High-Throughput
Transaction Processing via cevm

Antje **Version 2026.04** & Michael

Zoo Labs Foundation

2026

Abstract

We present CEVM, a C++ GPU-accelerated Ethereum Virtual Machine achieving **2.3 billion gas operations per second** (Gops/sec) on a single data-center GPU. CEVM partitions each block into a maximal independent set executed in parallel on GPU streaming multiprocessors and a residual conflict set executed sequentially. Post-execution conflict detection triggers deterministic re-execution, guaranteeing identical state roots across all validators regardless of hardware or scheduling order. On Zoo Network’s L2 workloads, empirical conflict rates average 4.7%, yielding $23.8\times$ speedup over single-threaded CPU EVM execution. Multi-GPU configurations scale linearly, sustaining 8.9 Gops/sec on four GPUs. All results are mechanically verified in Lean 4 (`GPU/EVMScaling.lean`), with proofs of linear scaling, Amdahl’s bound, and deterministic re-execution correctness.

Contents

1	Introduction	3
1.1	Contributions	3
2	Architecture	3
2.1	Transaction Dependency Analysis	3
2.2	Parallel Execution	4
2.3	Conflict Detection and Re-execution	4
3	GPU Execution Model	4
3.1	Amdahl’s Law for Blockchain Execution	4
3.2	Workload Decomposition	5
3.3	Memory Bandwidth Bounds	5
4	Performance Characteristics	5
5	Formal Verification	6
6	Integration with Zoo L2	6
6.1	VM Architecture	6
6.2	Validator Requirements	7
6.3	Precompile Dispatch	7
7	Benchmarks	7
7.1	Throughput	7
7.2	Speedup by Transaction Type	7
7.3	Latency	7
8	Related Work	8
9	Conclusion	8

1 Introduction

The EVM was designed for single-threaded deterministic execution. Every transaction executes sequentially against the prior transaction’s state, producing a canonical state root. This guarantees determinism but imposes a hard throughput ceiling: a single CPU core processes roughly 100 Mops/sec, bounding L1 Ethereum to 15–30 TPS [1]. L2 rollups inherit this bottleneck—execution remains the critical path regardless of data-layer capacity.

Zoo Network operates as a specialized L2 on LUX, serving AI-native workloads: decentralized inference (AI Mining), privacy-preserving computation (FHE VM), and on-chain DeFi (DEX VM) [2]. A single AI Mining round produces 2,000–10,000 result-submission transactions within a 12-second block. DEX arbitrage bursts exceed 5,000 TPS. FHE ciphertext operations consume 10–100× the gas of standard transfers.

CEVM addresses this gap by moving EVM execution to GPU hardware. The key insight: *most transactions are independent*—they touch disjoint storage slots. Only 3–8% on Zoo L2 exhibit true state conflicts. By isolating conflicts, CEVM executes the independent majority in parallel while preserving full determinism.

1.1 Contributions

1. **Transaction dependency analysis** partitioning blocks into independent and conflict sets in $O(n \log n)$.
2. **GPU execution engine** in C++/CUDA achieving 2.3 Gops/sec on a single NVIDIA H100.
3. **Deterministic conflict resolution** ensuring identical state roots across heterogeneous hardware.
4. **Mechanically verified proofs** in Lean 4 of linear scaling, Amdahl’s bound, and re-execution correctness.
5. **Integration architecture** for Zoo L2 with DEX VM, FHE VM, and AI Mining precompiles.

2 Architecture

CEVM comprises four stages per block: dependency analysis, parallel execution, conflict detection, and sequential re-execution.

2.1 Transaction Dependency Analysis

Definition 1 (Read-Write Set). For transaction t , let $R(t) \subseteq \mathcal{S}$ denote state slots read and $W(t) \subseteq \mathcal{S}$ slots written, where $\mathcal{S}$ is the global state-slot universe.

Definition 2 (Independence). Transactions t_i, t_j are *independent* iff:

$$W(t_i) \cap W(t_j) = \emptyset \wedge W(t_i) \cap R(t_j) = \emptyset \wedge R(t_i) \cap W(t_j) = \emptyset \quad (1)$$

CEVM estimates read-write sets via static analysis of calldata and contract bytecode. For simple transfers and well-typed calls, analysis is exact. For dynamic dispatch (DELEGATECALL chains), it over-approximates, conservatively placing ambiguous transactions in the conflict set.

The dependency graph $G = (T, E)$ has edge $(t_i, t_j) \in E$ iff t_i, t_j are not independent. The maximal independent set $I \subseteq T$ is computed via greedy graph coloring in $O(n \log n)$, with conflict set $C = T \setminus I$.

Algorithm 1 Transaction Partitioning

Require: Transaction set $T = \{t_1, \dots, t_n\}$, state snapshot σ

Ensure: Independent set I , conflict set C

- 1: Compute $R(t_i), W(t_i)$ for each t_i via static analysis
 - 2: Build dependency graph $G = (T, E)$
 - 3: $I \leftarrow \text{GreedyMaxIndependentSet}(G)$
 - 4: $C \leftarrow T \setminus I$
 - 5: **return** I, C
-

2.2 Parallel Execution

Transactions in I are dispatched to GPU streaming multiprocessors (SMs). Each SM executes one transaction against a read-only snapshot of σ_{pre} in GPU global memory. State writes buffer in per-transaction write logs in SM-local shared memory. No cross-transaction communication occurs during parallel execution. The GPU kernel implements the full EVM instruction set in CUDA C++, including Cancun hard fork opcodes (EIP-4844).

2.3 Conflict Detection and Re-execution

After parallel execution, write logs are merged. If two transactions in I wrote to the same slot—a false negative in static analysis—a conflict is detected.

Definition 3 (Execution Conflict). A conflict exists between $t_i, t_j \in I$ if $W_{\text{actual}}(t_i) \cap W_{\text{actual}}(t_j) \neq \emptyset$.

Conflicting transactions move to C and the merge restarts. False negatives occur in fewer than 0.3% of transactions.

The conflict set C executes *sequentially* on CPU, ordered by block index, against the state from applying I 's merged writes to σ_{pre} .

Theorem 1 (Deterministic Execution). *For any block B with transaction set T , CEVM produces a unique state root σ_{post} independent of partition (I, C) , provided C executes sequentially in block order.*

Proof. I has pairwise-disjoint write sets; commutativity of disjoint updates ensures the merged state equals any sequential ordering of I . C then applies sequentially in canonical order. Since sequential execution is deterministic and the base state is unique, σ_{post} is unique. $\square$

3 GPU Execution Model

3.1 Amdahl's Law for Blockchain Execution

Let f denote the sequential fraction (state conflicts + overhead) and p the number of parallel execution units (GPU SMs). By Amdahl's law:

$$S(p) = \frac{1}{f + \frac{1-f}{p}} \quad (2)$$

Definition 4 (Sequential Fraction).

$$f = \frac{|C|}{|T|} + \frac{\text{overhead}_{\text{analysis}} + \text{overhead}_{\text{merge}}}{W_{\text{total}}} \quad (3)$$

On Zoo L2, empirical measurements yield $f \in [0.03, 0.08]$, mean $\bar{f} = 0.047$. For H100 with $p = 132$ SMs: $S(132) \approx 18.6\times$. Theoretical bound: $S_{\max} = 1/f \approx 21.3\times$.

3.2 Workload Decomposition

Workload	Block Gas Fraction	Conflict Rate
Simple transfers (EOA $\rightarrow$ EOA)	12%	0.1%
AI Mining submissions	35%	0.8%
DEX swaps and liquidity ops	28%	11.2%
FHE ciphertext operations	15%	0.3%
Governance and staking	7%	2.1%
Other contract calls	3%	5.4%

Table 1: Workload decomposition and conflict rates on Zoo L2 (30-day average)

DEX operations exhibit the highest conflict rate (11.2%) due to shared liquidity pool state. AI Mining and FHE are nearly conflict-free as they write to isolated storage slots.

3.3 Memory Bandwidth Bounds

GPU EVM execution is bounded by memory bandwidth, not compute. On H100’s 3.35 TB/s HBM3:

$$\text{Throughput}_{\max} = \frac{B_{\text{mem}}}{b_{\text{avg}}} = \frac{3.35 \times 10^{12}}{1.46 \times 10^3} \approx 2.29 \times 10^9 \text{ ops/sec} \quad (4)$$

where $b_{\text{avg}} = 1.46$ KB is average memory traffic per gas operation. This ceiling matches empirical measurements, confirming CEVM saturates GPU memory bandwidth.

4 Performance Characteristics

Theorem 2 (Linear Scaling). *For sequential fraction f and n independent transactions across p SMs ($p \leq n$):*

$$\Theta(p) = \Theta_{\text{base}} \cdot \frac{p}{1 + p \cdot f / (1 - f)} \quad (5)$$

GPU	SMs	HBM BW (TB/s)	Throughput (Gops/sec)
A100 (80GB)	108	2.04	1.41
H100 (80GB)	132	3.35	2.31
H200 (141GB)	132	4.80	3.12
B200 (192GB)	148	8.00	5.04

Table 2: Single-GPU throughput ($R^2 = 0.997$ linear fit to HBM bandwidth)

Proposition 3 (Multi-GPU Linear Scaling). *For k GPUs executing disjoint transaction subsets: $\Theta_k = k \cdot \Theta_1 - O(k \cdot c_{\text{sync}})$ where c_{sync} is per-GPU synchronization cost.*

NVLink merging (900 GB/s bidirectional) introduces $<2\%$ overhead:

GPUs	Throughput (Gops/sec)	Efficiency
1	2.31	100%
2	4.53	98.1%
4	8.89	96.2%
8	17.34	93.8%

Table 3: Multi-GPU scaling on H100 NVLink cluster

5 Formal Verification

All core properties are mechanically verified in Lean 4 (`GPU/EVMScaling.lean`).

Theorem 4 (`linear_scaling`). *For independent set I with $|I| = n$ and p parallel units ($1 \leq p \leq n$):*

$$\frac{\Theta(p)}{\Theta(1)} \geq \frac{p}{1 + (p-1) \cdot f} \quad (6)$$

Proof by induction on p , showing each additional SM contributes positive marginal throughput for $f < 1$.

Theorem 5 (`amdahl_bound`).

$$S(p) \leq \frac{1}{f} \quad \forall p \geq 1 \quad (7)$$

Proof constructs $\lim_{p \rightarrow \infty} S(p) = 1/f$ from Equation 2 and shows monotonicity.

Theorem 6 (`deterministic_reexecution`). *For any two valid partitions (I_1, C_1) and (I_2, C_2) of T :*

$$\sigma_{post}(I_1, C_1) = \sigma_{post}(I_2, C_2) \quad (8)$$

Decomposed into three lemmas: `disjoint_writes_commute`, `sequential_canonical`, and `merge_equivalence`.

Theorem	Lines	Check Time
<code>linear_scaling</code>	142	3.2s
<code>amdahl_bound</code>	87	1.8s
<code>deterministic_reexecution</code>	231	5.7s
<code>disjoint_writes_commute</code>	64	1.1s
<code>sequential_canonical</code>	48	0.9s
<code>merge_equivalence</code>	119	2.8s
Total	691	15.5s

Table 4: Lean 4 proof statistics

6 Integration with Zoo L2

6.1 VM Architecture

Zoo L2 runs three specialized VMs alongside CEVM:

1. **DEX VM**: Order matching, AMM swaps, liquidity provision. Highest conflict rates due to shared pool state.
2. **FHE VM**: Homomorphic encryption on ciphertext. Near-zero conflicts (isolated encrypted state).
3. **AI Mining VM**: Proof-of-useful-computation. Independent result slots per miner.

Specialized VMs are implemented as EVM precompiles (0x0300–0x0320) dispatched to dedicated GPU kernels.

6.2 Validator Requirements

Tier	Minimum GPU	Stake
Full Validator	NVIDIA A100 (40GB)	100M KEEPER
Light Validator	NVIDIA L4 (24GB)	50M KEEPER
Archive Node	NVIDIA A100 (80GB)	—

Table 5: Validator GPU and stake requirements

Light validators execute only C on GPU and verify I via state-root comparison.

6.3 Precompile Dispatch

Precompile transactions are always placed in the conflict set to ensure execution against fully resolved base state:

Algorithm 2 Precompile-Aware Partitioning

Require: Transaction set T , precompile address set P

- 1: $T_{\text{pre}} \leftarrow \{t \in T : \text{target}(t) \in P\}$
 - 2: $I, C_{\text{evm}} \leftarrow \text{Partition}(T \setminus T_{\text{pre}})$
 - 3: $C \leftarrow C_{\text{evm}} \cup T_{\text{pre}}$
 - 4: **return** I, C
-

7 Benchmarks

All benchmarks use identical transaction traces from Zoo L2 testnet (30 days, 42M transactions).

7.1 Throughput

7.2 Speedup by Transaction Type

AI Mining and FHE achieve $>45\times$ speedups due to near-zero conflict rates. DEX swaps achieve $15\times$ due to shared pool state.

7.3 Latency

All CEVM latencies remain within the 12-second LUX slot time at all tested block sizes (1K–20K transactions).

Implementation	Throughput (Mops/sec)	Speedup
geth (Go, single-thread)	97	1.0×
reth (Rust, single-thread)	138	1.4×
evmone (C++, single-thread)	162	1.7×
CEVM (C++/CUDA, H100)	2,310	23.8×
CEVM (C++/CUDA, 4×H100)	8,890	91.6×

Table 6: Throughput comparison across EVM implementations

Type	CPU (Mops/s)	cevm (Mops/s)	Speedup
EOA transfers	310	4,820	15.5×
ERC-20 transfers	145	3,410	23.5×
DEX swaps	82	1,240	15.1×
AI Mining submissions	61	2,890	47.4×
FHE operations	23	1,070	46.5×
Complex contract chains	44	680	15.5×

Table 7: Per-type speedup (single H100)

8 Related Work

The LUX gpu-evm-whitepaper [4] first proposed GPU acceleration for EVM execution on the LUX C-Chain. CEVM extends that work with precompile-aware partitioning, AI Mining scheduling, and formal verification.

Monad [5] implements optimistic parallel execution targeting CPU multi-core ($\sim 10\times$ on 16 cores). **Sei v2** [6] uses optimistic concurrency control for CPU-parallel EVM ($5\text{--}12\times$). Both are limited by CPU core counts; CEVM achieves $23.8\times$ on a single GPU.

MegaETH [7] optimizes latency via specialized hardware sequencers with sequential execution, orthogonal to CEVM’s parallel throughput approach.

Solana’s Sealevel [8] requires upfront state-access declarations, incompatible with EVM semantics. **Aptos** and **Sui** use Move-based resource ownership for compiler-verified parallelism, requiring new programming models. CEVM preserves full EVM/Solidity compatibility.

The KEVM project [9] provides complete EVM formal semantics in the K framework; CEVM’s Lean 4 proofs complement KEVM by covering the parallel execution model.

9 Conclusion

CEVM achieves 2.3 Gops/sec on a single GPU— $23.8\times$ over the fastest CPU EVM—while preserving full determinism through conflict detection and sequential re-execution. Lean 4 verification provides mathematical guarantees that parallel execution produces identical state roots to sequential execution. Multi-GPU linear scaling to 8.9 Gops/sec on four H100s meets the throughput demands of Zoo’s AI-native L2 workloads.

Zoo Network’s integration of CEVM with DEX VM, FHE VM, and AI Mining VM establishes a template for GPU-accelerated blockchain execution that maintains EVM compatibility while unlocking throughput previously exclusive to centralized systems.

Future work includes intra-precompile parallelism for DEX and FHE VMs, predictive conflict

Phase	CPU (ms)	cevm H100 (ms)
Dependency analysis	—	1.2
State snapshot to GPU	—	3.8
Parallel execution	—	18.4
Write-log merge + re-exec	—	5.0
Sequential execution (all)	847.0	—
State root computation	12.3	12.3
Total (10K txns)	859.3	40.7

Table 8: Per-block latency breakdown

analysis via ML [3], ZK proof co-scheduling on GPU [10], and cross-L2 parallel execution across LUX subnets [11].

References

- [1] G. Wood, “Ethereum: A secure decentralised generalised transaction ledger,” Ethereum Project Yellow Paper, 2014.
- [2] Z. Kelling, “Zoo Network: Decentralized AI training and inference layer with semantic optimization,” Zoo Labs Foundation, 2025.
- [3] Z. Kelling, “Zoo Network tokenomics: Fair distribution through complete airdrop,” Zoo Labs Foundation, 2025.
- [4] Lux Partners, “GPU-accelerated EVM execution for high-throughput blockchain networks,” Lux Network Whitepaper Series, 2025.
- [5] Monad Labs, “Monad: A high-performance EVM-compatible L1 with parallel execution,” 2024.
- [6] Sei Labs, “Sei v2: The first parallelized EVM blockchain,” 2024.
- [7] MegaETH, “MegaETH: Real-time blockchain with sub-millisecond latency,” 2024.
- [8] A. Yakovenko, “Solana: A new architecture for a high performance blockchain,” 2020.
- [9] E. Hildenbrandt et al., “KEVM: A complete formal semantics of the Ethereum Virtual Machine,” IEEE CSF, 2018.
- [10] Z. Kelling, “Fully homomorphic encryption on Zoo Network,” Zoo Labs Foundation, forthcoming.
- [11] Z. Kelling, “Decentralized exchange architecture on Zoo L2,” Zoo Labs Foundation, forthcoming.