



Proof of AI: A Consensus Mechanism for Verifiable Compute Contribution

The execution-proof substrate for decentralized
AI research, and how a Zoo Layer-3 settles it

Antje Worring & Zach Kelling

Zoo Labs Foundation

June 2026

Abstract

We introduce **Proof of AI** (PoAI), a consensus mechanism for networks where participants contribute AI compute rather than hash power or stake. Bitcoin’s hash is *useless* work with $O(1)$ verification; PoAI replaces it with *useful* work—genuine AI execution—plus a proof that is a *side effect* of that execution and cheaply checkable, requiring *software only*: no trusted execution environment and no GPU attestation. The recomputable “hash everyone can check” becomes an activation-trace commitment that cannot be produced without performing the genuine matrix multiplications. We give one primitive realizing four execution-level proofs over a workload tag—Proof-of-Embedding, Proof-of-Inference, Proof-of-Training, and a privacy-preserving Proof-of-Contribution—by committing the computation trace and spot-checking its matmuls with Freivalds’ algorithm over the int8 path’s exact integer accumulator, which yields a bit-exact check with soundness error $1/p \approx 2^{-61}$ per vector (2^{-122} at $k=2$) and *zero* false rejections. We bind every proof to its task with a domain-separated, beacon-fresh, measured-model commitment, and prove anti-substitution and anti-replay for it. We position the scheme as Bitcoin’s shared, software-only difficulty—the Hamiltonian Market Maker *prices* heterogeneous compute while the execution proof *verifies* it, two orthogonal dials—and give a single profile gate mapping tiers (zkML, M -of- N confidential compute, optimistic-Freivalds, redundancy) to required proof strength, designed so that once the gate is engaged “no valid compute proof” means “no mint.” We motivate the construction with a concrete attack on a deployed settlement engine that today mints for a one-bit guess, because a signature proves authorship and a plurality proves agreement, but neither proves computation. We retain TEE attestation, multi-party quality scoring, and BLS aggregation as the strong and pricing layers; the execution proofs are what bind reward to genuine work.

Status, stated up front. This paper proves a *primitive*, builds the *system* around it, and specifies the *residual* that remains design-only; we mark which is which in the abstract rather than burying it, and Section 2 gives the per-component ledger before the construction begins. What is *built and tested* is now the whole load-bearing path, not just the core: (i) the exact-integer Freivalds matmul verifier and its keccak-Merkle activation-trace transcript, a proof-bearing forward (**ProvableLinear**) whose returned value *is* the committed exact-integer result, and a typed full-graph trace (**GraphTrace**) that binds attention, MoE, and—with one transpose op—a training step, all in the Rust engine (**hanzo-engine**: 27 unit tests + 11 adversarial red-team tests, 38/38 passing); (ii) the canonical Go verifier (**luxfi/crypto/poi**, 23/23) that the chain consumes, byte-identical to the prover under a pinned cross-language challenge vector; and (iii) the on-chain enforcement (**luxfi/standard** Proof-of-AI contracts: 87/87 across the eight Proof-of-AI suites), where **ComputeWitnessLib** runs Freivalds and the Merkle fold in pure Solidity, **OptimisticEvidence** pays a bond to a challenger *only* on a verified fraud, and the **AIMiner** mint gate *fails closed*: with no compute-verifier wired it reverts (**ProofGateUnavailable**), so a forged receipt root alone no longer mints. Two honest properties survive the build: a fabricated matmul is caught except with probability $\leq 2^{-122}$ at two challenge vectors because Freivalds runs over an exact integer accumulator, and an honest prover is *never* slashable. What remains *designed, not built* is the genuinely smaller residual: the RepOps pinned-reduction floating-point tier (the quantized integer tier is what ships), and emission coverage of architecture families beyond those exercised by the test suite (dense/SwiGLU, attention with a fixed-point integer softmax, an MoE block, a training step). The honest one-line summary: *the soundness primitive is proven, the unfakeable-mining system it enables is built and tested end to end, and the named residual is a floating-point tier and wider architecture coverage.*

The Zoo reading. ZOO Labs Foundation is the Layer-3 application and research layer (decentralized AI and decentralized science, governed through **zips.zoo.ngo**) that *consumes* this stack. Section 18 shows how a ZOO L3—BEUGA and its peers—registers in an on-chain directory (**L3Registry**) and settles PoAI-verified work down to the ZOO L2 and the Lux L1, so that proven inference, embedding, and training become the reproducibility substrate for decentralized science: a result is not “trust me, I ran the model” but a transcript whose root anyone can open and re-check.

Keywords: consensus, proof of useful work, AI compute verification, Freivalds’ algorithm,

1 Introduction

Blockchain consensus mechanisms face a fundamental challenge: how to fairly allocate block production rights and rewards. Proof of Work (PoW) solves this by requiring computational effort, but the computation serves no purpose beyond securing the chain—an estimated 150 TWh annually for Bitcoin alone. Proof of Stake (PoS) reduces energy waste but creates plutocratic dynamics where the wealthy accumulate more wealth.

The rise of AI compute networks presents an opportunity to align consensus with useful work. AI training and inference require substantial compute resources; if this compute could simultaneously secure a blockchain, we would achieve both useful work and decentralized consensus.

1.1 Challenges

Designing POAI requires solving several challenges:

1. **Verification:** How can validators confirm that claimed AI work was actually performed?
2. **Quality Assessment:** Not all AI compute is equally valuable—training a useful model differs from generating noise.
3. **Sybil Resistance:** Without proof of identity, adversaries could claim credit for others’ work.
4. **Efficiency:** Verification must be cheaper than the original computation.
5. **Incentive Compatibility:** Rational nodes must be incentivized to perform high-quality work.

1.2 Contributions

This paper makes the following contributions:

1. **PoAI Protocol:** A complete consensus mechanism combining TEE attestation, quality scoring, and BLS aggregation (Section 4)
2. **Compute Verification:** Novel techniques for verifying AI workloads without re-executing them (Section 5)
3. **Quality Scoring:** Multi-party evaluation protocol for assessing AI output value (Section 6)
4. **Security Analysis:** Proofs of Byzantine fault tolerance and incentive compatibility (Section 7)
5. **Evaluation:** Experimental results demonstrating scalability and efficiency (Section 8)
6. **Execution-level proofs (software-only), built end to end:** We close the gap that signature-plus-plurality leaves open—paying for a guess rather than for computation (Section 9)—with one primitive yielding four proofs (Section 11), a software-only Freivalds verifier with a soundness theorem (Section 12), a domain-separated task binding with anti-replay

and anti-substitution lemmas (Section 10), the Bitcoin shared-difficulty correspondence (Section 13), a single-gate tier matrix (Section 14), and a user-sovereignty trust spectrum (Section 15). The verifier no longer stands alone: a *running* forward pass emits a keccak-Merkle activation-trace transcript (`poi_transcript.rs`), a proof-bearing linear returns exactly its committed exact-integer result (`poi_forward.rs`), a typed graph binds the whole pass including a training step (`poi_graph.rs`), the Go verifier the chain consumes is byte-identical (`luxfi/crypto/poi`), and the on-chain gate verifies the witness in pure Solidity and slashes a bond only on proven fraud (`ComputeWitnessLib`, `OptimisticEvidence`, `AIMiner`).

7. **Settlement to a Zoo L3, and the DeSci substrate:** We show how a ZOO Layer-3 (BELUGA and its peers) consumes the execution proof: each L3 registers in an on-chain directory and settles PoAI-verified work down to the ZOO L2 and the LUX L1, turning proven inference, embedding, and training into the reproducibility substrate for decentralized science (Section 18). A scientific claim backed by a transcript root is re-checkable by anyone; “I ran the model” becomes a proof, not an assertion.
8. **A coverage *design* for modern LLMs, with a nonempty tested floor:** We show that “support every modern LLM” is achieved by *tiering* verification to the arithmetic—never by a floating-point Freivalds, which is unsound (a tolerance corridor) and unusable (cross-hardware false-reject) at once—and that floating-point GEMMs become bit-exact by pinning the reduction order (Section 16). We map the ZOO model zoo (dense, mixture-of-experts, linear-recurrent/state-space, multimodal vision, audio, and diffusion) to proof obligations and prove that beyond base matmul checking exactly *two* new core techniques suffice—a routing proof (*built and tested*, `poi_graph.rs`) and a sequential-chaining proof (*designed*, for SSM and diffusion latents)—Section 17. The tested floor is nonempty: a SwiGLU FFN, the attention spine with a deterministic integer softmax, an MoE block, and a training step all emit and verify today. We report the build status honestly and up front (Section 2, with the full ledger in Section 19): the verifier core, the transcript, the proof-bearing forward, the graph composition, and the on-chain enforcement are implemented and tested (38 Rust, 23 Go, 87 Solidity tests passing), while the RepOps floating-point tier and emission of the long-tail architectures remain designed.

The reader is owed the built-vs-designed boundary before the construction, not after: Section 2 states it immediately, and each execution-level section below restates the status of its own claims. In one line: *we prove a sound verification primitive, build and test the end-to-end emission and on-chain enforcement that make AI mining unfakeable, settle it to a ZOO L3, and name the floating-point tier and wider architecture coverage as the remaining work.*

2 What Is Built vs. What Is Designed

A reader should not have to reach the implementation section to learn which of this paper’s claims are backed by running code and which are specifications awaiting implementation. We therefore state the boundary here, before the construction, and repeat the discipline at each section that makes an execution-level claim. The full per-result ledger is Section 19; the two companion documents—the soundness companion [24] and the emission/enforcement companion [25]—carry the same boundary from the proof side and are consistent with this table. This section is that table, surfaced.

The honest one-line summary. The *soundness primitive* is proven and shipped; the end-to-end *execution-proof system* it enables is now built and tested, not a roadmap; and the residual is a

named, smaller pair (a floating-point tier, and emission coverage of architectures beyond the tested families). The matmul verifier that catches a fabricated output with probability $\geq 1 - 2^{-122}$ (at two challenge vectors) and never false-rejects honest integer work is implemented and unit-tested. The wire that the earlier draft of this paper marked “designed, not built”—a *running* forward pass committing a streamed activation-trace transcript that a challenger opens and Freivalds-checks against the committed operands—is built: it emits from a real quantized forward (`ProvableLinear`), composes across a whole typed graph (`GraphTrace`, including a training step), is verified on-chain in pure Solidity (`ComputeWitnessLib`), and gates a bond-slashing optimistic state machine (`OptimisticEvidence`). The properties this buys—“a fabricated matmul is provable” and “an honest prover is never slashable”—are theorems with passing adversarial tests behind each, across three languages with byte-identical commitments.

Three consequences a reviewer can grep for. A peer reviewer who reads the repositories will find exactly these, and we state them so they are unmistakable:

1. **The verifier now has callers.** The earlier draft’s caveat—“the verifier has no production callers”—no longer holds. `crypto/poi`’s `VerifyMulti` and the engine’s `poi.rs` are invoked end to end: `poi_transcript.rs` opens a beacon-selected matmul of a real `ProvableLinear/GraphTrace` forward and Freivalds-checks it, the Solidity `ComputeWitnessLib` runs the identical check on-chain, and `OptimisticEvidence.challenge` calls `provesFraud` to slash a bond. The `AIMiner` mint gate *reverts* when no compute-verifier is wired (`ProofGateUnavailable`), so the proofless-mint default the earlier draft flagged is closed: out of the box a proof is mandatory, not opt-in.
2. **The “exact, software-only” claim is about a kernel the engine ships in the proof path.** `ProvableLinear` computes $y = \text{dequant}(\tilde{x} \cdot W^\top)$ on a *whole-K exact-i64 accumulator* (`poi.rs::exact_matmul`), commits exactly those operands, and dequantizes the committed C —so the returned value *is* the committed computation, with no floating-point reduction in the committed trace and zero cross-hardware drift (the float-shadow unsoundness the earlier draft warned of is avoided by construction). The throughput cost of the exact accumulator is paid only on proof-bearing inferences; the fast `f8q8` path is unchanged for ordinary inference.
3. **Provable-today architecture coverage is no longer 0%.** The test suite emits and verifies a real SwiGLU FFN (`gate` $\rightarrow$ `up` $\rightarrow$ `down`), the attention spine ($QK^\top \rightarrow \text{scale} \rightarrow \text{deterministic integer softmax} \rightarrow \cdot V$), a mixture-of-experts block (`router` $\rightarrow$ `top-k` $\rightarrow$ `gather` $\rightarrow$ expert matmuls $\rightarrow$ `combine`), and a single SGD training step (`forward` + `backward` + `update`) as one typed graph—so the families with a committed, verified transcript today are nonempty. What remains a *designed* ceiling rather than a deployed floor is the breadth: the fp16/bf16/fp8 tier (RepOps) and the architecture families not yet wired to emit (the sequential-chaining obligation for state-space and diffusion models, and the long tail of the model zoo). “Full modern-LLM coverage” names the design’s reach; the tested set names its present floor, and the floor is no longer empty.

None of this weakens the mathematics. The Freivalds soundness theorem, the binding’s anti-replay/anti-substitution lemmas, the floating-point no-go, and the whole-graph composition theorem are correct as stated and are the Bitcoin-whitepaper-grade contribution of the paper. What the table fixes, relative to the earlier draft, is the *tense* of the system claims: we have proven a sound primitive, shipped its core, *and built and tested* the emission and on-chain enforcement that turn that primitive into an unfakeable settlement gate—leaving a named floating-point tier and wider architecture emission as the honest remaining work.

3 Background

3.1 Proof of Useful Work

Prior attempts at useful PoW include:

- **Primecoin** [1]: Finding prime number chains
- **Gridcoin** [2]: BOINC scientific computing
- **FoldingCoin**: Protein folding simulations

These approaches suffer from limited useful work domains and difficulty of verification. Our work extends to general AI computation.

3.2 Trusted Execution Environments

TEEs provide hardware-enforced isolation:

Definition 3.1 (Trusted Execution Environment). *A TEE is a hardware-protected region where:*

1. *Code and data are encrypted in memory*
2. *Execution cannot be observed or modified by the host OS*
3. *Remote attestation proves what code ran and its output*

Major TEE implementations:

- Intel SGX / TDX
- AMD SEV-SNP
- ARM TrustZone / CCA
- NVIDIA Confidential Computing

3.3 BLS Signatures

BLS signatures enable efficient aggregation:

Definition 3.2 (BLS Signature Aggregation). *Given signatures $\{\sigma_i\}_{i=1}^n$ from signers $\{pk_i\}$ on message m :*

$$\sigma_{agg} = \prod_{i=1}^n \sigma_i \quad (1)$$

Verification: $e(\sigma_{agg}, g) = e(H(m), \sum_i pk_i)$

This enables $O(1)$ verification of n signatures.




Figure 1: PoAI epoch structure.

4 PoAI Protocol

4.1 System Model

Participants:

- **Workers:** Perform AI computation within TEEs
- **Validators:** Verify attestations and score quality
- **Aggregators:** Combine attestations for on-chain submission

Assumptions:

- At least $2f + 1$ of $3f + 1$ validators are honest
- TEE hardware is not compromised (standard TEE assumption)
- Network is partially synchronous (bounded delay after GST)

4.2 Protocol Overview

PoAI operates in epochs:

1. **Work Phase:** Workers perform AI computation in TEEs
2. **Attestation Phase:** Workers submit TEE attestations
3. **Scoring Phase:** Validators evaluate work quality
4. **Aggregation Phase:** Attestations aggregated via BLS
5. **Finalization:** Block produced, rewards distributed

4.3 Work Phase

Workers execute AI tasks within TEEs:

The attestation A includes:

- TEE measurement (code hash)
- Input/output hashes
- Timestamp and nonce
- Hardware signature

Algorithm 1 Worker Execution

```
1: Input: Task  $T$ , model  $M$ , data  $D$ 
2: Output: Result  $R$ , attestation  $A$ 
3: Enter TEE enclave
4: Load verified model  $M$  (hash checked)
5:  $R \leftarrow \text{Execute}(M, T, D)$ 
6:  $h \leftarrow \text{Hash}(T \| M \| D \| R)$ 
7:  $A \leftarrow \text{TEE.Attest}(h)$ 
8: Exit enclave
9: return  $R, A$ 
```

Algorithm 2 Quality Scoring

```
1: Input: Attestation  $A$ , sample outputs
2: Output: Quality score  $q \in [0, 1]$ 
3: Verify TEE signature authenticity
4: if signature invalid then
5:   return  $q = 0$ 
6: end if
7: Sample subset of outputs for evaluation
8:  $q_{\text{correct}} \leftarrow \text{VerifyCorrectness}(\text{samples})$ 
9:  $q_{\text{useful}} \leftarrow \text{AssessUtility}(\text{samples})$ 
10:  $q_{\text{novel}} \leftarrow \text{CheckNovelty}(\text{samples})$ 
11:  $q \leftarrow \alpha q_{\text{correct}} + \beta q_{\text{useful}} + \gamma q_{\text{novel}}$ 
12: return  $q$ 
```

4.4 Attestation Phase

Workers submit attestations to the network:

$$A = (w, T, h_{\text{in}}, h_{\text{out}}, \tau, \text{sig}_{\text{TEE}}) \quad (2)$$

where:

- w : Worker identifier
- T : Task type (inference, training, extraction)
- $h_{\text{in}}, h_{\text{out}}$: Input/output hashes
- τ : Timestamp
- sig_{TEE} : TEE hardware signature

4.5 Scoring Phase

Validators evaluate work quality:

Quality components:

- q_{correct} : Output correctness (deterministic verification)
- q_{useful} : Utility for downstream tasks (benchmark evaluation)

- q_{novel} : Information gain over existing knowledge

4.6 Aggregation Phase

Validator scores are aggregated using weighted median:

$$Q(A) = \text{WeightedMedian}(\{q_v(A)\}_{v \in V}, \{s_v\}) \quad (3)$$

where s_v is validator v 's stake.

Attestations are BLS-aggregated:

$$\sigma_{\text{block}} = \prod_{A \in \text{block}} \text{BLS.Sign}(sk_v, H(A)) \quad (4)$$

4.7 Reward Distribution

Block rewards are distributed proportionally:

$$R_w = R_{\text{total}} \cdot \frac{\sum_{A \in \mathcal{A}_w} Q(A) \cdot C(A)}{\sum_{A \in \text{block}} Q(A) \cdot C(A)} \quad (5)$$

where:

- $\mathcal{A}_w$: Attestations from worker w
- $Q(A)$: Quality score
- $C(A)$: Compute cost (measured in FLOPs)

5 Compute Verification

5.1 TEE-Based Verification

The core verification relies on TEE attestation:

Theorem 5.1 (TEE Attestation Soundness). *Under the TEE security assumption (hardware not compromised), if attestation A verifies, then the claimed computation was performed within the enclave with inputs/outputs matching the hashes.*

5.2 Efficient Re-Execution

For tasks where TEEs are unavailable, we use probabilistic re-execution:

Theorem 5.2 (Detection Probability). *With sampling rate p and n claims, a cheating worker is caught with probability $1 - (1 - p)^n$.*

For $p = 0.1$ and $n = 100$, detection probability exceeds 99.99%.

5.3 Hybrid Verification

We combine TEE attestation with probabilistic re-execution:

$$V(A) = \begin{cases} \text{TEE.Verify}(A) & \text{if TEE available} \\ \text{ProbVerify}(A, p) & \text{otherwise} \end{cases} \quad (6)$$

Algorithm 3 Probabilistic Verification

```
1: Input: Claimed output  $y$ , task  $T$ , sampling rate  $p$ 
2: Output: Verified / Rejected
3: if Random()  $< p$  then
4:    $y' \leftarrow \text{ReExecute}(T)$ 
5:   if  $y' \neq y$  then
6:     Slash worker stake
7:     return Rejected
8:   end if
9: end if
10: return Verified
```

Algorithm 4 Multi-Party Quality Evaluation

```
1: Input: Attestation  $A$ , validator set  $V$ 
2: Output: Consensus quality  $Q$ 
3: Each validator  $v$  computes  $q_v(A)$  independently
4: Validators commit  $c_v = \text{Hash}(q_v \| r_v)$ 
5: Wait for all commits
6: Validators reveal  $(q_v, r_v)$ 
7: Verify  $c_v = \text{Hash}(q_v \| r_v)$  for all  $v$ 
8:  $Q \leftarrow \text{TrimmedMean}(\{q_v\}, 0.1)$ 
9: return  $Q$ 
```

6 Quality Scoring

6.1 Multi-Party Evaluation

Quality scoring uses multi-party computation to prevent manipulation:

6.2 Evaluation Metrics

Inference Tasks:

$$q_{\text{inf}} = \frac{\text{correct outputs}}{\text{total outputs}} \cdot (1 - \text{latency penalty}) \quad (7)$$

Training Tasks:

$$q_{\text{train}} = \Delta_{\text{val}} \cdot e^{-\lambda \cdot \text{cost}} \quad (8)$$

where Δ_{val} is validation accuracy improvement.

Experience Extraction:

$$q_{\text{exp}} = \text{novelty} \cdot \text{applicability} \cdot \text{correctness} \quad (9)$$

6.3 Gaming Resistance

The trimmed mean resists Byzantine manipulation:

Theorem 6.1 (Byzantine Resistance). *With $f < n/3$ Byzantine validators, the trimmed mean with 10% trim differs from the honest median by at most $2\sigma/\sqrt{n-2f}$.*

7 Security Analysis

7.1 Byzantine Fault Tolerance

Theorem 7.1 (Safety). *If at most $f < n/3$ validators are Byzantine, no two honest validators finalize conflicting blocks.*

Proof. Block finalization requires $2f + 1$ valid signatures. With $n = 3f + 1$ total validators and at most f Byzantine, any two sets of $2f + 1$ signers overlap in at least one honest validator. An honest validator signs at most one block per height. $\square$

Theorem 7.2 (Liveness). *In partial synchrony, if at most $f < n/3$ validators are Byzantine, honest transactions are eventually included.*

Proof. After GST, message delays are bounded. With $2f + 1$ honest validators, quorum can always be formed. Leader rotation ensures eventual honest leader. $\square$

7.2 Incentive Compatibility

Theorem 7.3 (Truthful Quality Reporting). *Under the multi-party evaluation protocol, reporting true quality is a dominant strategy.*

Proof. The trimmed mean ignores extreme values. Lying increases variance but cannot shift the result beyond honest range. Deviating validators are identified and slashed. $\square$

Theorem 7.4 (No Free Riding). *Workers cannot claim rewards without performing computation.*

Proof. TEE attestation binds output to enclave execution, and the software-only execution proof binds reward to genuine work with no hardware-trust assumption: a worker that claims an output it did not compute commits an activation-trace whose matmul is caught by the Freivalds challenge except with probability $\leq 2^{-122}$ (Section 12; built and tested, Section 19), and the on-chain mint gate fails closed without a valid proof (AIMiner, Section 18). Probabilistic challenge catches cheaters with high probability and slashing the bond exceeds the expected gain. The execution proof is what makes this hold for an untrusted, hardware-free prover, not only inside a trusted enclave. $\square$

7.3 Sybil Resistance

Sybil attacks are mitigated through:

1. **Stake Requirement:** Validators must stake tokens
2. **Work Requirement:** Workers must perform verified computation
3. **Hardware Binding:** TEE attestations bound to physical devices

8 Experimental Evaluation

Status note. The throughput, overhead, and detection figures in this section are illustrative projections of the consensus design under the stated configuration, not measurements from a running testnet; we mark them as such so no reader mistakes a target for an observed result, in keeping with the up-front-honesty discipline of Section 2. The figures that are measured live elsewhere in this paper and its companions: the execution-proof test suites (Section 19: 38 Rust, 23 Go, 87 Solidity tests passing, run for this paper) and the single-node verification latency of a challenged




Figure 2: Reward distribution by quality. PoAI incentivizes high-quality work.

matmul slice reported in the emission companion [25]. The asymmetry those establish—verifying a challenged matmul is an order cheaper than recomputing it (Section 12)—is the load-bearing performance claim; the consensus-level numbers below are the design’s intended operating point.

8.1 Testnet Configuration

- **Nodes:** 100 validators, 500 workers
- **Hardware:** Mix of Intel SGX and AMD SEV-SNP
- **Network:** Geographically distributed (5 continents)
- **Workloads:** Inference (70%), training (20%), extraction (10%)

8.2 Throughput and Latency

8.3 Verification Overhead

8.4 Quality Scoring Accuracy

8.5 Reward Distribution

9 The Motivating Attack: Minting for a Guess

The protocol of Sections 4–5 pays workers for AI computation. This section makes precise a gap that any such protocol must close before its rewards mean anything: *a signature proves authorship, a plurality proves agreement, and neither proves computation*. We state the gap as a concrete attack against a deployed settlement engine, because it is the reason the execution-level proofs of Sections 11–14 exist.

9.1 The settlement engine, as deployed

The Lux A-Chain (`aivm`) settles AI tasks with a commit–reveal quorum. For a task t naming a model specification `modelSpec` and a prompt, each selected operator o runs (nominally) the model, obtains an output, and publishes a commitment

$$c_o = \text{keccak}(t \parallel \text{modelSpecHash} \parallel \text{promptHash} \parallel \text{outputHash} \parallel \text{embeddingHash} \parallel o \parallel \text{nonce}), \quad (10)$$

then later reveals the preimage. Reveal verification recomputes c_o from the revealed fields and the caller’s address and checks equality; a mismatch is rejected. Settlement tallies the revealed `outputHash` values, takes the *plurality* group (the statistical mode), and—if that group’s size meets a threshold θ —declares its hash canonical and pays `rewardPerOperator` to every operator in it from escrow, slashing those who withheld a reveal.

Two properties of this design are worth quoting from the implementation itself, because they are stated there without euphemism. The settlement substrate’s engine documentation records that

“...the settlement RESULT—did $\geq$ threshold staked operators independently submit the same `output_hash` under the same MODEL_SPEC—is what A-Chain consensus agrees on; validators never run the model.”

and the address binding in (10) is documented as

“the anti-copy / anti-front-run control: a peer who observes operator A’s commit cannot replay it as their own”

The first sentence is the safety premise of the entire mechanism; the second names exactly what the commitment proves. We take both at face value.

9.2 What is—and is not—proven

The commitment (10) binds an output hash to an *author* (the address o is inside the keccak preimage and is re-checked on reveal) and, because `nonce` and o are mixed in, resists copying and front-running. The plurality rule binds the canonical answer to *agreement* (how many distinct authors published the same hash). Neither operation references the model weights, the prompt tokens beyond their hash, or any artifact of the forward pass. We isolate the consequence.

Proposition 9.1 (Computation is unconstrained by authorship and agreement). *Let Π be a settlement engine whose verdict is a function only of (i) address-bound commitments of the form (10) and (ii) the multiset of revealed output hashes via a plurality rule with threshold θ . Then for any prompt and any model specification, a coalition of θ operators that never evaluates the model can cause Π to settle—and to pay each coalition member—on an output hash h^* of the coalition’s choosing, provided h^* is the plurality among revealed hashes.*

Proof. Each coalition member o_i ($1 \leq i \leq \theta$) chooses the same target hash h^* , samples a fresh `noncei`, and computes c_{o_i} by (10) with `outputHash` := h^* and its own address o_i . Each c_{o_i} is a valid, distinct, non-replayable commitment: it is well-formed (a keccak image), it is bound to the correct author (so reveal verification’s recompute-and-compare succeeds), and it differs from every other member’s commitment (distinct addresses and nonces), so the anti-copy control raises no objection. On reveal, all θ members publish h^* ; their group has size $\theta \geq \theta$ and, by hypothesis, is the plurality. Π declares h^* canonical and pays each member. No member evaluated the model; h^* was never required to be the model’s output on the prompt. $\square$ $\square$

Proposition 9.1 is not exotic. In the degenerate but common case of a *decision* task—a yes/no classification, a moderation verdict, an “approve/reject” agentic action—the output space is one bit, so h^* is one of two hashes. The cost-minimizing strategy is then explicit and requires *no model at all*:

Corollary 9.2 (The one-bit guess). *For a binary-output task with prior $\Pr[\text{answer} = 1] = q \geq \frac{1}{2}$, the strategy “commit and reveal the hash of the modal answer 1, never run the model” settles correctly with probability q at zero inference cost, and—when a coalition of size θ adopts it—settles and mints with probability 1 regardless of the true answer.*

The corollary is the headline: under authorship-plus-agreement, the network can be made to mint for a one-bit guess. A signature scheme cannot fix this, because the defect is not forgery—every signature and every commitment in the attack is genuine. The missing predicate is *“this output is the result of evaluating the named model on the named input,”* and that predicate is exactly what an execution-level proof supplies.

9.3 Why the existing fallbacks do not close it

Three mechanisms in the deployed stack might be hoped to help; none does, for a structural reason worth stating.

- **Re-execution by validators** is the natural fix, and the A-Chain documentation contemplates an optional “ M -of- N redundant” mode in which validators re-run inference. But this is precisely what the deployed engine *does not* do (“validators never run the model”), and naive re-execution costs $M \times$ the compute per task, which is why it was not made the default. Section 12 replaces $O(n^3)$ re-execution with an $O(n^2)$ *check*, making the fix affordable.
- **TEE attestation** (the path of Section 5) does bind output to enclave execution, and closes the gap *when present*. But it requires confidential-compute hardware, excludes CPU-only and consumer-GPU providers, and reduces the trust base to a vendor’s remote-attestation root. It is a strong path, not an egalitarian one (Section 14).
- **A GPU “proof” kernel** exists in the A-Chain backend, but it is a TEE-attestation envelope check—it verifies a signature over a hardware measurement—and it is never invoked by the settlement or mint path. It is, by construction, subject to Proposition 9.1: it proves authorship of an attestation, not execution of a model.

The remainder of this paper *constructs, proves sound, and builds* the missing predicate as a *software-only* artifact—no enclave, no special hardware—that is a side effect of the genuine computation, that cannot be produced without performing the matmuls, and that any participant can both emit and check against a network-shared difficulty. The predicate is no longer only on paper: the verifier, its emission from a real forward/graph pass, the off-chain challenger, and the on-chain fail-closed mint gate are implemented and tested (38 Rust, 23 Go, 87 Solidity tests; Section 19). So the attack above describes the settlement of a system that does *not* install the compute-proof gate—which the deployed contracts now refuse to be: **AIMiner** reverts with **ProofGateUnavailable** when no verifier is wired, so a proofless mint is impossible out of the box. The residual is the floating-point tier and wider architecture emission, not the closing of this attack.

10 The Binding: One Commitment for Four Proofs

Before specifying *how* a worker proves it ran a model (Sections 11–12), we fix *what* the proof is about. The binding is a single 32-byte value, **reportData**, that names the task, the measured model, the input, the output, and the runtime under which the computation must have occurred. Every proof system—TEE, optimistic-Freivalds, zkML, or M -of- N —attests to the same **reportData**; the proof system is a backend, and **reportData** is the join key. This is the decomplexing move: *what was computed* (the binding) is separated from *how we believe it* (the backend), so the two vary independently.

10.1 Domain-separated two-stage construction

The binding is built in two keccak stages with disjoint domain tags, mirroring the deployed `ComputeProofLib` library. Both tags are raw UTF-8 byte strings used as keccak prefixes (no length framing), and both are versioned:

$$\text{DOMAIN_CHALLENGE} = \text{"lux/aivm/compute-challenge/v1"}, \quad (11)$$

$$\text{DOMAIN_REPORT} = \text{"lux/aivm/compute-report/v1"}. \quad (12)$$

The *challenge* commits to everything fixed *before* the worker runs:

$$\text{ch} = \text{keccak}(\text{DOMAIN_CHALLENGE} \parallel \text{taskId} \parallel \text{intentID} \parallel \text{modelSpecHash} \parallel \text{inputHash} \parallel \text{openBlockHash} \parallel \text{operator}), \quad (13)$$

and the *report* commits the challenge together with the produced output and the runtime under which it was produced:

$$\text{reportData} = \text{keccak}(\text{DOMAIN_REPORT} \parallel \text{ch} \parallel \text{workloadType} \parallel \text{modelSpecHash} \parallel \text{inputHash} \parallel \text{outputHash} \parallel \text{runtimeMeasurement}). \quad (14)$$

The wire encoding is byte-pinned so that the Solidity precompile, the Go attestation precompile, and the Rust engine produce identical digests: `taskId` is a 32-byte big-endian integer; `intentID`, `modelSpecHash`, `inputHash`, `openBlockHash`, `outputHash`, and `runtimeMeasurement` are 32-byte words; and `operator` is a 20-byte address, *not* left-padded. The deployed library realizes (13)–(14) verbatim over `abi.encodePacked`; this paper adds the single field `workloadType` to (14) (Section 11) so that one binding serves all four proofs without overloading any other field.

10.2 The fields, and what each defends

Each field in (13)–(14) neutralizes a specific cheat.

modelSpecHash — a *measured weights* Merkle root, not a label.

This is the keccak Merkle root over the loaded, quantized model bytes (Section 12), *computed at load time from the weights the worker actually holds in memory*, not a human-readable model name. Binding the measured root defeats *cheap-model substitution*: a worker cannot answer with a 1B model while claiming the fee schedule of a 70B model, because the proof backend re-derives activations against the committed root and a different model yields a different root and different activations.

inputHash.

Commits the exact prompt/embedding input. A worker cannot answer an easier question than the one asked.

outputHash.

Commits the produced output (token ids for inference, the embedding vector for embedding, the LoRA-delta root for training). This is what the proof must show is the genuine result.

openBlockHash — the unpredictability beacon.

The hash of a recent block that is *not known when the task is posted*. Because `ch` mixes in `openBlockHash`, the challenge—and hence the random vector the Freivalds check derives from it (Section 12)—cannot be predicted or pre-baked. This is the anti-replay and anti-precomputation root.

operator — per-operator domain.

The provider’s address. Because each operator’s `ch` and `reportData` are distinct, a proof produced by operator A does not verify for operator B ; honest work cannot be copied and re-sold under another identity (the same anti-copy property the A-Chain commitment of Section 9 sought, now extended to the *proof* rather than only the output).

runtimeMeasurement — the determinism pin.

A digest of the runtime and sampler configuration under which the output is reproducible: the kernel set, tokenizer, quantization mode, temperature, and seed (Section 12). Pinning it makes “re-derive and compare” a well-defined operation across heterogeneous hardware.

10.3 Security of the binding

We model `keccak` as a random oracle $\mathcal{H}$ and quantify the two guarantees the construction must provide: a forged report cannot be replayed onto a different task, and it cannot be silently re-pointed to a different model or input. Both reduce to collision/preimage resistance, but the domain separation and the beacon make the reductions tight and the attack windows closed.

Lemma 10.1 (Anti-substitution). *Model $\mathcal{H}$ as a random oracle. If two reports share a `reportData` value but differ in any of `modelSpecHash`, `inputHash`, `outputHash`, `workloadType`, or (through `ch`) `taskId`, `intentID`, `openBlockHash`, or `operator`, then the adversary has computed a collision of $\mathcal{H}$. After Q oracle queries the probability of producing any such pair is at most $\binom{Q}{2} 2^{-256} \leq Q^2 2^{-257}$.*

Proof. `reportData` = $\mathcal{H}(x)$ where x is the encoding of (14), which contains `ch` = $\mathcal{H}(y)$ with y the encoding of (13). The packed encoding is injective on the listed fields: every field has a fixed byte length (`operator` is 20 bytes, all other components are 32 bytes, `workloadType` is a fixed-width tag) and the domain prefixes are constants of fixed length, so distinct field tuples yield distinct byte strings x (resp. y). Two reports with equal `reportData` and unequal x are a direct collision of $\mathcal{H}$; if x is equal but the challenge inputs differ, then `ch` collided, i.e. a collision of $\mathcal{H}$ on y . A random oracle answers each fresh query with a uniform 256-bit string, so any fixed pair of distinct inputs collides with probability 2^{-256} ; a union bound over the $\binom{Q}{2}$ query pairs gives the stated bound. $\square$ $\square$

The disjoint domain tags do real work in Lemma 10.1: without them a `challenge` digest and a `reportData` digest could in principle be made to coincide and be cross-used (a type-confusion between the two stages). Prefixing each stage with its own constant forces the two oracle “slices” apart, so a challenge can never be mistaken for a report and vice versa. The same construction generalizes to any future proof family by adding a new `workloadType` tag rather than a new field, keeping the binding stable.

Lemma 10.2 (Anti-replay / freshness). *Suppose `openBlockHash` is the hash of a block whose contents are unpredictable until height H_{open} , and that a task posted at height $H_{\text{post}} < H_{\text{open}}$ fixes `openBlockHash` := `blockhash`(H_{open}). Then any `reportData`, and any random challenge vector derived from it, that an adversary commits to before height H_{open} is independent of the value that will be required, except with the probability that the adversary predicts `blockhash`(H_{open}). Under the random-oracle model for the block-hash beacon this prediction probability is 2^{-256} per guess.*

Proof. `ch` is a function of `openBlockHash` through $\mathcal{H}$ (Eq. 13), and `reportData` is a function of `ch` through $\mathcal{H}$ (Eq. 14); the Freivalds vector is in turn derived from `ch` (Section 12). A value committed before H_{open} is a function of information available before H_{open} , which by hypothesis is independent of `blockhash`(H_{open}). By the random-oracle property of the beacon, the only way a

pre- H_{open} commitment can equal the post- H_{open} requirement is if it embeds a correct guess of the unpredictable beacon output, an event of probability 2^{-256} per attempt. $\square$ $\square$

Lemma 10.2 is what prevents a prover from *pre-baking* a proof: it cannot precompute activations or a Freivalds response for a challenge it does not yet know, and it cannot replay a stale proof, because the beacon ties the challenge to a block that did not exist when a stale proof was made. Together with Lemma 10.1 and the per-operator field, the binding guarantees that a valid proof is *about this task, this model, this input, this operator, and this freshly-sampled challenge*—the necessary frame within which Section 12’s soundness theorem then guarantees the computation itself was performed.

11 One Primitive, Four Proofs

We now specify the work that the binding of Section 10 attests to. There are four AI workloads a node can be paid for—embedding, inference, training, and contribution—and we make a deliberate claim of *simplicity*: they are not four mechanisms but one. Each is the same scheme—*commit the computation trace, then spot-check the matmuls under a determinism contract*—with the binding’s `workloadType` field selecting which trace is committed and what the output hash ranges over. We tag the four

$$\text{workloadType} \in \{\text{EMBED}, \text{INFER}, \text{TRAIN}, \text{CONTRIBUTE}\} \quad (15)$$

and name the corresponding proofs PoE (Proof-of-Embedding), PoI (Proof-of-Inference), PoT (Proof-of-Training), and PoC (Proof-of-Contribution). The unifying observation is that the dominant cost of every modern neural workload is dense matrix multiplication, including the backward pass; so a single primitive that verifies matrix products (Section 12) covers the expensive part of all four, and the cheap remainder is committed and recomputed deterministically.

Status of this section. The four proofs are *specified* here and the machinery they share is *built*: each reduces to the Freivalds primitive plus the activation-trace commitment and forward-pass emission, all implemented and tested (`poi.rs`, `poi_transcript.rs`, `poi_forward.rs`, `poi_graph.rs`; Section 19). Proof-of-Inference and Proof-of-Embedding emit and verify today (a real SwiGLU forward and an MoE/attention graph are exercised by the suite); Proof-of-Training emits and verifies as a longer graph of the same ops plus a transpose (`test_training_step_is_just_a_graph`); the privacy-preserving Proof-of-Contribution adds the X-Wing prompt seal and attestation binding ([25]). The definitions below say what each proof *is* and what it establishes; the operational reading (“this output is certified to be the model’s output”) holds wherever the emitting forward is wired—which, for the families in the test suite, is now.

Definition 11.1 (Compute-proof primitive). *Fix the binding `reportData` of Section 10. A compute-proof for a workload is the tuple $(\text{modelWeightsRoot}, \text{activationTraceRoot}, \pi)$ where `modelWeightsRoot` = `modelSpecHash` is the measured weights root, `activationTraceRoot` is the root of the streamed trace of all intermediate tensors produced during the computation (Section 12), and π is the backend witness (a TEE quote, a zk proof, or—in the software-only case—nothing beyond the commitments, the verification being an interactive or on-challenge re-derivation). A verifier accepts iff (i) `reportData` matches the task binding, (ii) `modelWeightsRoot` is a governance-accepted model spec, and (iii) the backend attests that, under the pinned runtime, evaluating the model whose weights hash to `modelWeightsRoot` on the input whose hash is `inputHash` yields the trace committed by `activationTraceRoot` and the output committed by `outputHash`.*

11.1 PoE: Proof-of-Embedding

An embedding is a forward pass that ends in a pooled hidden vector rather than a decoded token. The prover runs the encoder, streams every intermediate tensor into the activation trace, and sets

$$\text{outputHash} = \text{keccak}(\text{quantize}(v)), \quad (16)$$

where v is the final embedding vector in its canonical fixed-point encoding. The proof is identical in structure to inference: the forward pass is a sequence of matrix products interleaved with normalization and activation; the matmuls are Freivalds-checked, the rest is recomputed. PoE is the cheapest of the four (no autoregressive loop) and is the workload behind retrieval, clustering, and the held-out evaluation used by PoC below.

11.2 PoI: Proof-of-Inference

Inference is autoregressive decoding. For a prompt of n tokens generating m output tokens, the prover performs $n + m$ forward passes (with key/value caching), streams the activation trace across all of them, and sets

$$\text{outputHash} = \text{keccak}(\text{id}_1 \parallel \text{id}_2 \parallel \cdots \parallel \text{id}_m), \quad (17)$$

the keccak of the produced token-id sequence. Under the determinism contract (temperature 0, fixed tie-break, pinned kernels and tokenizer; Section 12), the decode is a deterministic function of the prompt and the model, so a verifier re-deriving any sampled step obtains the identical token. The per-token logits are the output of the final matmul of each pass and are thus covered by the same Freivalds check as every other layer; the argmax that selects the token is an element-wise comparison the verifier recomputes from the committed logits. PoI is the proof *designed* to make the receipt of Section 9’s engine mean “this output is the model’s output” rather than merely “ θ operators agreed on this output”—a property it delivers once the trace it checks is actually emitted (Section 19), and not before.

11.3 PoT: Proof-of-Training

Training adds a backward pass and an optimizer step to the forward pass, and the key structural fact is that *backpropagation is also matrix multiplication*. For a linear layer with pre-activation $z = Wa$, the gradients are

$$\frac{\partial \mathcal{L}}{\partial W} = \frac{\partial \mathcal{L}}{\partial z} a^\top, \quad \frac{\partial \mathcal{L}}{\partial a} = W^\top \frac{\partial \mathcal{L}}{\partial z}, \quad (18)$$

both of which are dense matrix products of the same shape class as the forward matmul. Therefore the single Freivalds primitive of Section 12 covers the gradients exactly as it covers the forward activations: the prover streams the forward activations, the upstream gradients $\partial \mathcal{L} / \partial z$, and the weight gradients $\partial \mathcal{L} / \partial W$ into the trace, and the verifier spot-checks both (18) products. The optimizer step (for AdamW: bias-corrected first/second moments and an element-wise update) is element-wise and cheap, so it is committed and deterministically recomputed rather than Freivalds-checked. The output is the produced weight delta:

$$\text{outputHash} = \text{root}_{\text{BF16}}(\Delta W), \quad (19)$$

the gym’s canonical BF16 LoRA-delta root—a backend-independent serialization (u64 header length, JSON header of per-tensor name/shape/offsets, then BF16 body) whose digest is byte-identical across MLX, CUDA, ROCm, MPS, and CPU. Because the serialization is canonical, two honest provers training the same step on different hardware commit the same root, and a verifier on any backend can check it.

11.4 PoC: Proof-of-Contribution

PoC is the fourth proof and the privacy-preserving one. It is the proof a node earns for *improving a shared model with private data*: it is PoT that the delta was trained, plus a *Proof-of-Improvement* that the delta is useful, plus a privacy guarantee that the raw training data never left the contributor. It is the gym’s DeltaSoup—federated LoRA aggregation with Byzantine-robust trimmed means and contributor reputation—made verifiable. We define it as a conjunction.

Definition 11.2 (Proof-of-Contribution). *A PoC for a base model M and an accepted held-out evaluation set E (named by a committed hash `evalHash`, fixed by governance and beacon-bound so the contributor cannot choose it adversarially) is the triple $(\sigma_{\text{PoT}}, \sigma_{\text{impr}}, \sigma_{\text{priv}})$ where:*

1. σ_{PoT} (**you did the work**). *A PoT proof (Section above) that ΔW was produced by the claimed training computation, with `outputHash` = $\text{root}_{\text{BF16}}(\Delta W)$.*
2. σ_{impr} (**it is useful**). *A pair of PoE/POI proofs evaluating M and $M \oplus \Delta W$ on E , showing that the delta lowers loss (or raises a benchmark) by at least a governance margin δ : $\mathcal{L}_E(M \oplus \Delta W) \leq \mathcal{L}_E(M) - \delta$. Both evaluations are themselves execution-proven on the same committed E , so the improvement claim is not self-reported—it is the difference of two verified forward passes.*
3. σ_{priv} (**the data stayed private**). *A commitment to the training-data policy that reveals only that the data satisfied policy, not the data itself: the raw data never leaves the contributor; it is committed by hash (and optionally DP-noised with a committed noise schedule), and what is published is the delta plus a proof that the gradient was computed over data consistent with the committed policy. The contributor shares useful delta weights, never the corpus.*

A verifier accepts the contribution iff all three hold and the binding’s `workloadType` = CONTRIBUTE.

PoC is the “share data anonymously, fine-tune privately, publish only the useful delta” path of AI consensus, and it is where the protocol’s incentives point at the thing the world actually wants: not raw compute cycles, but *measured model improvement*. The improvement margin δ is the analogue, for training, of the correctness check for inference—it is what makes the work *useful* rather than merely *expended*. A contributor who submits a delta that does not improve E fails σ_{impr} and earns nothing, exactly as a miner who submits a non-improving guess fails Corollary 9.2’s implicit verification.

Remark 11.1 (Privacy is orthogonal to verifiability). The privacy guarantee σ_{priv} and the execution guarantees $\sigma_{\text{PoT}}, \sigma_{\text{impr}}$ are independent. The execution proofs show the delta is the result of a genuine training computation on data hashing to the committed value; the privacy mechanism ensures that value reveals nothing about the data. One can strengthen σ_{priv} (e.g. to a zero-knowledge proof that the committed data satisfied a content policy) without touching the Freivalds machinery, and vice versa. This is the separation-of-concerns that lets a deployment dial privacy and verification strength independently (Section 15).

11.5 Why four proofs and not four mechanisms

The payoff of Definition 11.1 is that the protocol’s verification surface is *one* primitive (Freivalds-checked matmuls plus deterministic recomputation of the cheap remainder) and *one* binding (Section 10), regardless of which workload a node sells. Adding a fifth workload—say, a diffusion sampling step, or a tool-augmented agentic trajectory—requires only a new `workloadType` tag and a specification of what the trace and output hash range over; it requires no new verifier and no change to the binding. This is the orthogonality discipline applied to proofs: each workload occupies its own tag, each is independently complete, and none is bolted onto another.

12 The Software-Only Execution Proof (Design)

This section constructs the artifact that is *designed* to play, for Proof-of-AI, the role the hash plays for Bitcoin: a value that is a *side effect* of the genuine computation, that is *cheap to verify* relative to producing it, and that *cannot be produced without performing the work*—but where the work is real AI inference and the verification requires *only software*: a model file and a CPU, no trusted execution environment, no GPU attestation. The construction has three parts: a commitment to the computation trace (so there is something to check against), a probabilistic verifier for the matrix products (Freivalds’ algorithm, made bit-exact by the integer quantization path), and a determinism contract (so that “re-derive and compare” is well-defined across heterogeneous hardware). We then prove soundness.

Status of this section. The probabilistic verifier (the Freivalds core), the determinism contract’s two nontrivial clauses, the activation-trace commitment of Section 12.1, and its emission from a real forward/graph pass are all *built and tested* (Section 19): a `ProvableLinear/GraphTrace` forward streams the trace, a challenger opens a beacon-selected matmul and Freivalds-checks it (`poi_transcript.rs`), and the on-chain witness re-checks it in pure Solidity (`ComputeWitnessLib`). The soundness theorems below are therefore both correct as mathematics *and* the specification the running code implements; “cannot be faked” is operational for the families the suite exercises. We mark the one place (the integer-accumulator remark) where the proof’s arithmetic assumption requires the exact whole- K i64 accumulator (`exact_matmul`, the proof path), which is stronger than the $f8q8$ fast path used for ordinary, non-proof-bearing inference.

12.1 The activation-trace commitment

A forward (or backward) pass produces a sequence of intermediate tensors $T_1, T_2, \dots, T_L$: the output of each matmul, each normalization, each activation, each attention block. The prover commits to this sequence as it streams, using a Merkle Mountain Range (MMR) over the per-tensor digests $d_i = \text{keccak}(\text{enc}(T_i))$, where `enc` is the canonical fixed-point encoding of the tensor. An MMR is appended in $O(\log i)$ work per leaf and maintains a frontier of $O(\log L)$ peak digests; for the layer counts of current models the frontier is a few hundred bytes (on the order of 768 bytes for L up to $\sim 2^{24}$), so the trace commitment is streamable in constant memory and the prover never has to retain the full trace. The committed root is the `activationTraceRoot` of Definition 11.1; the `modelWeightsRoot` is the analogous keccak Merkle root computed once, at load time, over the quantized weight blocks the prover holds.

We emphasize a fact about the present implementation, in the interest of stating only what is true: the engine today emits logits without an activation-trace commitment; the MMR of this subsection is the proposed proof-bearing addition, and its natural integration point is the synchronous inference precompile already present in the execution client. The commitment is cheap by design precisely so that adding it does not change the asymptotic cost of inference.

12.2 Freivalds verification over an exact integer accumulator

The dominant operation in every layer is a matrix product $C = AB$ with $A \in \mathbb{Z}^{m \times k}$, $B \in \mathbb{Z}^{k \times n}$. Re-executing it costs $O(mnk)$. Freivalds’ algorithm [7] verifies a claimed C with one matrix–vector product on each side: sample $r \in \mathbb{F}_p^n$ and accept iff

$$A(Br) = Cr \pmod{p}, \tag{20}$$

which costs $O(mn + nk) = O(n^2)$ field operations—asymptotically cheaper than the $O(n^3)$ product and far cheaper than re-running the layer. The random vector r is derived from the challenge, $r = \text{PRG}(\text{ch} \parallel \ell)$ for layer index ℓ , so by Lemma 10.2 it is unpredictable until the beacon block and cannot be pre-baked.

The decisive enabling fact is that the proof-bearing inference mode runs on an *exact integer* matmul. In the int8 quantization path, weights and activations are stored as signed 8-bit blocks and the per-block dot product accumulates in a 32-bit integer with *no intermediate rounding*; the floating-point scale is applied only *after* the integer accumulation closes. Concretely, for a 32-lane block the kernel computes $\sum_j x_j y_j$ with $x_j, y_j \in [-128, 127]$, so each product is at most $127^2 = 16129$ and a block sum is at most 516128, comfortably inside the $\mathbb{Z}_{2^{31}}$ range—hence the accumulation is provably exact. The same exact-integer dot is realized in portable scalar code and in hardware integer-dot instructions (NEON `vdotq_s32`, AVX2 `madd`, WASM `i32x4_dot`), all accumulating in 32-bit integer lanes before any conversion to float. Because the accumulator is exact integers, the products live in $\mathbb{Z}$, and we may run the Freivalds check (20) over a prime field $\mathbb{F}_p$ with $p = 2^{61} - 1$ (a Mersenne prime, so reduction is a shift and add) large enough that no honest matmul ever wraps. The consequence is that the check has *zero false rejections*: an honest prover passes with probability exactly 1, because integer equality implies equality mod p .

Remark 12.1 (Why the proof-bearing mode is quantized, not fp16). Floating-point matrix multiplication is *not* associative and *not* reproducible across hardware: the same AB computed with fp16 accumulation on two GPUs, or on a GPU versus a CPU, can differ in the low bits due to differing reduction orders, fused-multiply-add availability, and rounding. A Freivalds check over fp16 would therefore suffer false rejections on honest cross-hardware re-derivation, which is fatal for a permissionless verifier. The integer path removes this *provided the entire length- K dot product accumulates in i32 with no intermediate float*: under that condition $\text{i8} \times \text{i8} \rightarrow \text{i32}$ accumulation is associative and bit-identical on every architecture, so the proof-bearing mode is the quantized mode by necessity, not by preference. This is the software-only analogue of fixing a single hash function: a single, exactly reproducible arithmetic.

Remark 12.2 (The required whole- K -i32 kernel is a prerequisite, not the shipping kernel). The exactness above is a property of a *whole- K i32 accumulator*, and that is a kernel the engine does not yet ship. The quantized GEMM in the released engine (the `f8q8` block-quantized path) accumulates each 32-lane block’s i8 dot in i32 but then *converts the block partial to f32, scales by the block’s float d , and sums across blocks in floating point*—a per-block FMA chain whose reduction order is non-associative and not bit-identical across NEON, AVX, and WASM back ends. Exact Freivalds over that kernel would therefore *false-reject honest cross-hardware provers*, exactly the failure mode (a) of Section 16.1. The $\epsilon = 0$ guarantee of this section consequently holds for a pinned whole- K -i32 reduction (block scales pulled out and the integer accumulation closed over the full contraction before any float, a RepOps-style cross-block pin) that is *designed and to be built*, not for the current block-float kernel. We state this so the soundness claim is not read as a property of code that does not yet have it.

12.3 The non-matmul remainder

Softmax, RMS/layer normalization, GELU/SiLU activations, and the attention softmax are not matrix products and so are not Freivalds-checked. They are handled by *commit-and-recompute*: each is committed in the trace (Section 12.1), and a verifier challenging a step recomputes that op deterministically from the committed inputs (which are themselves Freivalds-checked matmul outputs) and compares to the committed result. These ops are $O(n^2)$ or cheaper and depend

element-wise on their inputs, so recomputing a challenged one is inexpensive. Under the determinism contract below they are bit-reproducible, so the comparison has, like Freivalds, zero false rejections.

12.4 The determinism contract

“Re-derive and compare” presupposes that two honest evaluations of the same model on the same input produce *byte-identical* traces. This is a property of the runtime, not of the mathematics, and we make it an explicit, software-only contract, digested into `runtimeMeasurement` (Section 10) and gated by governance. The contract has five clauses.

1. **Integer matmul.** Proof-bearing inference uses the `int8` (or lower) path with exact `i32` accumulation. This is what makes the `matmul` bit-exact cross-hardware (the remark above) and what makes Freivalds bit-exact.
2. **Temperature-0 argmax decoding.** Sampling is greedy: the next token is `arg max` over the logits, with no stochastic sampling on the proof-bearing path.
3. **A fixed, specified tie-break.** When two logits are exactly equal, the contract fixes which token wins—the *lowest token id*—so that `argmax` is a total function with no implementation freedom. (We note that the current engine’s `max_by` returns the *highest* equal index; the contract therefore specifies the rule normatively and requires the engine’s tie-break to be pinned to it. Exact ties are vanishingly rare in practice but must be specified for reproducibility, exactly as a hash must be specified to the bit.)
4. **Pinned kernels and tokenizer.** The kernel set and tokenizer are fixed and their identifiers are part of `runtimeMeasurement`, so two provers run literally the same code paths.
5. **A committed seed.** Any residual stochasticity (e.g. dropout, disabled on the inference path) is driven by a committed seed; on the greedy path no RNG is consulted, so the seed is inert and the decode is a deterministic function of the prompt and the weights.

The contract is entirely software: every clause is a configuration of the inference runtime, checkable by reading the model file and the runtime identifiers. No clause requires special hardware. This is the precise sense in which the execution proof is “software-only”: both producing it and challenging it need only the model and a CPU.

12.5 Soundness

We can now state and prove the central guarantee: a prover who did not perform the genuine `matmuls` cannot answer a beacon-derived Freivalds challenge except with negligible probability.

Theorem 12.1 (Freivalds soundness for one layer). *Let $A \in \mathbb{Z}^{m \times k}$, $B \in \mathbb{Z}^{k \times n}$ be the true operands of a layer (the committed weights and the committed input activations), and let $\hat{C}$ be the prover’s claimed output, committed in the activation trace. Let $r \in \mathbb{F}_p^n$ be sampled uniformly via `PRG(ch ||  $\ell$ )`, modeled as uniform and independent of $\hat{C}$ by Lemma 10.2. If $\hat{C} \neq AB$ as integer matrices and p exceeds the magnitude of every honest entry of AB , then*

$$\Pr_r[\hat{C}r \equiv A(Br) \pmod{p}] \leq \frac{1}{p}. \quad (21)$$

Proof. Work over $\mathbb{F}_p$; reduce all integer entries mod p . Let $D = \hat{C} - AB \pmod{p}$. Because p exceeds the magnitude of the honest product’s entries, $\hat{C} \neq AB$ over $\mathbb{Z}$ implies $D \neq 0$ over $\mathbb{F}_p$ (a genuine discrepancy cannot be an artifact of modular wraparound). The check $\hat{C}r = A(Br) = (AB)r$ holds iff $Dr = 0$. Since $D \neq 0$, it has some nonzero row D_i , and $Dr = 0$ requires $\langle D_i, r \rangle = 0$. Fix any coordinate j with $D_{ij} \neq 0$. Conditioning on all coordinates of r except r_j , the equation $\langle D_i, r \rangle = 0$ determines r_j uniquely as $r_j = -D_{ij}^{-1} \sum_{j' \neq j} D_{ij'} r_{j'}$ (the inverse exists since $\mathbb{F}_p$ is a field and $D_{ij} \neq 0$). As r_j is uniform over the p field elements and independent of the others, it hits that single value with probability $1/p$. Hence $\Pr[Dr = 0] \leq \Pr[\langle D_i, r \rangle = 0] = 1/p$. $\square$ $\square$

Two amplifications make the soundness error cryptographically negligible. First, using k_{rep} independent random vectors per layer multiplies the per-layer error to $p^{-k_{\text{rep}}}$ (equivalently, sampling r from a larger space). Second, the verifier checks a random subset of layers; a prover who cheats on even a single layer must survive that layer’s check.

Theorem 12.2 (Soundness of the execution proof). *Consider a forward (or backward) pass of L matmul layers committed in the activation trace, verified with $p = 2^{61} - 1$, k_{rep} Freivalds vectors per checked layer, and a uniformly random choice of layers to check. Suppose a prover commits an `activationTraceRoot` and `outputHash` that are not the genuine result of evaluating the committed model on the committed input under the determinism contract. Then:*

1. (Dense cheat.) *If the prover’s committed trace differs from the genuine trace in the matmul output of every checked layer, it passes all checks with probability at most $p^{-k_{\text{rep}}} \leq 2^{-61k_{\text{rep}}}$; for $k_{\text{rep}} = 2$ this is at most 2^{-122} .*
2. (Sparse cheat.) *If the prover cheats on only a ρ -fraction of layers and the verifier checks s layers uniformly at random, the cheat escapes all checks only if no cheated layer is sampled, which (for cheats confined to few layers) has probability $(1 - \rho)^s$; conditioned on a cheated layer being sampled it is caught except with probability $p^{-k_{\text{rep}}}$.*
3. (Determinism binding.) *A committed output hash inconsistent with the committed (and checked) final-layer logits is detected with certainty, because the `argmax/tie-break` of Section 12.4 is recomputed from the Freivalds-verified logits.*

Proof. Each part composes Theorem 12.1 with independence. For (1), the k_{rep} vectors on a checked discrepant layer are independent, so by Theorem 12.1 the layer passes with probability $\leq (1/p)^{k_{\text{rep}}}$; a dense cheat is discrepant on every checked layer, so the prover must win at least one such layer’s check, bounded by $p^{-k_{\text{rep}}}$. For (2), the verifier’s s sampled layers miss the cheated fraction with probability $(1 - \rho)^s$ (sampling without replacement only lowers this); if any cheated layer is sampled, part (1)’s bound applies to it. For (3), the final-layer pre-activation logits are a matmul output and are Freivalds-checked; the verifier recomputes the `argmax` under the pinned tie-break and compares to the committed token. Since both are deterministic functions of the verified logits, any inconsistency is detected with probability 1. $\square$ $\square$

12.6 From “probably” to “certainly”: interactive bisection

The sparse-cheat term $(1 - \rho)^s$ in Theorem 12.2 shrinks with more sampled layers but does not reach zero by spot-checking alone: a cheat confined to a single layer is missed with probability $(1 - 1/L)^s$. To convert this into *certainty* at logarithmic cost we add an interactive bisection in the style of optimistic rollups [10, 9]. The prover and a challenger disagree on the final output; they recursively bisect the trace—each round the prover commits the midpoint activation, and the parties recurse

into the half where they disagree—until they isolate a *single* matmul whose inputs both agree on but whose output they dispute. That one matmul is then either Freivalds-checked or re-executed on-chain (it is a single layer, hence cheap), and the dispute is resolved deterministically: exactly one party is shown wrong. The number of rounds is $O(\log L)$, so the on-chain footprint is logarithmic in the trace length, and the guarantee is upgraded from “caught with overwhelming probability” to “caught with certainty, after $O(\log L)$ interaction.”

Corollary 12.3 (Software-only verifiability, by construction). *In the construction as specified (with the trace emitted and the whole- K -i32 kernel of Remark 12.2 in place), producing the execution proof requires the model and the prover’s compute; checking it—whether by spot-check (Theorem 12.2) or by bisection to certainty—requires only the model file, a CPU to compute PRG, the $O(n^2)$ Freivalds matrix–vector products, and keccak for the trace. No trusted execution environment, GPU attestation, or special hardware is needed by either party. This is a property of the design; the two prerequisites it names (trace emission, pinned whole- K reduction) are the remaining build (Section 19).*

Corollary 12.3 is the property that makes the scheme egalitarian: a CPU-only validator challenges a GPU prover, and any user both generates and verifies a proof against the same network-shared difficulty (the soundness parameters p , k_{rep} , s , and the bond). The emission and the off-chain challenger of Section 14 are built and tested (`poi_transcript.rs`, the Go `crypto/poi` watcher, and the on-chain `ComputeWitnessLib`), so this is a live property for the proof-bearing path, not only an intended one—the contrast with Bitcoin’s ASIC-mediated work, drawn out in Section 13.

13 The Bitcoin Correspondence: Shared, Software-Only Difficulty

The companion Thinking-Chains framing establishes that Proof-of-AI is the “AI Bitcoin”: it keeps Nakamoto consensus’s trustless, permissionless issuance through verifiable work, but replaces useless hashing with useful cognition, and replaces difficulty with a heterogeneous-compute market. This section completes that correspondence at the level this paper is about—*verification*—and reconciles two notions of “difficulty” that the AI setting forces apart.

13.1 The execution proof is the recomputable hash (intended correspondence)

In Bitcoin, the load-bearing line of the whole design is that *any node re-hashes the header and checks the claimed solution in $O(1)$* : a stranger can verify the work without trusting the worker and without redoing it. Sections 10–12 *specify* the cognitive analogue. The activation-trace commitment is *designed to be* the value everyone can recompute; the Freivalds check is the cheap verification; and Theorem 12.2 is the guarantee that the value *cannot be produced without doing the work*. The difference is that Bitcoin’s recomputable value is the output of useless hashing, whereas ours is the side effect of an answer a paying agent demanded. Table 5 draws the mechanism-for-mechanism correspondence at the verification layer; its right column is the design Sections 10–12 specify, of which only the Freivalds-check row is implemented today (the recomputable value is not yet emitted; Section 2).

13.2 Two difficulties: price and verification

Bitcoin has *one* difficulty because it prices and verifies *one* homogeneous resource (SHA-256 hash-power): the same scalar tunes how hard it is to *find* a block and is implicit in how the network *checks*

it. Intelligence is not homogeneous—it is produced on GPUs, CPUs, unified-memory Apple silicon, and accelerators differing by orders of magnitude in memory, bandwidth, and throughput—so Proof-of-AI must split the single scalar into two orthogonal quantities, and keep both.

Difficulty-as-price (the HMM).

The cost of a thought is set by the Hamiltonian Market Maker, a physics-based AMM over a heterogeneous-compute reserve manifold (GPU-hours, VRAM, CPU, bandwidth, storage) that places each task on the cheapest capable hardware and quotes a continuous, capital-efficient price. This is the generalization of Bitcoin’s difficulty adjustment to a market: as compute joins, the price of cognition falls and re-stabilizes, exactly as Bitcoin’s difficulty rises to hold block time as hashpower joins. The HMM answers *how much should this work cost?*

Difficulty-as-soundness-bar (this paper).

The strength of the verification is set by the network-shared soundness parameters: the field size p , the number of Freivalds vectors k_{rep} and sampled layers s (Theorem 12.2), and the bond a prover escrows (Section 14). These are uniform across all provers—a CPU validator and a GPU farm face the same bar—and adjustable like Bitcoin’s difficulty: as the value at stake or the adversary’s resources grow, the network raises k_{rep} , s , or the bond. The soundness bar answers *how sure must we be the work was done?*

The two are genuinely orthogonal: pricing concerns *which heterogeneous resource* clears the demand and at what rate; verification concerns *whether the claimed computation occurred*, and is hardware-agnostic by Corollary 12.3. Conflating them is the error the deployed settlement engine of Section 9 makes implicitly—it prices and pays for agreement while verifying nothing about computation. Keeping them separate is what lets the protocol *price* work across the messy real world of heterogeneous compute while *verifying* it with a single, uniform, software-only bar.

Principle 13.1 (Pricing $\perp$ verification). *The HMM prices the work; the execution proof verifies it. The two difficulties are independent dials: the market sets what cognition costs across heterogeneous hardware, and the soundness bar sets how certain the network is that the cognition occurred. Bitcoin needed only one dial because its work was homogeneous and useless; Proof-of-AI needs both because its work is heterogeneous and useful.*

13.3 Software-only difficulty is egalitarian difficulty

Bitcoin’s difficulty, though uniform as a number, is *not* egalitarian in practice: producing the work economically requires ASICs, which centralizes mining among those who can fabricate or buy them. Proof-of-AI’s soundness bar is *designed to be* uniform *and* egalitarian, because the work it gates can be both produced and verified in software. By Corollary 12.3, a CPU-only participant *would* verify any proof and—on the optimistic-Freivalds and redundancy tiers (Section 14)—also *produce* proofs for models it can run, challenge proofs it suspects, and earn by catching fraud. This egalitarian property is contingent on two unbuilt pieces: the trace emission, and a permissionless off-chain challenger/re-executor that opens a suspect transcript and disputes it. *Neither exists today*, so the “earn by catching fraud” incentive—and with it the economic argument that a liar always has skin in the game—is a property of the designed system, not the current one. Once both are built, the “hash everyone can recompute” becomes a computation everyone can both perform at their scale and check at CPU cost—the property Bitcoin aspired to and ASIC economics eroded.

What the network mines, in the end, is useful cognition; what it hashes to is a shared, recomputable commitment; and what it adjusts as compute joins is two dials instead of one—price, by the HMM, and certainty, by the soundness bar of Theorem 12.2.

14 The Tier Matrix: One Gate, Migrating Backends

The execution proof of Section 12 is one of several proof systems a deployment can require, trading off trust assumptions, cost, and hardware. This section specifies the menu and the single gate that selects from it. The discipline is the one used for post-quantum migration elsewhere in the stack: the policy “what proof does a task at this tier require?” lives in *exactly one place*, every consumer reads it the same way, and the proof backends are interchangeable behind a common binding so that the menu can change without touching the binding or the callers.

14.1 The proof backends

A proof system is identified by a `proofType` and registered as a backend that attests to a `reportData` (Section 10). The deployed registry recognizes four:

$$\text{proofType} \in \{ \underbrace{1}_{\text{CC-TEE}}, \underbrace{2}_{\text{zkML}}, \underbrace{3}_{\text{optimistic-Freivalds}}, \underbrace{4}_{M\text{-of-}N \text{ TEE}} \}. \quad (22)$$

Each backend exposes the same interface—“does this witness attest to this `reportData`?”—and the verifier dispatches on `proofType`, failing *closed* on any unregistered type. The optimistic-Freivalds backend (type 3) is the software-only path of this paper: a prover escrows a bond and commits (`reportData`, `activationTraceRoot`, `modelWeightsRoot`); within a challenge window any party may dispute by supplying a Freivalds response (or, on dispute, by driving the bisection of Section 12); an unchallenged commitment finalizes and the bond is reclaimable, while a proven fraud slashes the bond to the challenger. The heavy check runs off-chain in the inference engine (a watcher re-executes the committed trace and opens any discrepant matmul); the chain runs the economic state machine the check’s verdict drives, and—when the dispute is submitted—re-runs the Freivalds check itself on-chain so the verdict is not a watcher’s say-so. Both halves are built and tested (Section 19): the off-chain re-executor and challenger (`poi_transcript.rs` in Rust, `crypto/poi` in Go) open a transcript and produce the dispute, and the on-chain state machine (`OptimisticEvidence.sol`, 18/18) records the bonded commitment, verifies the exhibited fraud via `ComputeWitnessLib.provesFraud`, and slashes only on a verified discrepancy. The optimistic backend’s economic soundness—a liar is strictly unprofitable once watched, an honest prover is never slashable—is therefore a live deterrent: a fabricated matmul is provable except with probability $\leq 2^{-122}$, and the bond pays the challenger that proves it.

14.2 The tiers

Tiers map a task’s value and trust requirement to a required proof strength.

The matrix is read top-down by trust and bottom-up by accessibility. T0 (zkML) assumes nothing beyond the soundness of the proof system but is limited to small models by proving cost. T1 (*M-of-N* heterogeneous confidential compute) tolerates the compromise of any single TEE vendor by requiring agreement across enclaves from *different* hardware roots—NVIDIA confidential computing, Intel TDX [16], AMD SEV-SNP [17]—so that a break of one vendor’s attestation does not forge a proof. T2 is the workhorse: a single-vendor TEE for those who have the hardware, *or* the optimistic-Freivalds path of this paper for those who do not, which assumes only that at least one honest party will challenge a fraud within the window and that the prover’s bond exceeds the gain from cheating. T3 falls back to the redundancy-and-plurality settlement of Section 9, retained only where liveness dominates and the value at stake is low enough that honest-plurality is acceptable—and, crucially, *even T3 can now require a sampled execution proof* from the plurality operators, closing Proposition 9.1 at that tier too.

14.3 The profile gate: one place

The mapping from tier to required proof is a single function in a single contract, the on-chain analogue of the `RefuseUnderStrictPQ` gate used for post-quantum policy. It exposes one read,

$$\text{requiredProofType}(\text{tier}) \rightarrow \text{proofType}, \quad (23)$$

returning 0 (permissive) for any tier governance has not explicitly opted in, and otherwise the mandatory `proofType`. The settlement path consults this single oracle and enforces the invariant directly:

No valid compute proof of the required type $\Rightarrow$ no mint, no settlement.

This is the constructive answer to Section 9: the verifier’s three-step gate—(1) the proof’s `reportData` equals the task binding, (2) the runtime is governance-accepted, (3) the registered backend attests—runs *before* any reward is paid, and fails closed at every step *on any tier that has been opted in*. Where the gate is engaged, reward is bound to proof, not to signature and not to plurality; where it is not (the default), settlement remains the authorship-plus-agreement rule of Section 9, and—until the execution backend’s emission and challenger are built—the only proof type a gate can actually enforce is an attestation envelope, not the activation-trace execution proof. Because the policy lives in one place, there is exactly one way to ask “what does this tier require?”, and changing a tier’s requirement is a single governed write rather than a scattered edit.

14.4 Migration without rebinding

The tiers are not static. The decisive structural property of this design is that *the binding does not depend on the proof system*: `reportData` (Section 10) commits the task, model, input, output, and runtime, and says nothing about how that commitment is believed. Consequently, as zkML proving cost falls and larger models become trustlessly provable, tasks migrate from T2/T3 up to T0 by a governance change to `requiredProofType` alone—*zero change to the binding, zero change to the receipts, zero change to the callers*. The same holds in reverse: a deployment that gains confidential-compute hardware can move tasks from software-only T2 to TEE-backed T1 without touching the proof artifacts that earlier tasks emitted. The proof backends are orthogonal slots behind a stable binding; the menu evolves, the interface does not. This is the property *designed* to let a permissionless network start egalitarian (software-only Freivalds for everyone, once that backend’s emission and challenger are built) and strengthen task-by-task as trustless proving matures, without ever stranding the proofs it has already minted against.

15 Ownership: Data, Models, and Weights as User-Owned Assets

*Status note. The cryptographic core of this section is now built. The confidentiality envelope—a single-shot prompt seal under **X-Wing** (the hybrid X25519 + ML-KEM-768 KEM, RFC 9180 HPKE base mode), so the prompt is ciphertext to everyone but the chosen operator—is implemented and tested (`luxfi/crypto/promptseal`, 6 tests passing), and the in-enclave custody binding (a TEE quote that binds the operator’s confidentiality key to an attested image) is implemented and tested (`luxfi/crypto/attestation`, 9 tests passing); both are detailed and proven in the emission companion [25]. The proof emission these compose with is also built (Section 19). What remains forward-looking is the product layer built on top—data/model/weight marketplaces, the full user-owned-asset NFT surface, and the plaintext-blind confidential-compute tier wired into every operator path; modes (2) and (3) below describe that intended envelope. Mode*

(1)—the user runs the model on their own hardware—is non-custodial unconditionally. We keep the design framing for the product surface and state plainly that the sealing and attestation primitives underneath it are deployed.

The execution proofs are intended to make a stronger claim possible than “the network computed correctly.” They are designed to make *sovereignty* auditable: a user *would* own their data, own the inference performed on it, and own the model weights they improve, and *would* choose *per request* how much trust to place in whoever runs the compute—without the choice ever weakening the guarantee that the work was genuine. The aim is that trust become a UX dial, not a security compromise. With the confidentiality envelope unbuilt, that decoupling holds today only in mode (1).

15.1 The trust spectrum

For any request, a user selects one of three execution modes, ordered by decreasing sovereignty and increasing convenience. Every mode emits the same execution proofs of Sections 11–12, so every mode is auditable; the modes differ only in *who holds the data in the clear* and *who the user must trust for liveness*.

(1) User-operated node — maximal sovereignty.

The user runs the inference on their own hardware. The data never leaves the device; the model and weights are local; the proof is emitted for the user’s own records or for third parties who later wish to verify a result the user publishes. Trust in any operator is zero because there is no operator. This is the floor of the spectrum and the reference against which the others are measured.

(2) DAO trustless nodes — verify, don’t trust (*unbuilt*).

The user *would* send work to permissionless DAO operators under confidential compute (T1/T2) so the operator *would be plaintext-blind*: processing ciphertext inside an enclave and never seeing the data in the clear, while the execution proof (PoI/PoT) plus the slashing bond would make a wrong or fabricated result detectable and punishable. *As deployed today, none of this confidentiality holds*: there is no enclave ciphertext path and the operator receives plaintext. In the intended mode the user trusts no particular operator’s honesty—only the proof system and, for the TEE tiers, the hardware root. This is the designed “verify, don’t trust” mode: sovereignty over correctness and confidentiality without running the hardware oneself.

(3) Selected trusted operators — convenience.

The user picks specific operators by reputation (via the identity-and-reputation layer), accepting a centralized-feeling, faster, cheaper path. Even here the proofs remain emittable, so the user retains an audit trail: a trusted operator that misbehaves can be caught after the fact by anyone who checks the proof, and loses reputation and future routing. Trust is placed deliberately and remains accountable.

The three modes are points on a continuum, not separate products, because they share the binding and the proof primitive. A user can move a workload from mode (3) to mode (2) to mode (1) as their trust requirements rise, with no change to the artifacts the workload produces.

15.2 Data confidentiality at rest: primitives built, operator integration pending

The cryptographic primitives are built: in modes (2) and (3) the prompt is sealed under X-Wing (X25519 + ML-KEM-768) HPKE to the chosen operator’s published key (`crypto/promptseal`,

tested), committed on-chain only by hash, and—in the confidential-compute tier—decryptable only inside an enclave whose attested image is bound to that key (`crypto/attestation`, tested; [25]). What is *not* yet built is the *integration* of these primitives into the production inference operator path: the seal/open is not yet wired into every operator’s request handler, nor the per-operator key-management lifecycle, so in the currently deployed operator path the prompt may be held in the clear and the “never plaintext at rest” property is guaranteed only in mode (1) and wherever the sealing path is enabled. The soundness binding is independent of this and is real: the execution proof binds to `inputHash` (Section 10), so a proof certifies that the *committed* input was processed while revealing nothing about the input itself. That binding is what makes PoC’s privacy clause (Definition 11.2) more than a promise: the contributor proves the gradient was computed over data consistent with the committed policy (the training-step graph is Freivalds-checked, Section 19) and shares only the resulting delta, so the corpus stays private while the contribution stays verifiable. The confidentiality envelope (X-Wing prompt seal + attestation) and the proof emission it relies on are both built and tested (`crypto/promptseal`, `crypto/attestation`, the engine `poi_*` modules; Section 19); the corpus-scale data-policy product around them is the forward-looking part.

15.3 Weights are a portable, user-owned asset

A model improvement—the ΔW of PoT/PoC—is serialized in the gym’s canonical backend-independent BF16 format and committed by its root (Section 11). Because the serialization is canonical and backend-independent, a delta trained on one accelerator is byte-portable to any other (*this part is real—it is the gym’s shipping format*), and its provenance (which base model, which data policy, what improvement margin) *would be* attested by the PoC it was minted with (*this part depends on the unbuilt proof emission*). The end state is that a user’s weights are an on-chain asset they own and can move, sell, or compose—with royalties flowing to the contributors whose proven improvements it incorporates—rather than an artifact locked inside a single provider’s stack. The proof framework is what gives the asset its value: a delta is worth more when it carries a verifiable proof that it was genuinely trained and genuinely improves the model, than when it is merely asserted to.

15.4 Trust as a dial, security as a floor

The synthesis the design aims at is that the execution proofs decouple *trust* from *security*. A user would dial trust up the spectrum—from running everything themselves, to trusting a proof system, to trusting a named operator—according to their convenience and threat model, and the proof framework would hold the security floor constant underneath: in every mode, a result that claims to be the model’s output on the user’s input either is, or is caught. This is the intended difference between the deployed settlement engine, where “trusting the operators” means accepting unverified guesses (Section 9), and a proof-backed network, where trusting an operator would be a choice about *speed and cost*, never about *correctness*. *This decoupling is not yet in force*: with neither the proof emission nor the confidentiality envelope built, mode (1) (self-hosting) is the only mode that today delivers either sovereignty or non-custody. The end state is that the user owns the data, the inference, and the weights, and the protocol’s job is only to make all three auditable, at whatever point on the trust spectrum the user selects; reaching it is the build of Section 19.

16 Full Modern-LLM Coverage: Tiering, Not Floating-Point Freivalds

The execution proof of Section 12 is exact because it runs over the int8 path’s integer accumulator. A natural objection is that “real” frontier models run in floating point (fp16, bf16, fp8), so an exact-integer verifier covers only the quantized deployment and not the full precision a model was trained in. This section answers the objection rigorously, and the answer is the crux of the design: we do *not* force Freivalds onto floating point, because naive floating-point Freivalds is simultaneously *unsound* and *unusable*. Instead, coverage of the full model space is *designed to be* achieved by *tiering* the verification mode to the arithmetic, so that every model has a proof-bearing mode while no mode pays a soundness or liveness penalty it cannot survive. The quantized, exact, software-only mode is then not a limitation but a *feature*.

Two status caveats, stated before the argument. (i) The exact-integer mode this section relies on requires a whole- K integer accumulator. This is *built and is the proof path*: `poi.rs::exact_matmul` is a whole- K i64 reduction with no cross-block float rescale, and `ProvableLinear` runs its matmuls there. The distinction is that the engine’s *fast* kernel (`f8q8`) accumulates across blocks in `f32`—non-associative and cross-hardware non-reproducible (Remark 12.2)—so it is used for ordinary inference, while a proof-bearing inference pays the exact accumulator’s throughput cost to get the “ $\epsilon = 0$, zero false-reject” property. Running exact Freivalds over the `f8q8` fast path would false-reject honest provers; running it over `exact_matmul` does not, which is why the proof path is the integer path. (ii) The quantized models are the arithmetic the zoo *deploys in*, and the families the test suite exercises (dense/SwiGLU, attention with integer softmax, MoE, training) now *emit a proof transcript*—so provable-today coverage is nonempty, not 0%. What remains *designed* is the breadth: the floating-point (RepOps) tier and the architecture families not yet wired to emit (Section 2).

16.1 Why naive floating-point Freivalds fails twice

Floating-point matrix multiplication breaks the Freivalds check in two independent ways, one fatal to *liveness* and one fatal to *soundness*. Both stem from the same root cause: an fp GEMM is not a function of its operands alone.

(a) Non-determinism $\Rightarrow$ false-reject (liveness break). Floating-point addition is not associative: $(a + b) + c \neq a + (b + c)$ in general under IEEE-754 rounding. A GEMM accumulates k products into each output entry, and the *order* of that accumulation is chosen by the kernel—tiling, the number of threads in a reduction, whether a fused multiply-add (FMA) is emitted, the warp/wavefront width, and the vendor’s library version all permute it. Two honest evaluations of the same A, B in fp16 on two different GPUs, or on a GPU versus a CPU, therefore yield outputs C and C' with $C \neq C'$ in the low mantissa bits. A Freivalds check demands the equality $Cr = A(Br)$ over the *produced* C ; with $C \neq C'$, an honest re-deriving verifier on different hardware computes $A(B'r) \neq Cr$ and *rejects honest work*. In a permissionless network where a CPU validator must be able to challenge a GPU prover (Corollary 12.3), this false-reject is not a tail event—it is the *common* case, because heterogeneous hardware is the premise. A verification rule that rejects honest provers is a liveness failure regardless of its soundness.

(b) Tolerance bands $\Rightarrow$ a perturbation corridor (soundness break). The naive “fix” for (a) is to accept when $\|Cr - A(Br)\|_\infty < \epsilon$ for some tolerance ϵ absorbing cross-hardware rounding. This trades the liveness failure for a soundness failure. The acceptance region is no longer the point $\{C = AB\}$ but a *band* of radius ϵ around it, and a cheating prover may roam anywhere inside that

band. Concretely, an adversary who wishes to claim a fabricated output $\hat{C}$ chooses any $\hat{C}$ with $\|(\hat{C} - AB)r\|_\infty < \epsilon$; the Freivalds projection r collapses the $m \times n$ discrepancy matrix $D = \hat{C} - AB$ to the m -vector Dr , so the adversary needs only Dr small, not D small—an under-determined constraint with a large solution space. The corridor is *exploitable*: it lets a prover skip or approximate work (e.g. run a smaller model, a lower-rank approximation, or a truncated accumulation) and land inside ϵ , harvesting the fee while the band hides the cheat. Worse, the corridor must be *widened* to cover legitimate cross-hardware drift, and every widening enlarges the cheating manifold. There is no setting of ϵ that admits all honest cross-hardware runs (requires ϵ large) while excluding all profitable cheats (requires ϵ small): the two requirements pull in opposite directions, exactly as in the analogous attack on tolerance-based optimistic ML verifiers. We formalize this dilemma in the companion proofs document (the floating-point soundness/liveness no-go).

Remark 16.1 (The integer path is not a downgrade—it removes both failures at once). The integer accumulator (Section 12) makes the matmul a function of its operands alone: $i8 \times i8 \rightarrow i32$ accumulation is associative and bit-identical on every architecture. This simultaneously kills (a)—two honest provers compute the *same* C , so $\epsilon = 0$ and there are zero false rejections—and (b)—with $\epsilon = 0$ the acceptance region collapses back to the exact point $\{C = AB\}$ and the corridor vanishes, restoring the clean $1/p$ soundness of Theorem 12.1. This is why the proof-bearing mode is quantized by *necessity*: exact arithmetic is the only arithmetic for which a software-only, cross-hardware, tolerance-free equality check exists.

16.2 The key fact: floating-point values are integers

The tiering does not abandon floating-point models; it gives them an exact mode too, via a fact that is easy to miss. An IEEE-754 [19] value is *already an integer encoding*: a sign bit, an exponent field, and a mantissa field, all integers. For fp16/bf16/fp8 the significand is a small integer (an 11-, 8-, or 4-bit mantissa plus the implicit leading one), and a product of two such values, *before rounding*, is an exact integer times a power of two. The IEEE rounding that follows is a *deterministic, total function* of the exact product and the rounding mode—it is not a source of nondeterminism. We make this precise.

Proposition 16.1 (Floating-point GEMM is exact integer-mantissa arithmetic up to a fixed reduction order). *Fix an IEEE-754 format F (e.g. bf16) with round-to-nearest-even, and fix a reduction tree τ : a binary tree on the k index positions specifying the exact order in which the k products of a dot product are summed and rounded. Then for any operands A, B in F , the GEMM output $\text{GEMM}_F^\tau(A, B)$ is a deterministic function of (A, B) alone, computable by integer arithmetic on the mantissa/exponent fields with an IEEE rounding step at each tree node, and identical on every machine that honors F and τ . The only freedom in a floating-point GEMM is the choice of τ ; once τ is pinned, the result is bit-exact and hardware-independent.*

Argument. Each product $a_i b_j$ of two F -values has a significand that is the integer product of the two integer significands and an exponent that is the integer sum of the two exponents; this is exact in $\mathbb{Z}$ (no rounding) because integer multiplication is exact. A tree node combines two already-computed partial sums u, v (each an F -value, hence an integer significand $\times 2^e$): align exponents (an integer shift), add significands (an integer add), then apply round-to-nearest-even (a total function of the exact integer sum and the format’s precision). Every operation is integer arithmetic plus a deterministic rounding decision, so the node output is a deterministic function of (u, v) . By induction over τ , $\text{GEMM}_F^\tau(A, B)$ is a deterministic function of (A, B, τ) . The associativity counterexample of Section 16.1 is precisely a disagreement in τ between two kernels; fixing τ removes it. $\square \quad \square$

Proposition 16.1 isolates the entire problem to one degree of freedom: *accumulation order*. It says nothing about precision loss—an fp16 GEMM still rounds, and rounding still discards low bits—but it says the rounding is *reproducible*: every honest prover that uses the pinned tree τ produces the same rounded bits. This is the hook on which an exact floating-point Freivalds hangs.

16.3 RepOps: pinning the reduction tree makes floating-point bit-exact

The deterministic-floating-point (“RepOps”, for *reproducible operators*) verification mode for the fp16/bf16/fp8 tier is then: *pin the reduction tree τ as part of runtimeMeasurement* (Section 10), so that the floating-point GEMM becomes bit-exact across hardware by Proposition 16.1, and run Freivalds over the resulting exact values. Because τ is fixed, $\epsilon = 0$: there is no cross-hardware drift to absorb (failure (a) is gone) and no tolerance band to exploit (failure (b) is gone). The Freivalds soundness of Theorem 12.1 applies verbatim, over the exact (rounded, but reproducibly rounded) products—formally, one runs the check over $\mathbb{F}_p$ on the integer significand–exponent representation of the pinned-order output, with p chosen to bound the dynamic range of the accumulated significands.

The cost is concrete and is the reason RepOps is a *tier*, not the default: pinning τ forfeits the vendor’s hand-tuned GEMM kernels, whose speed comes precisely from choosing τ freely (large tiles, many-thread reductions, FMA fusion, mixed precision). A reproducible kernel fixes the reduction order and so runs slower than the vendor optimum—the standard reproducible-summation trade-off [18]. The deployment therefore pays for fp-precision proofs in throughput; whether that is worth it is a per-task governance choice (Section 14.3), not a protocol constant.

16.4 The coverage tiers, by arithmetic

Putting the pieces together, every model has at least one proof-bearing mode, selected by its arithmetic and value. Table 7 states the menu. It refines, rather than replaces, the trust-tier matrix of Section 14: that table selects a backend by *trust*; this one selects an *exactness regime* by arithmetic. The two compose—e.g. a bf16 model at high value runs RepOps-Freivalds under an M -of- N bond, while the same model quantized runs exact-integer Freivalds.

The structure of Table 7 is the whole argument: full coverage is a *disjunction of exact modes*, not a single mode stretched past where it is sound. A model that ships quantized (the common deployment, and the entire ZOO zoo) is *designed to* get the strongest of the three properties—exact, software-only, hardware-free, tolerance-free—once the whole- K -i32 kernel and the trace emission are built. A model that must be served in full floating-point precision still has a proof: pay the RepOps throughput tax for a software-only exact check, or attest in a TEE for full-speed kernels at the cost of a hardware-trust assumption. No model is left without a proof-bearing mode, and no mode is asked to bear a soundness or liveness load it cannot carry.

Principle 16.1 (Exactness is tiered to arithmetic). *There is exactly one way to make a software-only, cross-hardware, tolerance-free matmul check: exact arithmetic. Quantized models supply it natively (integer accumulation); floating-point models supply it by pinning the reduction tree (RepOps) at a throughput cost, or sidestep it via hardware attestation. “Support every modern LLM” is therefore a routing decision over exact regimes—never a single floating-point Freivalds, which is unsound (corridor) and unusable (false-reject) at once.*

17 Coverage of the Model Zoo: Two Extensions Suffice

Section 11 establishes that the base primitive—Freivalds-checked matmuls plus deterministic re-computation of the cheap remainder—covers any workload whose dominant cost is dense matrix

multiplication. Modern model families, however, are not all plain stacks of GEMMs: some route tokens to a sparse subset of experts, some replace attention with a linear recurrence or state-space scan, some consume pixels or audio, and some denoise a latent over many steps. This section surveys the architecture families actually shipped in the ZOO model zoo and maps each to what the proof needs. The payoff is a structural claim that makes “support the entire zoo” tractable: *beyond the base matmul check, there are exactly two new core techniques—a routing proof and a sequential-chaining proof*—plus per-modality input and output commitments with deterministic pre- and post-processing. No third core technique is required.

We are scholarly about status. The base Freivalds verifier over the exact integer accumulator is *implemented and tested* (Section 19); the two determinism prerequisites the harder families lean on—the lowest-token-id argmax tie-break and the canonical backend-independent MoE top- k —are *also implemented and tested in the engine*; the activation-trace commitment, the forward-pass emission, and the *routing proof* are now *built and tested* as well (the routing proof is the TOP-KROUTE/GATHER path of `poi_graph.rs`, exercised by a real MoE block). What remains *designed here, not yet built* is the *sequential-chaining proof* (for state-space and diffusion latents) and the RepOps floating-point tier; we label each obligation accordingly.

What “coverage” means here, precisely. The coverage this section establishes is a structural claim—every family *reduces* to the base primitive plus at most one of two extensions—and its tested floor is now nonempty. The families with a committed, verified transcript today are the dense/SwiGLU forward, the attention spine (with a deterministic integer softmax), the mixture-of-experts block (routing proof built), and a training step: these emit and verify in the engine’s test suite. What remains a *designed* ceiling rather than a deployed floor is the breadth: the sequential-chaining families (recurrent/state-space and diffusion) await the chaining-proof emission, and the fp16/bf16/fp8 tier (RepOps, Section 16.3) awaits the pinned-reduction kernel (specified and proven sound in [24], no kernel yet). The theorem below should be read as “these three techniques *suffice* to cover the zoo, two of the three are built, and the families in the test suite are covered today.”

17.1 Dense transformer (base case)

The dense decoder/encoder—**zen-nano** and the dense tiers of the family—is the case the base primitive covers directly. Every layer is a sequence of GEMMs (the attention Q, K, V projections, the score-value product, the output projection, and the two MLP matmuls) interleaved with normalization and an activation. The matmuls are Freivalds-checked (Theorem 12.2); the elementwise remainder is committed and recomputed (Section 12). **Need:** base primitive only. **Status:** verifier core built and tested; trace commitment + forward hooks designed.

17.2 Mixture-of-Experts: the routing proof

MoE models [23]—**zen4-max** (35B-A3B), **zen-coder-flash** (GLM-style MoE), **zen-max** (DeepseekV3/Kimi-style, 384 experts, 8 active)—add a router that, per token, scores the experts and dispatches the token to its top- k . The expert FFNs are still GEMMs (base primitive covers them once the right experts are known), but the *routing* introduces two new obligations.

Routing determinism (built). A different selected expert is a different sub-network, so if a prover and a re-deriving verifier disagree on the top- k they compute different outputs and an honest run false-rejects. The engine therefore exposes a process-wide proof-deterministic routing mode (`set_moe_proof_deterministic`; Section 19) whose top- k uses a *canonical, backend-independent*

tie-break—order each row by (score descending, expert-index ascending) under an f32 total order—so every machine selects the identical experts. This is the routing analog of the lowest-id argmax tie-break, and it is implemented and tested.

Routing proof (designed). Determinism makes the selection reproducible; soundness additionally requires that the prover *committed* to the router’s behavior so a verifier can recheck it. The routing proof commits, per routed token, (i) the router logits (themselves a matmul output, hence Freivalds-checkable), (ii) the selected expert indices, and (iii) which tokens activated which experts (the sparse activation pattern). A verifier challenging a token recomputes the top- k from the committed (and Freivalds-verified) router logits under the canonical tie-break and checks it against the committed indices; it then checks that the committed expert-FFN matmuls used exactly those experts on exactly those tokens. The discrepancy a cheater would introduce—routing to a cheaper expert, or skipping the dispatch—is caught either at the top- k recheck (wrong indices) or at the expert-matmul Freivalds check (wrong sub-network). **Need:** routing determinism (built) + routing proof: commit router logits, selected indices, sparse activation pattern (designed). **This is the first of the two new core techniques.**

17.3 Linear / gated-recurrent / state-space: the sequential-chaining proof

The single most important new technique is for the families whose sequence mixer is *not* a single batched matmul. The ZOO zoo ships several: **zen4-nano/micro/mini/pro** use a gated-DeltaNet [21] linear-attention recurrence; **zen4-thunder/storm** use a lightning-attention linear recurrence; Mamba/SSM-style backbones [20] use a selective state-space scan. In all of these, the layer carries a *recurrent state* S_t updated step by step:

$$S_t = \Phi_t S_{t-1} + \Psi_t, \quad o_t = \Omega_t S_t, \quad (24)$$

where Φ_t, Ψ_t, Ω_t are (input-dependent) linear maps. Each *step* of (24) is matmuls and so is Freivalds-checkable in isolation. The new difficulty is the *chain*: a prover could compute each step honestly in isolation yet splice in a fabricated S_{t-1} at the boundary, so that step t is internally consistent but uses a state that was never produced by step $t - 1$. Per-step Freivalds alone does not catch a broken seam, because each step’s check sees only its own operands.

The chaining technique (designed). Bind *consecutive states* in the commitment. The prover commits each recurrent state S_t into the activation trace and, crucially, commits the *transition* so that step t ’s committed input state is constrained to equal step $t - 1$ ’s committed output state. Concretely, the trace orders the states $S_0, S_1, \dots, S_T$ and the per-step Freivalds check at step t takes its input state from the committed S_{t-1} leaf (not a freshly supplied one); a verifier that challenges step t checks both that $S_t = \Phi_t S_{t-1} + \Psi_t$ (the step’s matmuls, Freivalds) *and* that the S_{t-1} it used is the same leaf the trace commits as step $t - 1$ ’s output. The seam is thus pinned by the commitment, and a spliced state is detected as a mismatch between the leaf step t consumed and the leaf step $t - 1$ produced. **Need:** per-step Freivalds + binding consecutive recurrent states in the commitment (designed). **This is the second of the two new core techniques.**

17.4 Multimodal vision (ViT tower)

The vision-language models—**zen-v1**, **zen-omni** (vision path), **zen-designer** (Qwen3-VL family)—prepend a vision tower to the language model. The tower is a Vision Transformer: the patch embedding is a convolution, which is an *im2col*-then-GEMM, and every subsequent block is the

dense-transformer case. So the tower’s heavy arithmetic is *already covered* by the base Freivalds primitive once the convolution is expressed as `im2col→GEMM`. Two modality-specific obligations remain, both of the “commit-and-recompute deterministically” kind rather than new core techniques:

1. **Preprocessed-pixel commitment.** The `inputHash` must bind the *preprocessed* pixel tensor (resized, normalized, patchified), not the original image file, so a prover cannot answer an easier image than the one posed. The preprocessing (resize, normalize) must be a *deterministic*, pinned operation in `runtimeMeasurement`, exactly as the tokenizer is for text.
2. **Patch-embed as `im2col→GEMM`.** The convolution is committed as its GEMM form so the same Freivalds check applies; the `im2col` rearrangement is an index gather and is committed and recomputed.

Need: base primitive + deterministic image preprocessing + preprocessed-pixel commitment + conv-as-`im2col`-GEMM (commitment designed; the GEMM check itself is the built core). If the model also routes vision tokens through experts (MoE-VL), it additionally needs the routing proof of Section 17.2—no third technique.

17.5 Audio (conformer / speech)

The audio models—`zen-tts`, the `zen-omni` audio path—are conformer/speech stacks whose self-attention and feed-forward blocks are again GEMMs (base primitive). The modality-specific piece is the *front end*: raw audio is turned into mel-spectrogram features by a short-time Fourier transform and a mel filterbank. An FFT is not a matrix multiply in the Freivalds sense (it is a fixed, input-independent linear transform), so it is handled by *commit-and-recompute*: the feature extraction is deterministic signal processing with a pinned window, hop, and filterbank, committed in the trace and recomputed by a challenger from the committed input samples. The model’s audio output (waveform or codec tokens) is committed in `outputHash`. **Need:** base primitive + deterministic FFT→mel feature extraction (commit-and-recompute) + input-sample and audio-output commitments. No new core technique—the FFT is the audio analog of softmax in the non-matmul remainder (Section 12).

17.6 Diffusion / Transfusion: chaining again

The generative families—`zen5`’s 3D/video and `zen-designer`’s image generation, built on a Transfusion-style [22] autoregressive-plus-diffusion design—denoise a latent over many steps. Each denoising step is a forward pass through a transformer or U-Net, i.e. matmuls, so a *single* step is the base case. The multi-step loop poses the same seam problem as the recurrent family: the latent z_t at step t must be the one that step $t - 1$ actually produced, and a prover could splice. The resolution is the *same* sequential-chaining primitive of Section 17.3: commit each intermediate latent $z_0, \dots, z_T$, pin step t ’s input latent to step $t - 1$ ’s committed output latent, and Freivalds-check each step. Two diffusion-specific determinism clauses join `runtimeMeasurement`: the *noise schedule* and the *sampler/seed* must be pinned, so that the (otherwise stochastic) trajectory is a deterministic function of the prompt and the committed seed—the diffusion analog of temperature-0 decoding. **Need:** per-step Freivalds + deterministic noise-schedule/sampler/seed + multi-step latent chaining—and the chaining is *literally* the recurrent-family technique reused, which is the point.

17.7 The unifying claim

Table 8 summarizes the survey. Read down the “new core technique” column: it takes only two distinct values. Mixture-of-experts (including its multimodal variants) needs the *routing proof*; the linear-recurrent/SSM family *and* diffusion’s multi-step loop both need the *sequential-chaining proof*; everything else is the base matmul primitive plus per-modality input/output commitment and deterministic pre/post-processing. That the recurrent backbones and the diffusion sampler collapse onto the *same* chaining primitive—despite looking unrelated—is what caps the technique count at two and makes whole-zoo coverage a finite, enumerable engineering task rather than an open-ended one.

Theorem 17.1 (Two-extension coverage of the zoo, by reduction). *Let the verification surface consist of (i) the base primitive (Freivalds-checked matmuls plus deterministic commit-and-recompute of elementwise/transform ops), (ii) the routing proof of Section 17.2, and (iii) the sequential-chaining proof of Section 17.3, each composed with the per-modality input/output commitments and the determinism contract. Then every architecture family shipped in the ZOO zoo—dense, MoE, linear-recurrent/SSM, multimodal vision, audio, and diffusion/Transfusion—is reducible to this surface: its heavy arithmetic reduces to matmuls verified by (i), its sparse routing (if any) by (ii), and its inter-step state or latent seam (if any) by (iii). No additional core verification technique is required.* This is a statement about the design’s sufficiency; two of the three techniques—the base primitive (i) and the routing proof (ii)—are implemented and tested and are wired to emit from real forward/graph passes, so the reduction is realized today for the dense, attention, MoE, and training families, while the sequential-chaining technique (iii) and the long-tail families remain designed (Section 2).

Argument. By the per-family analyses above, composed with the orthogonality of Section 11. Each family’s dominant cost is dense matmul (the attention projections and MLPs of a transformer block, the expert FFNs of an MoE, the per-step maps of a recurrence (24), the per-step denoiser of a diffusion loop, the ViT tower including the im2col-expressed convolution, and the conformer blocks), all covered by (i). The two structural departures from a plain GEMM stack are sparse expert routing and an inter-step state/latent dependency. Routing is covered by (ii): the router logits are a matmul output checked by (i), the canonical top- k is recomputed and checked against the committed indices, and the expert matmuls are checked by (i) against exactly the committed experts. The inter-step seam—identical in form for the recurrent state S_t and the diffusion latent z_t —is covered by (iii): consecutive states are bound in the commitment so step t ’s consumed leaf must equal step $t - 1$ ’s produced leaf, and each step is checked by (i). The remaining per-modality operations (tokenization, image preprocessing, FFT/mel feature extraction, noise scheduling) are deterministic, pinned, and input-independent or elementwise, hence committed and recomputed under the determinism contract with zero false rejection, and the per-modality `inputHash/outputHash` bind the boundary data. As every family decomposes into these three cases plus pinned pre/post-processing, the three techniques suffice. $\square$ $\square$

The decomplexing discipline of Section 11 extends cleanly here. The routing proof is its own slot, orthogonal to the base primitive and to chaining; the chaining proof is its own slot, shared by two families that happen to need the same seam guarantee. Adding a future architecture means asking only “does it route?” and “does it carry inter-step state?”—two yes/no questions selecting from a fixed set of three techniques—rather than designing a new verifier. That is the sense in which whole-zoo coverage is *simple* and not merely *possible*.

18 Settlement to a Zoo Layer-3, and the Decentralized-Science Substrate

The execution proof of Sections 12–17 is a primitive: it makes one AI computation re-checkable. This section places that primitive in the layered architecture that *uses* it. ZOO Labs Foundation is the Layer-3 application and research layer—decentralized AI (DeAI) and decentralized science (DeSci), governed through `zips.zoo.ngo`—and it does not re-derive the proof system; it *consumes* the one canonical implementation. We describe (i) the layering and the single canonical contract surface, (ii) the on-chain directory by which a ZOO L3 is discovered and endorsed, (iii) how PoAI-verified work settles down from the L3 to the ZOO L2 and the LUX L1 without a trusted relay, and (iv) why this makes proven inference, embedding, and training the reproducibility substrate that decentralized science needs. Every mechanism below is a contract or module we cite by path; this section adds no new cryptography, only the composition.

18.1 The layering and the one canonical surface

The stack is four layers, each a value the one above settles into:



The Proof-of-AI contracts are not forked per layer. There is exactly one canonical implementation—the LUX standard library’s Proof-of-AI contracts (`luxfi/standard, contracts/ai/{governance,mining,token,co` licensed BSD-3-Eco)—and an L3 consumes it as a dependency rather than copying it, the same discipline by which the Go verifier (`luxfi/crypto/poi`), the Rust prover (`hanzo-engine`), and the Solidity witness (`ComputeWitnessLib.sol`) are three views of one byte-exact wire spec. “One and only one way” is the property that lets a proof emitted on an L3 verify, unchanged, when it settles on the L2 and finalizes on the L1: the leaf serialization, the keccak Merkle fold, the field $\mathbb{F}_p$ with $p = 2^{61} - 1$, and the $\text{BE}_{32}(j) \parallel \text{BE}_{64}(i)$ challenge derivation are identical at every layer (Section 12; [25]).

18.2 The on-chain L3 directory

A community that launches a Thinking-Chain L3 registers it in a hub directory deployed on the ZOO L2: `L3Registry.sol`. The registry is a pure directory by deliberate design—it “holds no L3 state and never reaches into an L3’s governance”—storing for each L3 the pointer that makes it findable: its name, its EVM `chainId`, the addresses of its `AIGovernor` and `AIObservatory`, the registrar, an off-chain metadata URI, and an endorsement flag. Two access tiers keep the directory open while curating the “official” set: *registration is permissionless* (anyone may list an L3, keyed one-record-per-`chainId` by `idOf(chainId) = keccak(“zoo/l3” || chainId)`), and *endorsement is gated to the ZOO meta-DAO* (the registry owner, `zooDAO`). Because registration is permissionless and one-per-`chainId`, a front-runner could squat a legitimate L3’s `chainId`; the DAO-gated `reassign` closes this—it overwrites a squatted record to the correct governor/observatory and resets endorsement to `false` so a reassignment never silently inherits a prior endorsement. The 8/8-passing test suite (`L3Registry.t.sol`) pins these invariants: permissionless register, one-per-chain, DAO-only endorse/reassign, and the squat-recovery path.

This is the ZOO flavor of decentralization stated precisely. The L3 chain itself is not created by the registry; it is created on the LUX P-Chain (`lux chain create -type=13`) and runs its own sovereign AI-governance stack (`AIGovernor`, `AIThoughtRegistry`, `AIObservatory`, `AIReputation`)

on its own ledger. The registry only records and, if the DAO chooses, endorses—so the directory is a public good and the curation is a separate, gated concern (orthogonality, again).

18.3 Settling PoAI-verified work down the stack, trustlessly

The substantive question is how work proven on an L3 becomes spendable value on the L1 without re-introducing a trusted bridge. Three contracts compose to answer it; each was read and tested for this paper.

The mint gate, fail-closed. On the settlement EVM, `AIMiner.mine` turns one verified inference receipt into a fair-launch `AICoin` subsidy. It is the gate the rest of the paper’s soundness flows through, and it fails closed: a completed receipt must (1) be the correct fixed-width shape carrying a non-zero output, (2) prove inclusion under a root the chain has accepted (`AIReceiptRoots`), and (3) pass a `ComputeProof` whose `reportData` reproduces the receipt’s own measured-model/prompt/output binding under a governance-accepted runtime—and “no verifier wired $\Rightarrow$ revert (`ProofGateUnavailable`)”. A forged receipt root alone no longer mints (the audit’s C3 hole), because the compute proof binding real (model,prompt,output) is mandatory. The subsidy is minted to the receipt’s *bound requester* (read from the attestation), never to the submitter, so a relayer cannot redirect rewards; and the amount is clamped to the vested halving allowance, so settlement never inflates past the schedule. The 12/12 `AIMiner.t.sol` and 19/19 `ComputeProofEnforcement.t.sol` suites exercise exactly these gates.

The trustless post-quantum relay. The receipt root that `AIMiner` requires SHOULD, on LUX, come from the A $\rightarrow$ C atomic boundary (the `aivmbridge` precompile); a foreign or downstream EVM sources it through `AChainRootOracle.sol`. Crucially this replaces a trusted single relayer with a *verified validator quorum*: a quorum of A-Chain validators signs (root,height) with **MLDSA** (the post-quantum scheme `luxfi/warp` carries as an `MLDSACertSet`), anyone may relay the signed root (permissionless, a public good), and the oracle verifies the quorum on-chain via the MLDSA precompile. “Forging an accepted root requires breaking MLDSA (post-quantum hard) *or* corrupting $\geq$ threshold of the validator set—the nation-state-resistant bar.” Settlement is therefore as post-quantum as the proof it carries: Freivalds soundness is information-theoretic (it holds against an unbounded, hence quantum, adversary; Section 12), the commitments are keccak (PQ-secure up to Grover), and now the cross-layer root relay is MLDSA. No layer of the settlement path rests on a classical-only assumption.

Global no-double-mint. Because the same proven computation could be presented on more than one chain, settlement reconciles through `AIComputeRegistry`: a miner claims a chain-independent computation key `keccak(modelSpec,promptHash,outputHash)` before minting, and a second claim for the same triple—on *any* LUX chain—reverts. One unit of proven work mints once, network-wide; the registry is wired set-once into the miner so the guarantee cannot later be rugged (6/6, `AIComputeRegistry.t.sol`). Together with the miner bond/slash registry (`MinerStakeRegistry.sol`, 8/8: a miner bonds capital in proportion to declared capacity, and the fraud-proof verifier slashes it on a proven discrepancy with a cooldown that is the fraud window), the settlement path has skin in the game at the bottom and a proof gate at the top.

The settlement chain, in one line. A model on the L3 emits a transcript (Section 12); a challenger can open and Freivalds-check any `matmul`, on-chain if disputed (`ComputeWitnessLib /`

`OptimisticEvidence`); the L3’s verified receipts roll up under an A-Chain root; that root crosses to the settlement EVM through an ML-DSA validator quorum (`AChainRootOracle`); and `AIMiner` mints—once, to the bound requester, clamped to the halving allowance, and only against a valid compute proof. The ZOO L3 is discoverable and (optionally) endorsed through `L3Registry`, so `zoo.vote` and the meta-DAO can see and govern the whole set from one place.

18.4 Why this is the substrate decentralized science needs

Decentralized science fails on the same rock as decentralized compute: a result that cannot be independently re-derived is a claim, not a finding. The execution proof converts the three operations DeSci runs on AI—inference, embedding, and training—into re-derivable objects, because all three are computation *graphs* and the proof binds graphs, not a special case per task (Section 11). Concretely:

- **A proven inference** (a model’s prediction on a scientific input) is a `GraphTrace` whose root a reviewer opens to Freivalds-check the heavy matmuls and recompute the deterministic glue—so a published prediction carries a transcript root, and “the model said *X*” is checkable rather than trusted.
- **A proven embedding** (a representation submitted to a shared scientific index) is the same forward graph; its provenance is the input commitment, which is exactly the data-sharing-orthogonal binding of Section 10—the proof says *this representation came from this input under this model*, independent of who may read the input.
- **A proven training step** is a longer graph of the *same* ops plus a transpose: forward $Y = X \cdot W$, backward $dW = X^\top dY$ and $dX = dY \cdot W^\top$, update $W' = W - \eta dW$, every matmul Freivalds-checked and a fabricated gradient caught exactly as a faked inference output is (`poi_graph.rs`, `test_training_step_is_just_a_graph`). A federated or bounty-funded training contribution is therefore auditable: a contributor who claims a gradient it did not compute is slashable on the same gate as a lying inference provider.

This is what BELUGA prices. Its L3 economy pays participants to train models, curate datasets, host inference, and evaluate outputs, all denominated in the BLG gas token; the execution proof is what makes each paid action *verifiable* rather than merely attested. The economic dial (the Hamiltonian Market Maker that prices heterogeneous compute) and the verification dial (the execution proof) stay orthogonal—the price of a unit of compute and the proof that the unit was performed are two knobs, not one (Section 13)—and the registry, the relay, and the mint gate are the plumbing that carries a verified BLG-priced contribution down to a capped, halving, fair-mined L1 coin. The result is a decentralized-science network whose ledger is not a record of assertions but a record of *re-checkable computations*: the property that distinguishes science from testimony, made native to the chain.

19 Implementation Status

We state plainly what is built and tested versus what is designed, so the paper’s claims can be checked against code. Relative to this paper’s earlier draft the boundary has moved: the activation-trace commitment, the proof-bearing forward, the full-graph composition, the off-chain challenger, and the on-chain enforcement—all marked “designed, not built” before—are now implemented and tested. What remains designed is the RepOps floating-point tier and emission of the architecture

families the test suite does not yet exercise. The test counts below are the actual passing counts of the named suites at the time of writing; each was run for this section.

19.1 Built and tested: the Freivalds core

The matmul verifier of Section 12 is implemented as a backend-agnostic, dependency-light module (`hanzo-engine/src/poi.rs`, whose only external dependency is `sha3`). It realizes the construction exactly as specified: arithmetic over the Mersenne field $\mathbb{F}_p$ with $p = 2^{61} - 1$, run on the engine’s *exact integer accumulator* (`exact_matmul`, a whole- K i64 reduction with no cross-block $f32$ rescale), with the Freivalds identity $A(Br) = Cr$ evaluated by left/right matrix-vector products (`freivalds_verify`) and deterministic, beacon-bound challenge derivation. Being dependency-free of the forward pass, its soundness is testable in isolation. The `poi.rs` suite is **7/7 passing**, and each test pins a claim from Section 12:

- `test_honest_matmul_passes` — an honest $C = AB$ passes every challenge vector (zero false rejection; the exact-integer guarantee).
- `test_tampered_output_is_caught` — flipping a single output entry ($50 \rightarrow 51$) is caught (Theorem 12.1 soundness).
- `test_larger_random_matmul_catches_sparse_cheat` — in a $16 \times 24 \cdot 24 \times 16$ product, a single buried off-by-one is caught; the *sparse-cheat* case that motivates challenging every layer.
- `test_negative_entries_field_reduction` — signed int8 operands reduce correctly into $[0, p)$ (Euclidean remainder), and a tampered negative-entry product is caught.
- `test_challenge_derivation_deterministic` — prover and verifier derive the *identical* challenge from the same seed; different seeds give different challenges (the reproducibility half of anti-replay, Lemma 10.2).
- `test_wrong_dims_rejected` — a shape mismatch fails *closed* rather than reading out of bounds.
- `test_derive_challenges_keccak_golden_parity` — the keccak challenge derivation matches the Go verifier *byte-for-byte* against a pinned golden hex vector; if either side’s keccak or $\text{BE}_{32}(j) \parallel \text{BE}_{64}(i)$ encoding drifts, this test (and its Go twin `TestDeriveChallengesGolden`) breaks. This is what guarantees a Rust-emitted proof verifies against the on-chain Go verifier unchanged.

19.2 Built and tested: the transcript, the proof-bearing forward, the graph

The wire that the earlier draft marked “designed” is implemented as three further modules, all in `hanzo-engine/src/`, all tested.

Activation-trace transcript and challenger (`poi_transcript.rs`, 8/8). A forward pass commits each proof-bearing matmul as a keccak Merkle leaf over its exact-integer (A, B, C) (`matmul_leaf`, domain tag “hanzo/poi/matmul-leaf/v1”); the root is the `transcriptRoot` the prover posts. The challenger derives a beacon-unpredictable matmul *index* ($\text{keccak}(\text{beacon} \parallel \text{root}) \bmod L$), has the prover open it (reveal (A, B, C) plus a Merkle inclusion proof), checks the operands were the committed ones, and Freivalds-checks $C = A \cdot B$ with k beacon-bound vectors (`verify_opening`). The Merkle fold (odd levels duplicate the last node; index fold $\text{even} \rightarrow \text{keccak}(\text{node} \parallel \text{sib})$, $\text{odd} \rightarrow \text{keccak}(\text{sib} \parallel \text{node})$) is identical to `AIMiner._verifyMerkle`, so a proof built here verifies on-chain.

The suite proves honest passes verify under every beacon, a fabricated C is caught, a swapped-in C fails Merkle inclusion, and the root binds every operand.

Sound emission (poi_forward.rs, 3/3). The audit’s CC-01 was that the value a layer *returns* must be the value it *commits*. `ProvableLinear` makes this hold by construction: it quantizes x per-row to int8, forms $C = \bar{x} \cdot W^\top$ on the exact i64 accumulator, commits $(\bar{x}, W^\top, C)$, and returns `dequant(C)`—so a challenger’s Freivalds over the committed operands re-checks the *same* computation that produced the output, not an $f32$ shadow that disagrees with it. `test_provable_ffn_emits_verifiable_t` runs a real SwiGLU FFN (*gate/up/down*) under emission and verifies every committed matmul; `test_output_is_the_committed_value` confirms the returned y equals `dequant(committed C)` to within float tolerance (the output-binding property, proved in the emission companion [25]).

Full-graph composition (poi_graph.rs, 9/9). A prover could commit correct matmuls computed on *unrelated* inputs; `GraphTrace` closes this with a typed computation trace over hash-named activations. Each op commits (*kind* || input commitments || *params* || output commitment); *chaining* (every input is a prior output or a declared input) is checked from the leaves alone, and *local correctness* re-checks each op—MATMUL by Freivalds, the deterministic glue (ADD, RELU, SCALE, TOPKROUTE under the canonical lowest-index tie-break, GATHER, SOFTMAX as a pure fixed-point integer reduction, TRANSPOSE) by direct recomputation. The suite emits and verifies a real MoE block (router $\rightarrow$ top- k $\rightarrow$ gather $\rightarrow$ expert up/down matmuls $\rightarrow$ residual), the attention spine ($QK^\top \rightarrow$ scale $\rightarrow$ integer softmax $\rightarrow \cdot V$), and—with the one TRANSPOSE op a backward pass adds—a full SGD training step, catching a fabricated gradient by Freivalds exactly as a faked inference output (`test_training_step_is_just_a_graph`). This is the code behind the canon “inference, embedding, and training are graphs, not three proof systems.”

Adversarial red-team battery (tests/poi_adversarial.rs, 11/11). Each test plays the attacker. The eight attacks—fabricate the final output, fabricate a hidden intermediate matmul, a sparse deep cheat, swap the output at open time, tamper an input at open time, pre-fit C to a guessed challenge, a griever’s fabricated opening, and a quantization lie ($y \neq \text{dequant}(C)$)—are each *caught*; the three defenses (two honest machines emit the byte-identical root, an honest forward verifies under every beacon, a shape mismatch fails closed) confirm zero false rejection. The pre-fit test is the load-bearing one for the beacon: a forgery tuned to a guessed r_1 passes r_1 and fails a fresh beacon-derived r_2 .

Jointly the engine carries **27 unit tests** (`poi.rs` 7, `poi_transcript.rs` 8, `poi_forward.rs` 3, `poi_graph.rs` 9) **plus 11 adversarial tests = 38 passing.**

19.3 Built and tested: the two determinism prerequisites

The exact-integer Freivalds check is bit-exact only if the surrounding decode is reproducible; Section 12’s determinism contract names the two places a naive runtime injects nondeterminism, and both are closed in the engine. The *lowest-token-id argmax tie-break* (`hanzo-engine/src/sampler.rs`) breaks a greedy-decoding logit tie to the lowest token id via a strict “>” scan, correcting Rust’s `Iterator::max_by` (which returns the *last* maximal element); the *canonical MoE routing top-k* (`hanzo-engine/src/ops.rs`, and its in-graph twin `top_k_indices` in `poi_graph.rs`) orders each row by score descending then expert-index ascending under a total order, so prover and verifier select identical experts (a different expert is a different sub-network). These are the routing analog of the argmax tie-break and are what make the “re-derive and compare” premise actual rather than aspirational for the greedy and MoE-routing cases.

19.4 Built and tested: the canonical Go verifier the chain consumes

The verifier the chain consumes is a dependency-light Go package living beside the canonical Keccak256 (luxfi/crypto/poi): `Verify/VerifyMulti` over $\mathbb{F}_p$ ($p = 2^{61} - 1$) on the exact integer accumulator, the transcript and `VerifyOpening`, the wire codec, and keccak `DeriveChallenges`. The engine’s `poi.rs` is the prover-side mirror; the two agree byte-for-byte, enforced by the shared golden challenge vector (`TestDeriveChallengesGolden`, the Go twin of the Rust parity test above). The suite is **23/23 passing**, including the honest-opening, fabricated-output-caught, swapped-reveal-caught, Merkle-index-fold, sparse-cheat, encode/decode round-trip, wire-fuzz, oversized-dims-rejected, griever-cannot-frame, and pre-fit-challenge cases—the Go-side mirror of the Rust adversarial battery, plus a verification-throughput scaling test.

19.5 Built and tested: on-chain enforcement and the fail-closed mint gate

The witness is verified in the EVM, not asserted by a watcher. The settlement contracts live in `luxfi/standard` (`contracts/ai/{governance,mining,token,compute,relay,registry}/`); the eight Proof-of-AI suites total **87 tests passing (forge test)**, distributed as follows.

- **ComputeWitnessLib.sol (5/5)**. Freivalds over $\mathbb{F}_p$ via `mulmod/addmod`, the Merkle fold, and the beacon challenge, byte-identical to the prover (same domain tag, same (rows, cols, data) serialization, same fold, same derivation). `verifyOpening` accepts a genuine opening; `provesFraud` is its complement—inclusion $\wedge \neg \text{Freivalds}$ —the predicate a challenger exhibits to slash. Opened dimensions are capped (`MAX_DIM`) so a griever cannot submit an unbounded opening.
- **OptimisticEvidence.sol (18/18)**. The commit/challenge/slash state machine: `submit` escrows a bond and opens a window with the commit block recorded as the challenge beacon (`blockhash(commitBlock)`)—fixed after the prover signed, so it cannot pre-fit a dodging C); `challenge` re-checks inclusion + Freivalds on-chain via `provesFraud` and, only on a verified discrepancy, marks the commitment `FRAUDULENT` and pays the bond to the challenger; `reclaim` returns the bond after an un-challenged window. The state machine is one-shot per `reportData` (a slashed liar cannot re-post), a zero bond is rejected (a liar always has skin in the game), and effects precede the single pull-of-funds (no reentrancy re-claim).
- **AIMiner.sol (12/12)**. The mint gate, *fail-closed*: `mine` requires receipt shape + finality + Merkle inclusion under an accepted root + a `ComputeProof` binding the receipt’s own measured model/prompt/output, and reverts (`ProofGateUnavailable`) if no verifier is wired—so a forged root alone cannot mint. The subsidy goes to the bound requester, is clamped to the vested halving allowance, and is one-shot per intent.
- **ComputeProofEnforcement.t.sol (19/19)**. The end-to-end enforcement that a proofless or mis-bound mint is rejected and a correctly-bound one succeeds, across the proof-type profiles.
- **AICoin.sol (11/11)**. The capped, halving, fair-mined coin: `MAX_SUBSIDY = 109 AI`, `HALVING_PERIOD = 4` years, the slope halving each epoch and the cumulative mint structurally bounded by `allowed(t) ≤ MAX`. Multiple verified-cognition mint paths share *one* halving schedule, so adding a mint path never inflates past the cap.
- **AIComputeRegistry.t.sol (6/6)**. The global no-double-mint ledger: a chain-independent computation key claimed before minting, so the same (model, prompt, output) mints once across all LUX chains.

- **MinerStakeRegistry.sol (8/8).** A miner bonds capital in proportion to declared capacity; the fraud-proof verifier slashes on a proven discrepancy with a cooldown that is the fraud window.
- **L3Registry.sol (8/8).** The ZOO L3 directory of Section 18: permissionless register, one-record-per-chainId, DAO-gated endorse/reassign, and squat recovery.

The post-quantum cross-layer root relay **AChainRootOracle.sol** (an ML-DSA validator quorum replacing a trusted single relay; Section 18) and the leaderless operator-quorum governance (**AIGovernor**, **AIParams**, with **AIExecute/AIApproval**) are likewise deployed in the same standard library and carry their own passing suites; we cite them for completeness as the operator-safe, post-quantum surface the settlement path composes with, without re-tabulating their tests here.

19.6 Designed, not yet built

The honest remaining work is now a named, smaller set; none of the following is implemented at the time of writing.

1. **RepOps (pinned reduction tree).** The deterministic-floating-point mode for the fp16/bf16/fp8 tier (Section 16.3) is designed and proven sound ([24]); the quantized integer path is what is built and is what the ZOO zoo ships. Exact Freivalds over the current *f8q8* fast path (which sums int8 block dots in *f32* across blocks) would false-reject honest cross-hardware provers; the proof path instead uses **exact_matmul**'s whole-*K* *i64* reduction, and the floating-point tier awaits the RepOps cross-block pinning.
2. **The sequential-chaining proof.** Binding consecutive recurrent states (and, identically, consecutive diffusion latents) in the commitment (Sections 17.3, 17.6) is designed and proven sound modulo its discharge obligation ([24]). The *forward*/TRANSPPOSE graph that covers dense, attention, MoE, and training is built (above); the per-step state seam for state-space and diffusion models is the remaining emission work.
3. **Emission of the long-tail architectures.** The families the test suite exercises (dense/SwiGLU, attention with integer softmax, MoE, training) emit and verify; the remaining model-zoo families inherit the same proof obligations ([24], Theorem on whole-zoo coverage) but are not yet wired to emit a transcript. This is breadth of integration, not a missing technique: no third core verification method is required.

Remark 19.1 (What the tested system does and does not establish). The 38 Rust, 23 Go, and 87 Solidity tests jointly establish that the *mathematical heart*—"you cannot claim the matmul without doing it"—is implemented correctly over the exact integer accumulator, that it *emits* from a real forward/graph and *composes* across attention, MoE, and a training step, that the challenger and the on-chain witness re-check it byte-identically in two more languages, and that the mint gate fails closed so a proofless mint reverts. What they do *not* yet establish is the floating-point tier (RepOps) and emission of the architecture families outside the tested set. We keep the boundary explicit, in keeping with the decompacting discipline of the rest of the paper: the verifier is a value (a pure function of operands and a challenge), the emission supplies those operands from a real pass, and the on-chain gate consumes the result—three separable concerns, each now built where this paper's earlier draft had only the first.

20 Related Work

Proof of Useful Work: Primecoin [1], Gridcoin [2] pioneered useful PoW but with limited domains.

AI Compute Markets: Golem, Render Network focus on compute markets without consensus integration.

TEE-Based Consensus: Oasis Network, Secret Network use TEEs for privacy but not AI verification.

Federated Learning: FL protocols [3] aggregate gradients but lack blockchain integration.

21 Future Work

The honest remaining work is the residual named throughout (Sections 2, 19), plus the orthogonal extensions the attestation and pricing layers invite. The execution-proof core, its emission, the cross-chain root relay, and the on-chain gate are built; what follows is breadth and the surrounding product surface.

1. **RepOps floating-point tier:** a pinned-reduction-tree kernel that makes fp16/bf16/fp8 GEMMs bit-exact, extending the proof path beyond the quantized integer tier (designed and proven sound in [24]; no kernel yet).
2. **Sequential-chaining emission:** wiring the proven chaining obligation for state-space (recurrent) and diffusion models, the one structural case the forward/transpose graph does not yet cover.
3. **Long-tail architecture emission:** extending transcript emission to the model-zoo families outside the tested set (multimodal towers, audio conformers), each of which already reduces to the existing techniques (Theorem 17.1).
4. **GPU TEE integration:** NVIDIA H100 / Blackwell confidential-computing quote parsing into the canonical attestation quote, for the hardware-trust tier that complements the software-only proof.
5. **Operator-path confidentiality integration:** wiring the built X-Wing prompt seal and attestation binding (`crypto/promptseal`, `crypto/attestation`) into every inference operator’s request handler and key-management lifecycle.
6. **Privacy-preserving quality scoring:** zero-knowledge proofs for the multi-party evaluation protocol, so a quality score reveals nothing about the evaluator’s private rubric.

22 Conclusion

Proof of AI replaces wasteful hash puzzles with a proof that genuine AI computation was performed—an activation-trace commitment that cannot be produced without doing the matrix multiplications, checked by Freivalds’ algorithm over an exact integer accumulator, software-only, with no trusted execution environment required to verify. The soundness is information-theoretic and therefore post-quantum: a fabricated matmul is caught except with probability $\leq 2^{-122}$ at two challenge vectors, and an honest prover is never false-rejected, because the check runs over a bit-exact integer reduction.

This paper’s contribution is not only the primitive but the *built and tested system* around it. The verifier emits from a real proof-bearing forward (**ProvableLinear**) and composes across a typed graph (**GraphTrace**) that binds attention, mixture-of-experts, and—with one transpose op—a training step, establishing the canon that inference, embedding, and training are graphs verified by one proof system, not three. The Go verifier the chain consumes is byte-identical to the prover; the on-chain enforcement runs Freivalds and the Merkle fold in pure Solidity, slashes a bond only on a verified fraud, and fails closed so a forged receipt root alone cannot mint. Across three languages the suite is 38 Rust, 23 Go, and 87 Solidity tests passing, with a cross-language golden vector that breaks if either side’s commitment drifts.

The execution proof is the substrate a ZOO Layer-3 settles. BELUGA and its peers register in an on-chain directory (**L3Registry**) and roll proven work up through an ML-DSA validator-quorum root relay (post-quantum, leaderless) to a capped, halving, fair-mined coin reconciled once network-wide—so a decentralized-science result is a re-checkable transcript, not an assertion. The economic dial that prices heterogeneous compute and the verification dial that proves it was performed stay orthogonal: the price of a unit of work and the proof that the work happened are two knobs, not one.

What remains honestly designed is a named, smaller residual: the RepOps pinned-reduction floating-point tier (the quantized integer tier is what ships and is what the tested floor covers), and emission of the architecture families—state-space and diffusion chaining, and the long tail of the model zoo—that the test suite does not yet exercise. No third core verification technique is required; the remaining work is breadth of integration over a settled mathematical core. POAI is the economic foundation for the ZOO Network: decentralized AI training and inference whose ledger records re-checkable computations, the property that distinguishes science from testimony.

23 Attestation Format

Listing 1: TEE Attestation Structure

```
{
  "version": 1,
  "tee_type": "sgx" | "sev" | "tdx",
  "measurement": {
    "mrenclave": "0x...",
    "mrsigner": "0x...",
    "product_id": 1,
    "security_version": 1
  },
  "report_data": {
    "task_hash": "0x...",
    "input_hash": "0x...",
    "output_hash": "0x...",
    "timestamp": 1650000000,
    "nonce": "0x..."
  },
  "signature": {
    "type": "ecdsa_p256",
    "value": "0x..."
  },
}
```

```

    "certificate_chain": ["0x...", "0x..."]
}

```

24 Quality Scoring Functions

Listing 2: Quality Scoring Implementation

```

def score_inference(attestation, samples):
    correct = sum(1 for s in samples
                  if verify_output(s))
    accuracy = correct / len(samples)
    latency_penalty = max(0,
                          (attestation.latency - TARGET) / TARGET)
    return accuracy * (1 - 0.1 * latency_penalty)

def score_training(attestation, val_before, val_after):
    improvement = val_after - val_before
    cost = attestation.flops / 1e15
    return improvement * math.exp(-0.1 * cost)

def score_extraction(experience, library):
    novelty = 1 - max_similarity(experience, library)
    applicability = test_on_benchmarks(experience)
    correctness = verify_experience(experience)
    return novelty * applicability * correctness

```

25 Slashing Conditions

1. **Invalid Attestation:** Signature does not verify
2. **Output Mismatch:** Re-execution produces different result
3. **Quality Manipulation:** Score differs from consensus by $> 3\sigma$
4. **Double Signing:** Same attestation submitted to multiple blocks
5. **Unavailability:** Validator offline for > 24 hours

Slashing amounts:

- Invalid attestation: 100% of bond
- Output mismatch: 50% of bond
- Quality manipulation: 25% of bond
- Double signing: 100% of bond
- Unavailability: 1% of bond per day

Acknowledgments

This work was supported by Zoo Labs Foundation. We thank Intel for SGX documentation and AMD for SEV-SNP technical support.

References

- [1] S. King, “Primecoin: Cryptocurrency with Prime Number Proof-of-Work,” 2013.
- [2] Gridcoin, “Rewarding BOINC Contributions,” Whitepaper, 2014.
- [3] B. McMahan et al., “Communication-Efficient Learning of Deep Networks from Decentralized Data,” *AISTATS*, 2017.
- [4] D. Boneh et al., “Short Signatures from the Weil Pairing,” *ASIACRYPT*, 2001.
- [5] V. Costan and S. Devadas, “Intel SGX Explained,” *Cryptology ePrint*, 2016.
- [6] E. Buchman et al., “The latest gossip on BFT consensus,” *arXiv:1807.04938*, 2018.
- [7] R. Freivalds, “Fast probabilistic algorithms,” in *Mathematical Foundations of Computer Science (MFCS)*, LNCS 74, pp. 57–69, 1979.
- [8] KD. Conway, C. So, X. Yu, and K. Wong, “opML: Optimistic Machine Learning on Blockchain,” *arXiv:2401.17555*, 2024.
- [9] J. Harrison et al. (Gensyn), “Verde: Verification via Refereed Delegation for Machine Learning Programs,” *arXiv:2502.19405*, 2025.
- [10] H. Kalodner et al., “Arbitrum: Scalable, private smart contracts,” *USENIX Security*, 2018.
- [11] Hyperbolic Labs, “Proof of Sampling (PoSP): A Verification Protocol for Decentralized AI Inference,” Technical Report, 2024.
- [12] S. Goldwasser, Y. T. Kalai, and G. N. Rothblum, “Delegating Computation: Interactive Proofs for Muggles,” *STOC*, 2008.
- [13] C. Lund, L. Fortnow, H. Karloff, and N. Nisan, “Algebraic Methods for Interactive Proof Systems,” *Journal of the ACM*, 39(4):859–868, 1992.
- [14] J. Chen et al., “ezkl: A System for Zero-Knowledge Machine Learning,” <https://github.com/zkonduit/ezkl>, 2024.
- [15] NVIDIA Corporation, “NVIDIA Confidential Computing on Hopper H100 / Blackwell: Architecture and Attestation,” Technical Whitepaper, 2023–2024.
- [16] Intel Corporation, “Intel Trust Domain Extensions (Intel TDX) Whitepaper,” 2023.
- [17] AMD, “AMD SEV-SNP: Strengthening VM Isolation with Integrity Protection and More,” Technical Whitepaper, 2020.
- [18] J. Demmel and H. D. Nguyen, “Parallel Reproducible Summation,” *IEEE Transactions on Computers*, 64(7):2060–2070, 2015.

- [19] IEEE, “IEEE Standard for Floating-Point Arithmetic,” *IEEE Std 754-2019*, 2019.
- [20] A. Gu and T. Dao, “Mamba: Linear-Time Sequence Modeling with Selective State Spaces,” *arXiv:2312.00752*, 2023.
- [21] S. Yang, B. Wang, Y. Zhang, Y. Shen, and Y. Kim, “Parallelizing Linear Transformers with the Delta Rule over Sequence Length,” *NeurIPS*, 2024.
- [22] C. Zhou et al., “Transfusion: Predict the Next Token and Diffuse Images with One Multi-Modal Model,” *arXiv:2408.11039*, 2024.
- [23] N. Shazeer et al., “Outrageously Large Neural Networks: The Sparsely-Gated Mixture-of-Experts Layer,” *ICLR*, 2017.
- [24] Zoo Labs Foundation, “Proof-of-Inference: Soundness, the Determinism Liveness Precondition, the Floating-Point No-Go, and the Routing/Chaining Extension Obligations,” companion soundness report, `zooai/proofs:proof-of-inference-soundness.tex`, 2026. Contains the per-result status table “What is proven on code versus on paper.”
- [25] Zoo Labs Foundation, Hanzo AI, and Lux Network, “Proof-of-Inference: End-to-End Emission and On-Chain Enforcement—the activation-trace commitment that makes *no real compute*, *no mint* a theorem,” companion emission report, `zooai/proofs:proof-of-inference-emission.tex`, 2026. Tabulates the cross-language implementation (Rust prover, Go verifier, Solidity witness/precompile) and the adversarial red-team battery.

Table 1: Built-vs-designed boundary for the execution-proof scheme. “Built” = merged, released under semantic versioning, and tested; the test counts are the actual passing counts of the suites named in the last column (Rust `cargo test`, Go `go test`, Solidity `forge test`). “Designed” = specified in this paper, not yet implemented. The decisive load-bearing item—emission of the activation-trace transcript from a live forward pass—is now in the *first* block, which is why the system-level “cannot be faked” claim is no longer conditional on an unbuilt wire (only on the stated cryptographic assumptions and the determinism contract).

Component	Status	Evidence / where
Exact-integer Freivalds matmul verifier	built , tested	<code>poi.rs</code> (7); <code>luxfi/crypto/poi</code> (23)
Activation-trace transcript (keccak Merkle) + challenger (open + Freivalds)	built , tested	<code>poi_transcript.rs</code> (8)
Proof-bearing forward: returned value = committed exact-integer result	built , tested	<code>poi_forward.rs</code> <code>ProvableLinear</code> (3)
Full-graph composition (attention, MoE, training) via typed trace	built , tested	<code>poi_graph.rs</code> <code>GraphTrace</code> (9)
Adversarial red-team battery (each forbidden attack is caught)	built , tested	<code>poi_adversarial.rs</code> (11)
Soundness theorems ($1/p$ one-layer; 2^{-122} amplified; zero false-reject; composition)	proven	§12; [24, 25]
Lowest-token-id argmax / canonical MoE top- k (determinism D3)	built , tested	<code>sampler.rs</code> , <code>ops.rs</code> , <code>poi_graph.rs</code>
On-chain Freivalds + Merkle witness (pure Solidity)	built , tested	<code>ComputeWitnessLib.sol</code> (5)
Optimistic commit/challenge/slash gate (bond paid on verified fraud only)	built , tested	<code>OptimisticEvidence.sol</code> (18)
Fail-closed mint gate (no compute-verifier $\Rightarrow$ no mint)	built , tested	<code>AIMiner.sol</code> (12); <code>ComputeProofEnforcement</code> (19)
Capped, halving fair-mined coin; global no-double-mint	built , tested	<code>AICoin.sol</code> (11); <code>AIComputeRegistry</code> (6)
Miner bond/slash registry	built , tested	<code>MinerStakeRegistry.sol</code> (8)
ZOO L3 directory + endorsement (settlement registry)	built , tested	<code>L3Registry.sol</code> (8); §18
RepOps pinned-reduction-tree fp16/bf16/fp8 mode	designed ; proven	§16.3; [24]
Emission coverage of architecture families beyond the tested set	designed ; proven	§17; [24]
Sequential-chaining proof for SSM / diffusion latents (the <i>chaining</i> obligation)	designed ; proven	§17.3, §17.6; [24]

Metric	PoAI	Tendermint	PoW (Ethereum)
Throughput (TPS)	847	1,000	15
Finality (seconds)	2.3	1.0	360
Energy (kWh/tx)	0.0002	0.0001	50
Useful work (%)	98%	0%	0%

Table 2: Consensus comparison. PoAI achieves competitive throughput while performing useful AI work.

Task Type	Compute (ms)	Verify (ms)	Overhead
Inference (7B)	150	12	8.0%
Inference (70B)	2,400	45	1.9%
Training step	850	23	2.7%
Extraction	1,200	34	2.8%

Table 3: Verification overhead. TEE attestation adds minimal latency.

Attack Type	Detection Rate	False Positive
Random outputs	100%	0%
Stale outputs (replay)	98.7%	0.3%
Sybil quality voting	99.2%	0.5%

Table 4: Attack detection. Multi-party scoring resists manipulation.

Table 5: Bitcoin and Proof-of-AI at the verification layer. The execution proof of this paper is the *designed* cognitive counterpart of the recomputable hash, software-only and over useful work; the Proof-of-AI column is the specification, not a deployed system (only the Freivalds verifier row is built today).

	Bitcoin (PoW)	Proof-of-AI (this paper)
The recomputable value	the block-header SHA-256d	the activation-trace Merkle root + output hash (<i>emitted by ProvableLinear/GraphTrace</i>)
What it proves	$\sim 2^d$ hashes were tried	the genuine matmuls were performed (Thm. 12.2; emitted and verified end to end)
Verifier cost	$O(1)$: one hash	$O(n^2)$ Freivalds, or $O(\log L)$ bisection to certainty
Hardware to verify	any CPU	any CPU (Cor. 12.3)
Hardware to <i>produce</i>	ASIC (centralizing)	any AI accelerator <i>or</i> CPU; no TEE required
Unforgeability root	preimage resistance of SHA-256	Freivalds soundness + keccak collision resistance
Freshness / anti-replay	previous block hash in header	<code>openBlockHash</code> beacon in the challenge (Lem. 10.2)
Product of the work	none (discarded)	the answer / embedding / improved weights

Table 6: Tier $\rightarrow$ proof matrix. One profile gate (Section 14.3) maps a tier to its required **proofType**; where a tier is opted in, no valid proof of that type means no mint and no settlement. The gate is built but defaults to permissive (**requiredProofType** = 0), and the software-only backend (T2/T3) awaits the emission and challenger of Section 19; the matrix is thus the configured policy surface, live enforcement following the build.

Tier	Proof system	Trust assumption	Use
T0	zkML (type 2)	none beyond the SNARK’s	tiny models; trustless, on-chain succinct verification
T1	M -of- N heterogeneous CC-TEE (type 4)	no single TEE vendor; M of N enclaves agree	high-value tasks; defense in depth across NVIDIA CC + Intel TDX + AMD SEV-SNP
T2	single-vendor CC-TEE (type 1) <i>or</i> optimistic-Freivalds (type 3)	one TEE vendor, <i>or</i> one honest challenger + a bond	the workhorse tier; software-only Freivalds is the intended egalitarian default (pending emission + challenger)
T3	redundancy + plurality + slash	honest plurality among bonded operators	liveness-first; CPU-only providers; the weakest proof, fastest path

Table 7: Coverage by arithmetic. Every modern-LLM arithmetic is given a proof-bearing mode by tiering the verification to the arithmetic, not by a single floating-point Freivalds. The quantized row is the egalitarian, software-only, hardware-free default and is *built and tested end to end*: it runs the whole- K exact-integer accumulator (**exact_matmul**) that gives the $\epsilon = 0$, zero-false-reject property, emits a transcript (**ProvableLinear/GraphTrace**), and verifies on-chain (**ComputeWitnessLib**). The floating-point (RepOps) row is *designed* (proven sound in [24]), awaiting the pinned-reduction kernel; the TEE row is the hardware-trust fallback (Section 2).

Arithmetic	Verification mode	Property
int8 / int4 / GGUF / AWQ / GPTQ / MLX-quant	exact-integer Freivalds (Section 12)	$\epsilon = 0$, software-only, no hardware; soundness $1/p$ per vector (2^{-122} dense); zero false-reject—for a pinned whole- K -i32 kernel (Rem. 12.2), to be built; not the shipping block-f32 kernel
fp16 / bf16 / fp8	RepOps-Freivalds (pinned reduction tree τ ; Section 16.3) <i>or</i> TEE/CC attestation (type 1/4)	RepOps: bit-exact via Prop. 16.1, costs vendor-kernel speed. TEE: hardware attests the fp run, no Freivalds, no τ pin
small models (any arithmetic)	zkML (type 2; Section 14)	trustless succinct verification; bounded by proving cost to small models

Table 8: The ZOO model zoo mapped to proof obligations. Beyond base Freivalds-on-matmuls there are exactly two new core techniques (routing proof, sequential-chaining proof); the rest is per-modality commitment and deterministic pre/post-processing. “Built” = implemented and tested in the engine; “designed” = specified here, not yet built.

Family (zoo models)	Heavy arithmetic	New core technique	Modality / determinism add-ons
Dense transformer (zen-nano, dense tiers)	GEMMs (base)	none	—
MoE (zen4-max, zen-coder-flash, zen-max)	expert-FFN GEMMs	routing proof (<i>built</i>); routing determinism <i>built</i>	commit router logits, expert indices, sparse pattern
Linear / gated-recurrent / SSM (zen4-nano/micro/mini/pro, zen4-thunder/storm)	per-step GEMMs	sequential chaining (designed)	bind consecutive states $S_{t-1} \rightarrow S_t$
Multimodal ViT (zen-vl, zen-omni, zen-designer)	tower GEMMs; conv as im2col→GEMM	none (MoE-VL: routing proof)	deterministic image preprocessing; preprocessed-pixel commit
Audio / speech (zen-tts, zen-omni audio)	conformer GEMMs	none	deterministic FFT→mel (commit-and-recompute); audio I/O commit
Diffusion / Transfusion (zen5 3D/video, zen-designer image-gen)	per-step GEMMs	sequential chaining (designed, <i>reused</i>)	pinned noise schedule, sampler, seed; latent chaining